

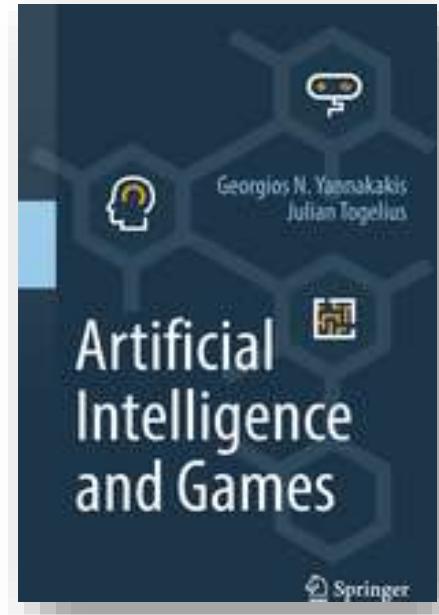
Artificial Intelligence and Games

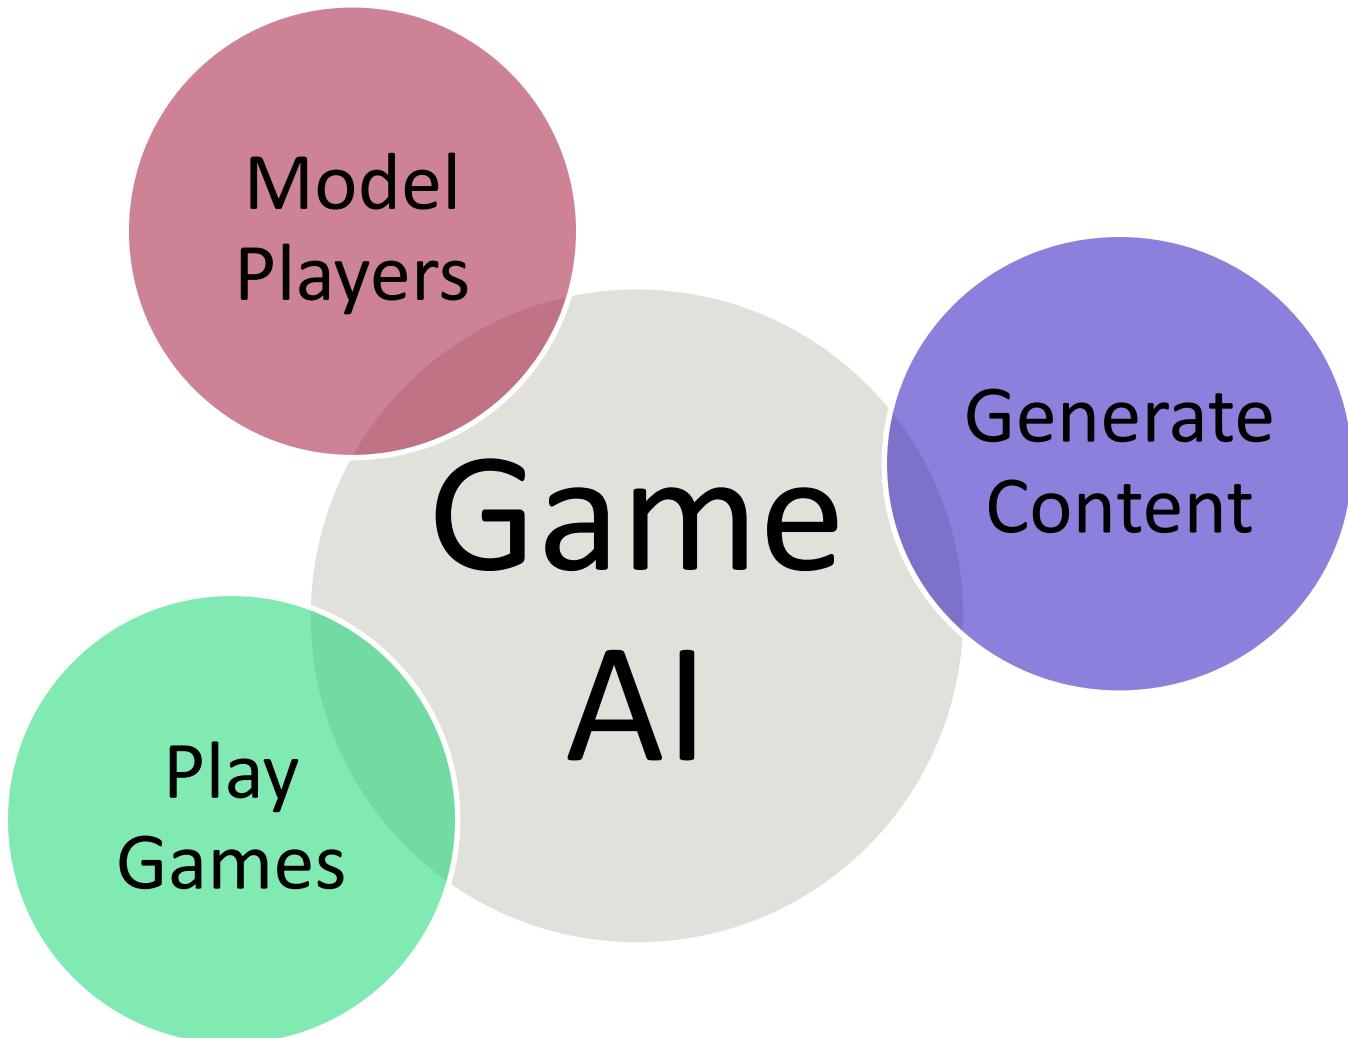
Generating Content

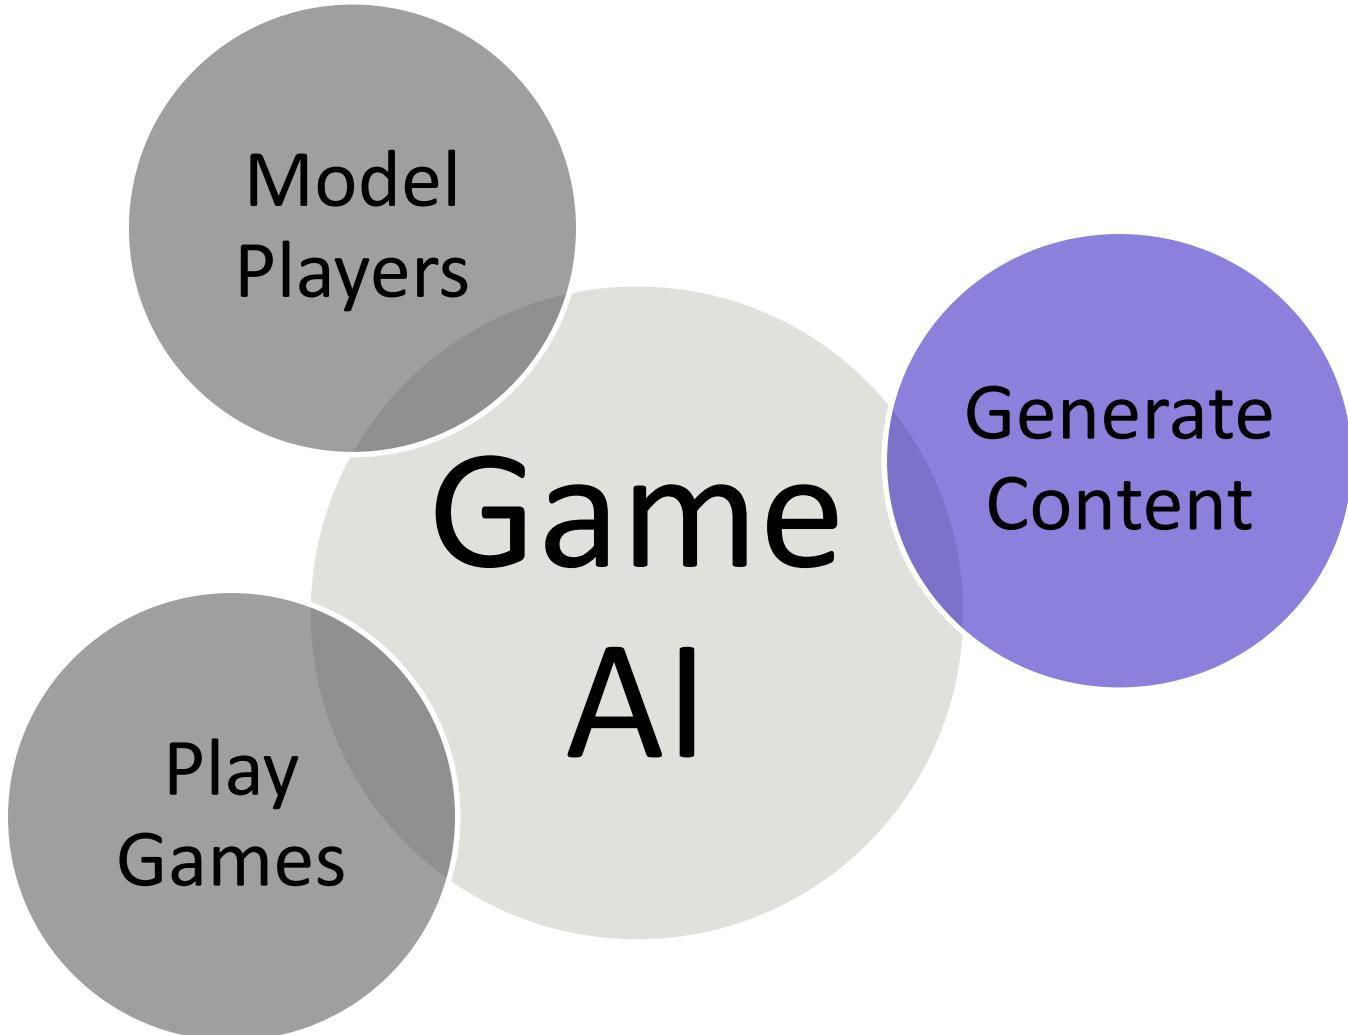


Georgios N. Yannakakis
@yannakakis

Julian Togelius
@togelius







Artificial Intelligence and Games

A Springer Textbook | By Georgios N. Yannakakis and Julian Togelius



Springer

About the Book Table of Contents Lectures Exercises Resources

About the Book

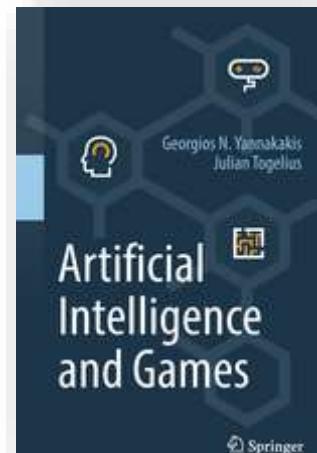
Welcome to the Artificial Intelligence and Games book. This book aims to be the first comprehensive textbook on the application and use of artificial intelligence (AI) in, and for, games. Our hope is that the book will be used by educators and students of graduate or advanced undergraduate courses on game AI as well as game AI practitioners at large.

Final Public Draft

The final draft of the book is available [here!](#)

Your readings from **gameaibook.org**

Chapter: 4

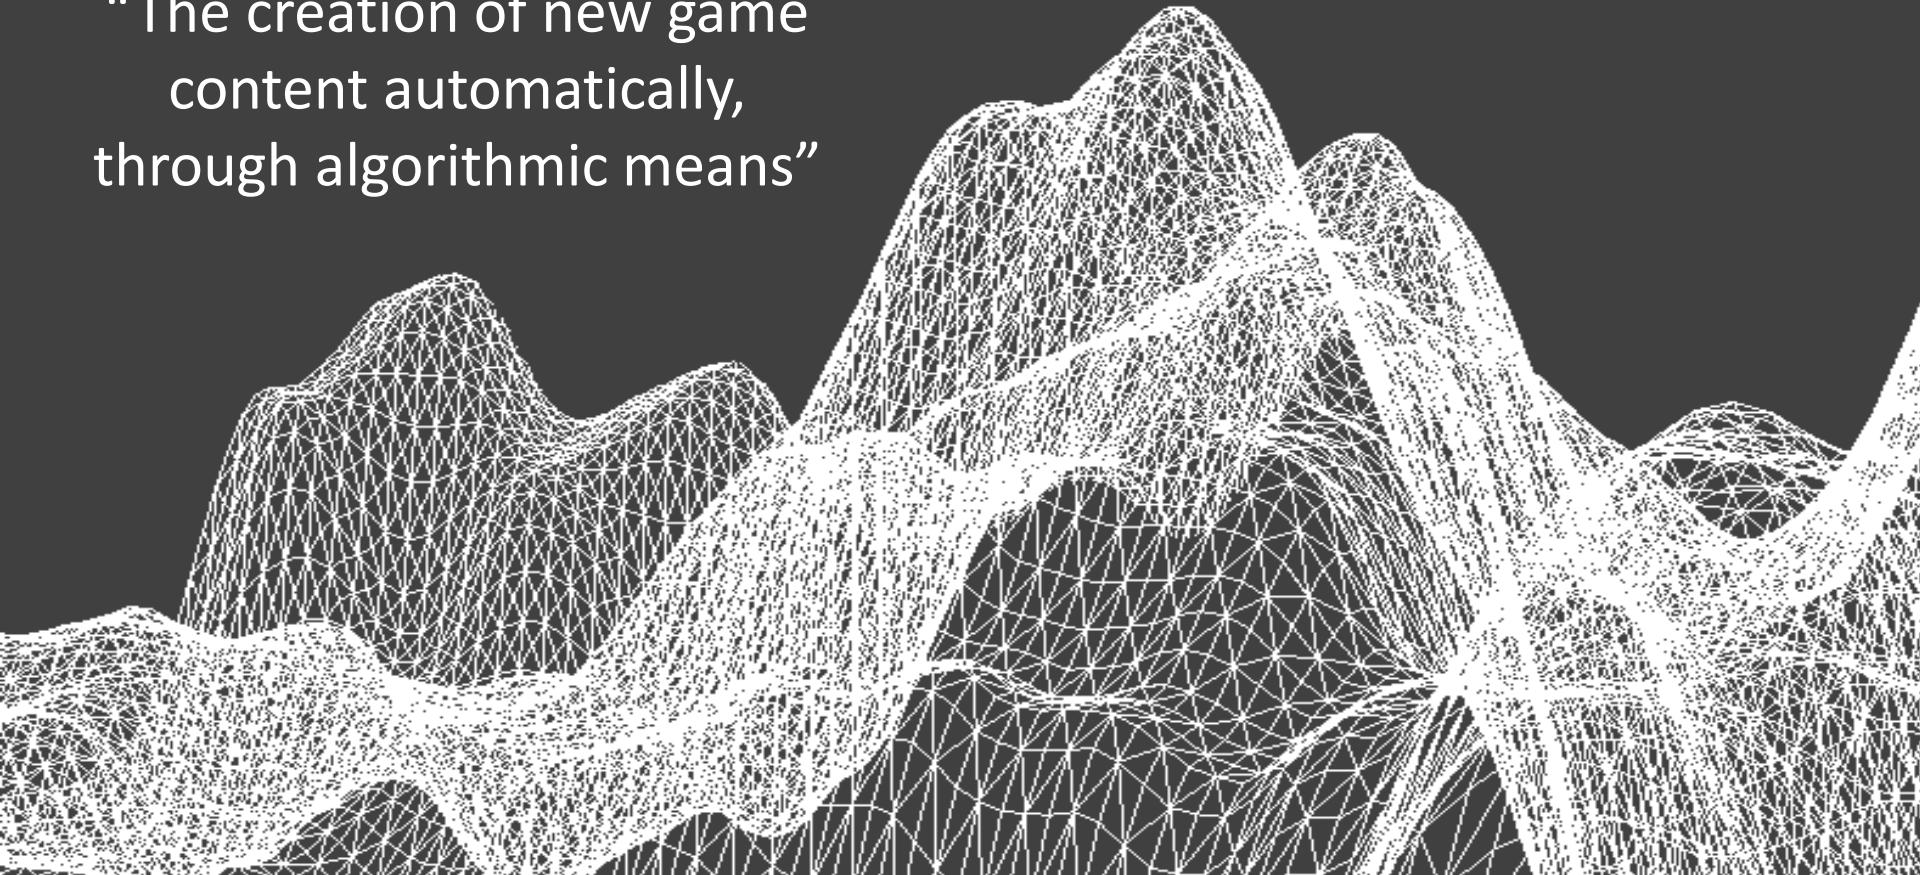


Overview



- Procedural Content Generation (PCG)
 - What is it?
 - Why we need it?
- Constructive Approaches
- Search-Based PCG
- Machine Learning PCG
- Mixed-Initiative PCG
- Experience-driven PCG

“The creation of new game content automatically, through algorithmic means”



What is **Procedural Content Generation?**



What is Game Content?



- Content can be:
 - NPC behavior (aspects)
 - Quest/story/narrative
 - Camera profiles
 - Audiovisual settings
 - Levels/maps/tracks
 - Items
 - Game mechanics
 - Reward schedules
 - ...
 - **Everything** together?
- Content is the game context
- Content has differing quality (ability to elicit player experience)



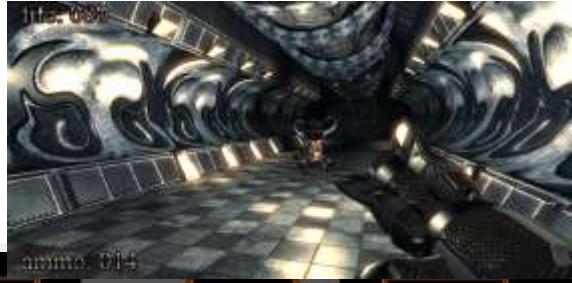
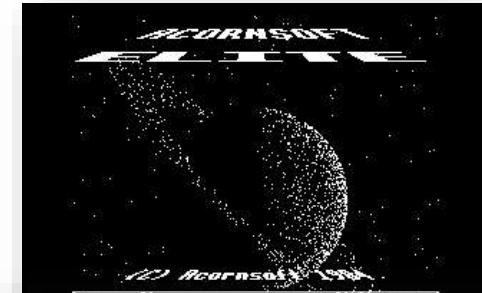
COURTESY BRICKARTIST.COM

Content Types



- Maps
- Levels
- Weapons / items
- Game rules
- Stories
- ...
- Everything together?

PCG in Industry



PCG in Academic



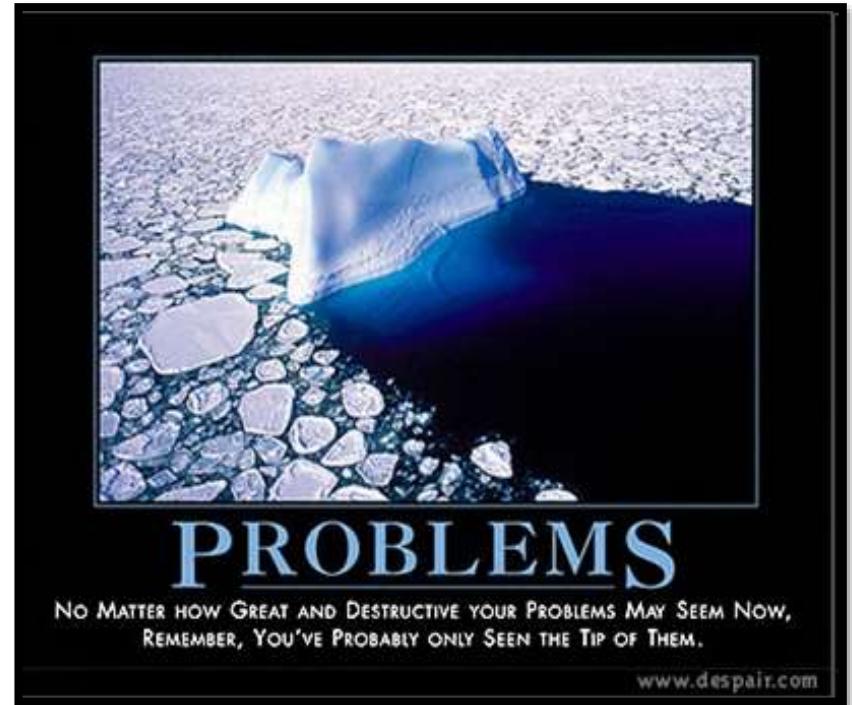
- An IEEE Task Force
- A book: pcgbook.com
- A number of PCG dedicated sessions/workshops since 2010:
FDG, IEEE CIG, AIIDE, AAAI, IJCAI,
- A paradigm shift:
 - ML-based PCG
 - Mixed-Initiative PCG
 - Computational Game Creativity
 - Experience-driven PCG

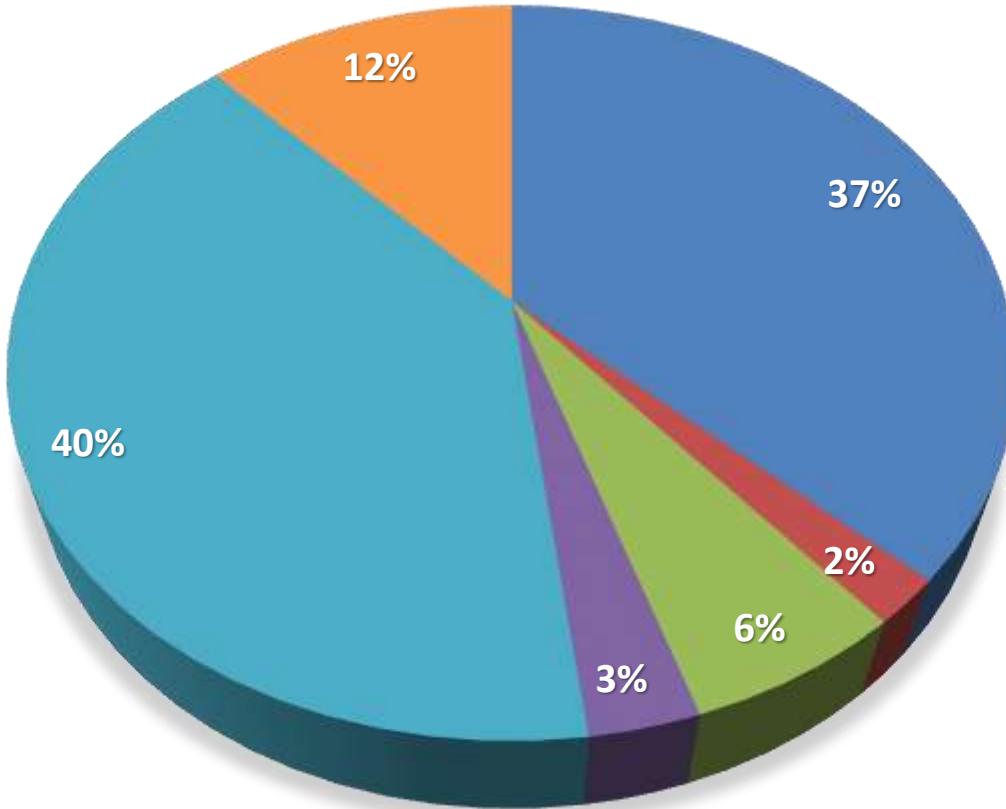
What is the Problem?



- Content costs! (money and time)
- Games end when finished
- Game worlds have bounds
- Human imagination (creativity) is limited
- Designer is not always present

In sum, there is **content shortage**



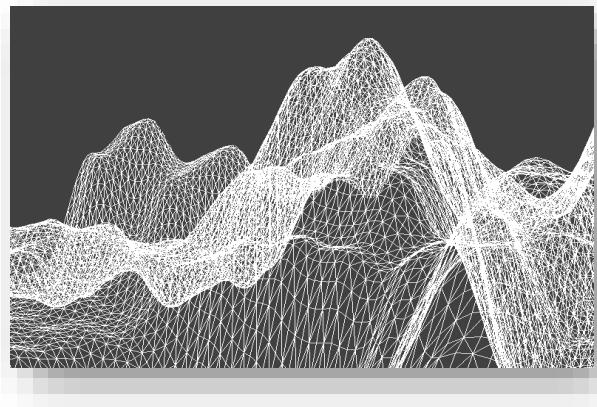


■ Art ■ Manufacturing ■ Other ■ Debugging ■ Marketing ■ Programming

What Can PCG Do?



- Can we drastically cut game development costs by creating content automatically from designers' intentions?
- Can we create games that adapt their game worlds to the preferences of the player?
- Can we create endless games?
- Can the computer circumvent or augment limited human creativity and create new types of games?
- Can we understand game design through formalising the design process?



What Are the Trade-offs

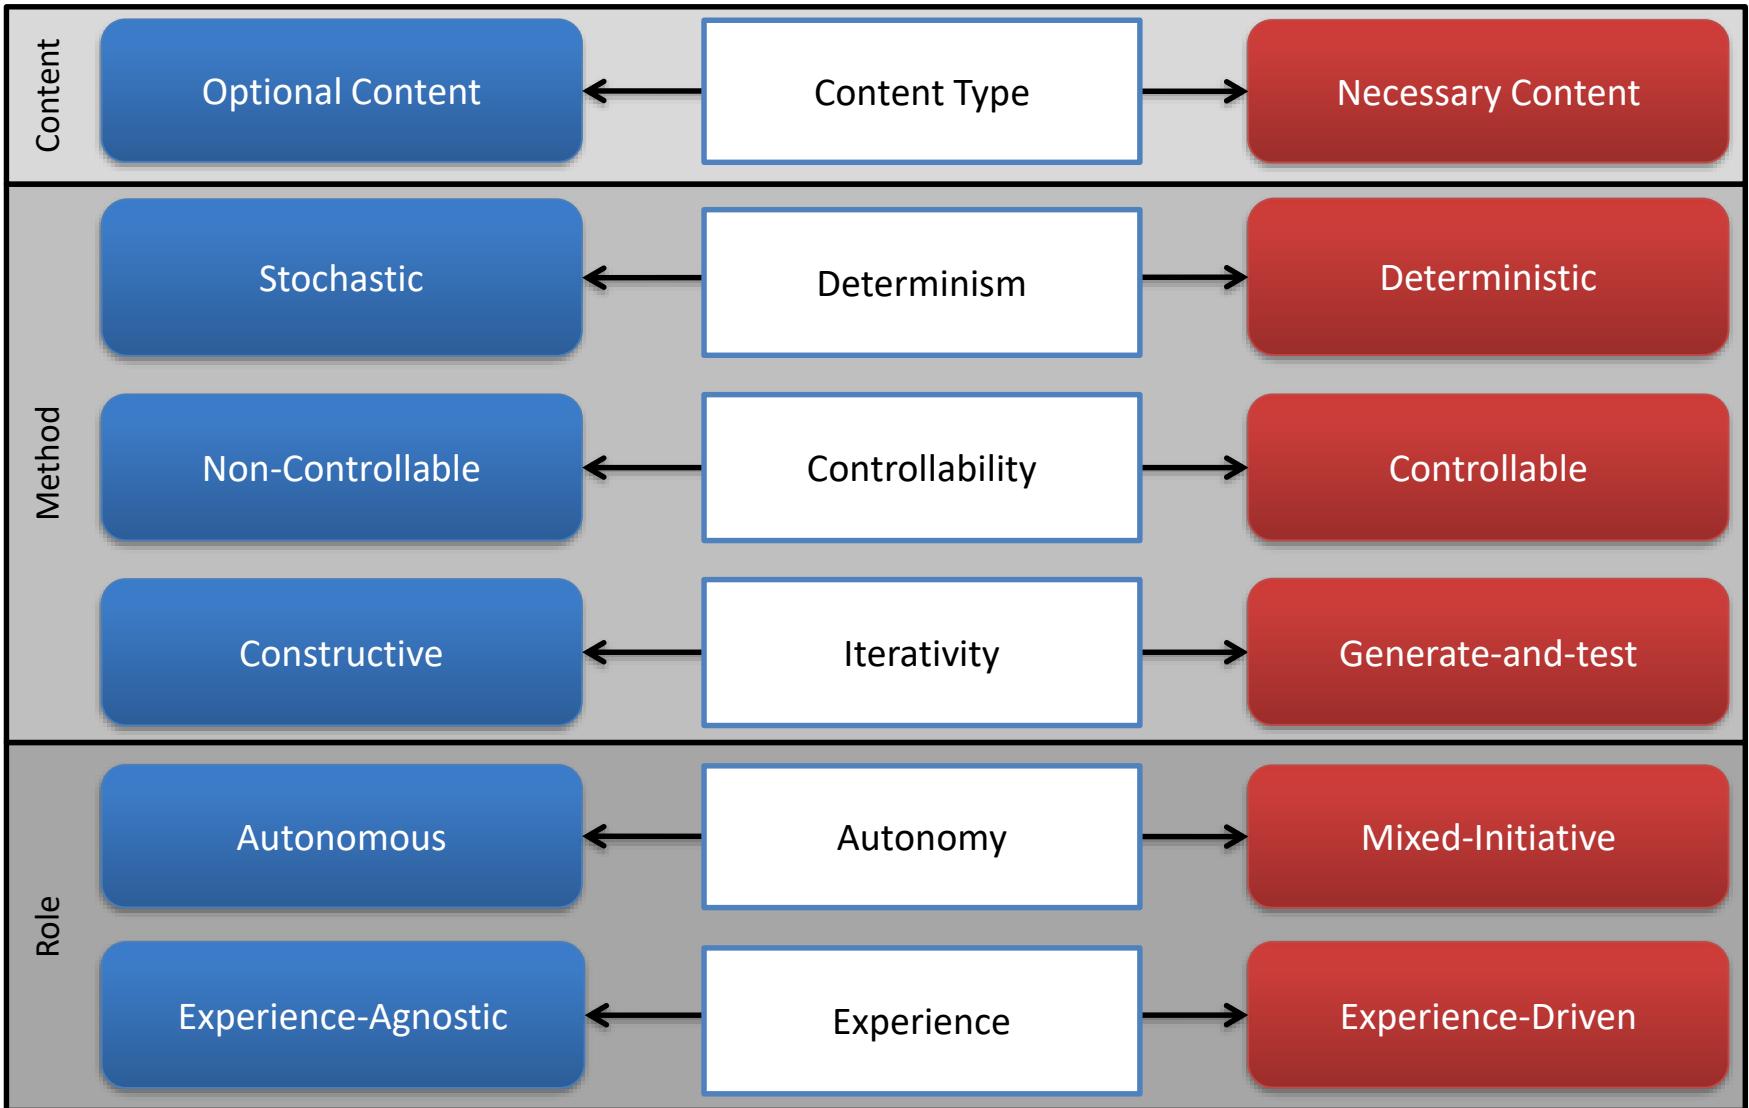


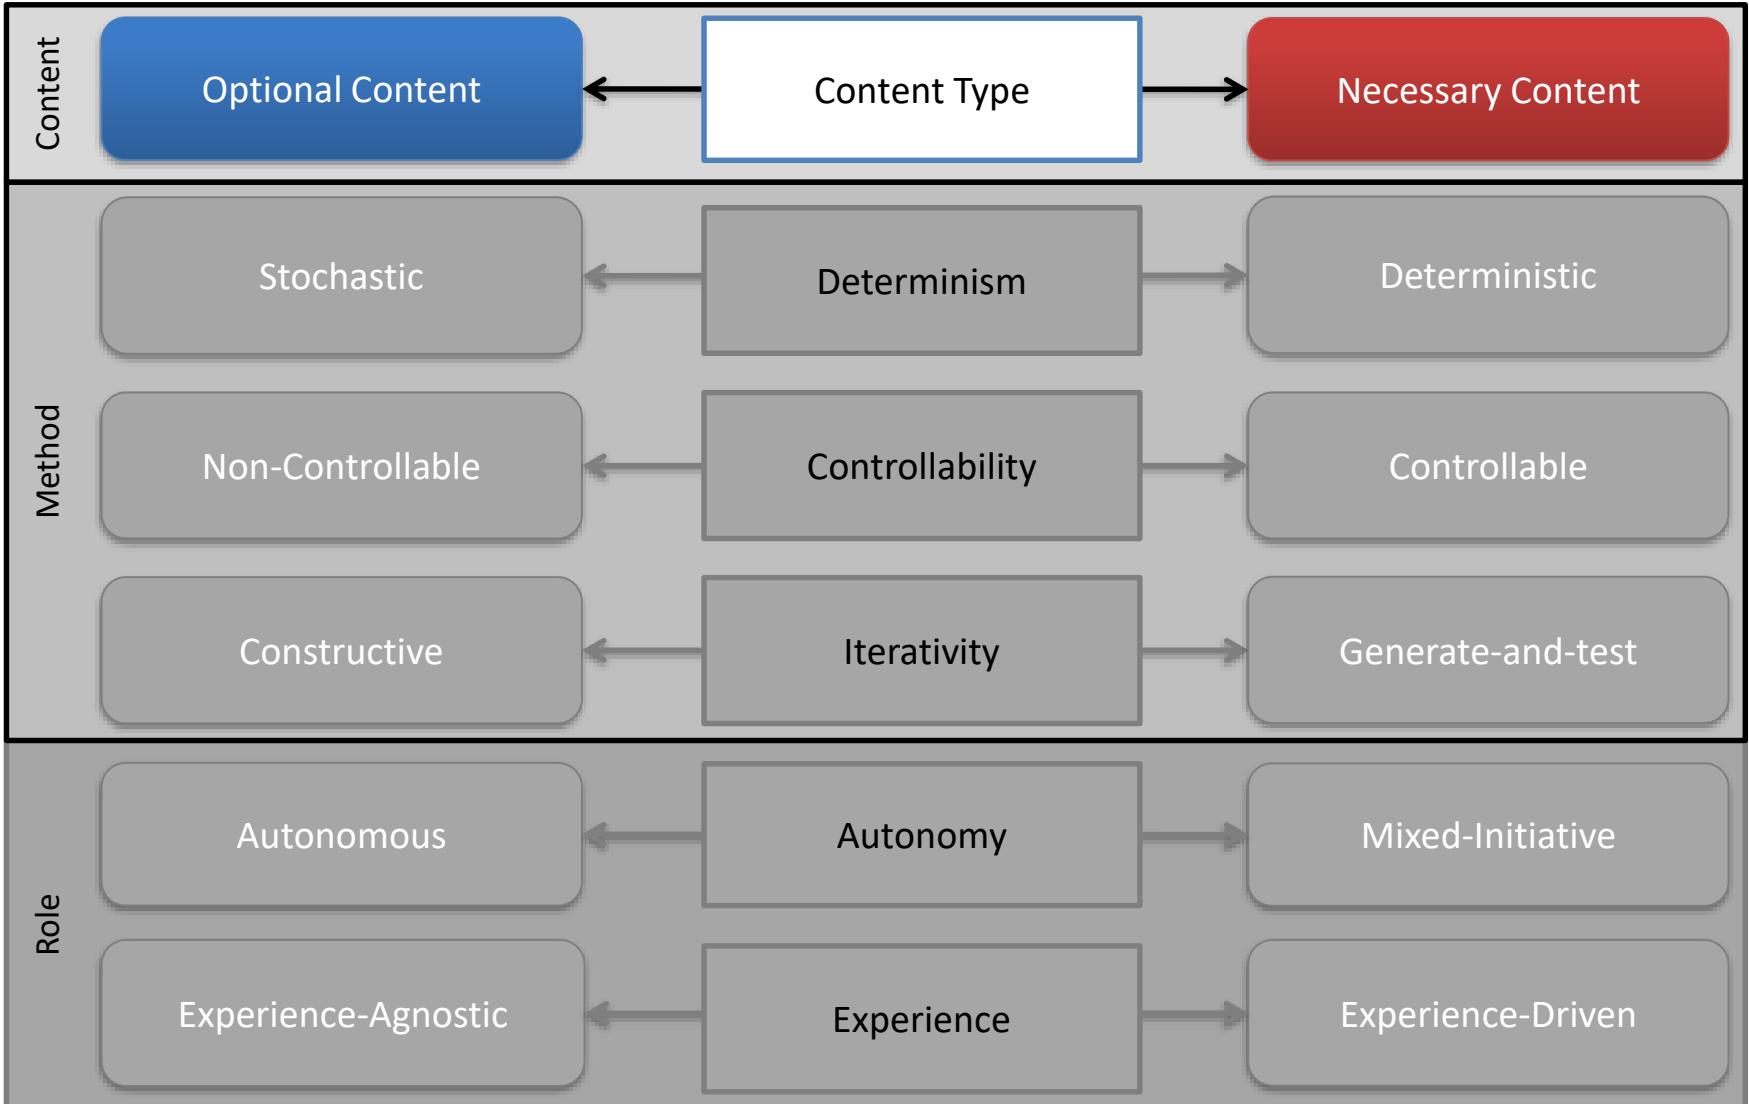
- *Speed*
Real-time? Or design-time?
- *Reliability*
Catastrophic failures break gameplay
- *Controllability*
Allow specification of constraints and goals
- *Diversity*
Content looks like variations on a theme
- *Creativity*
Content looks “computer-generated”

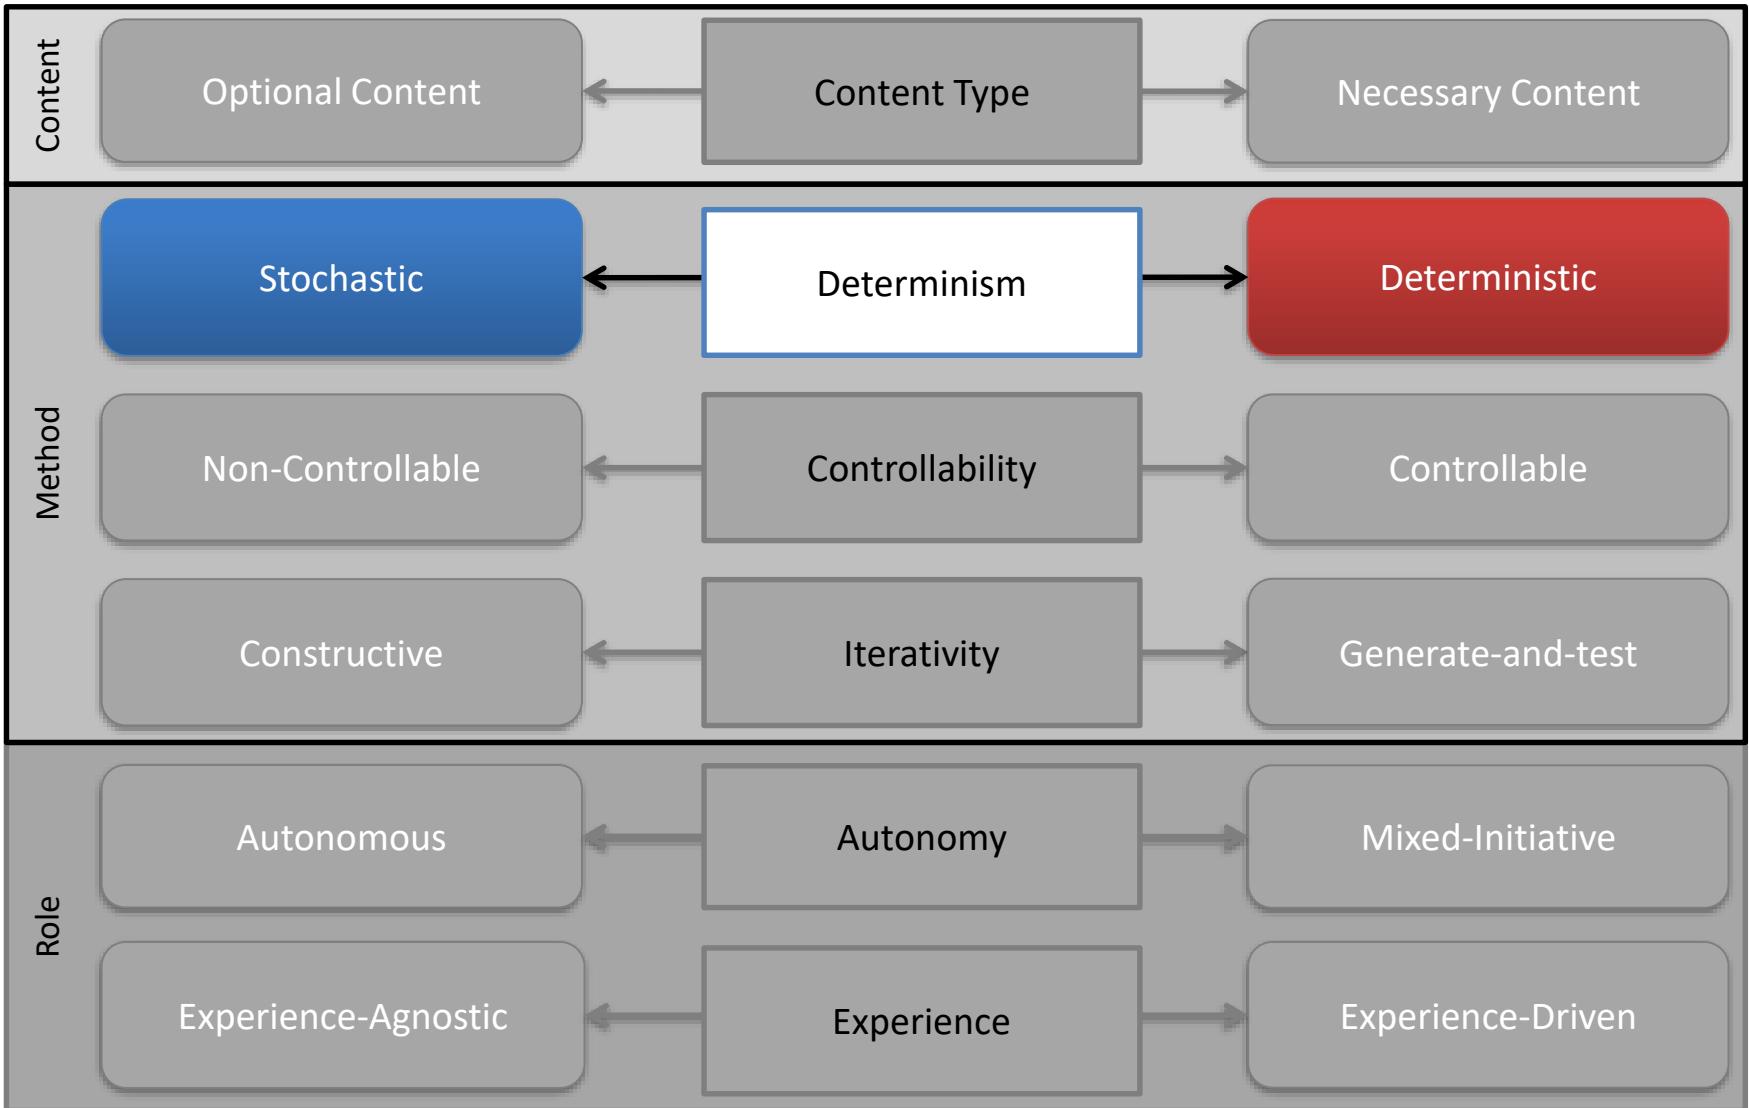


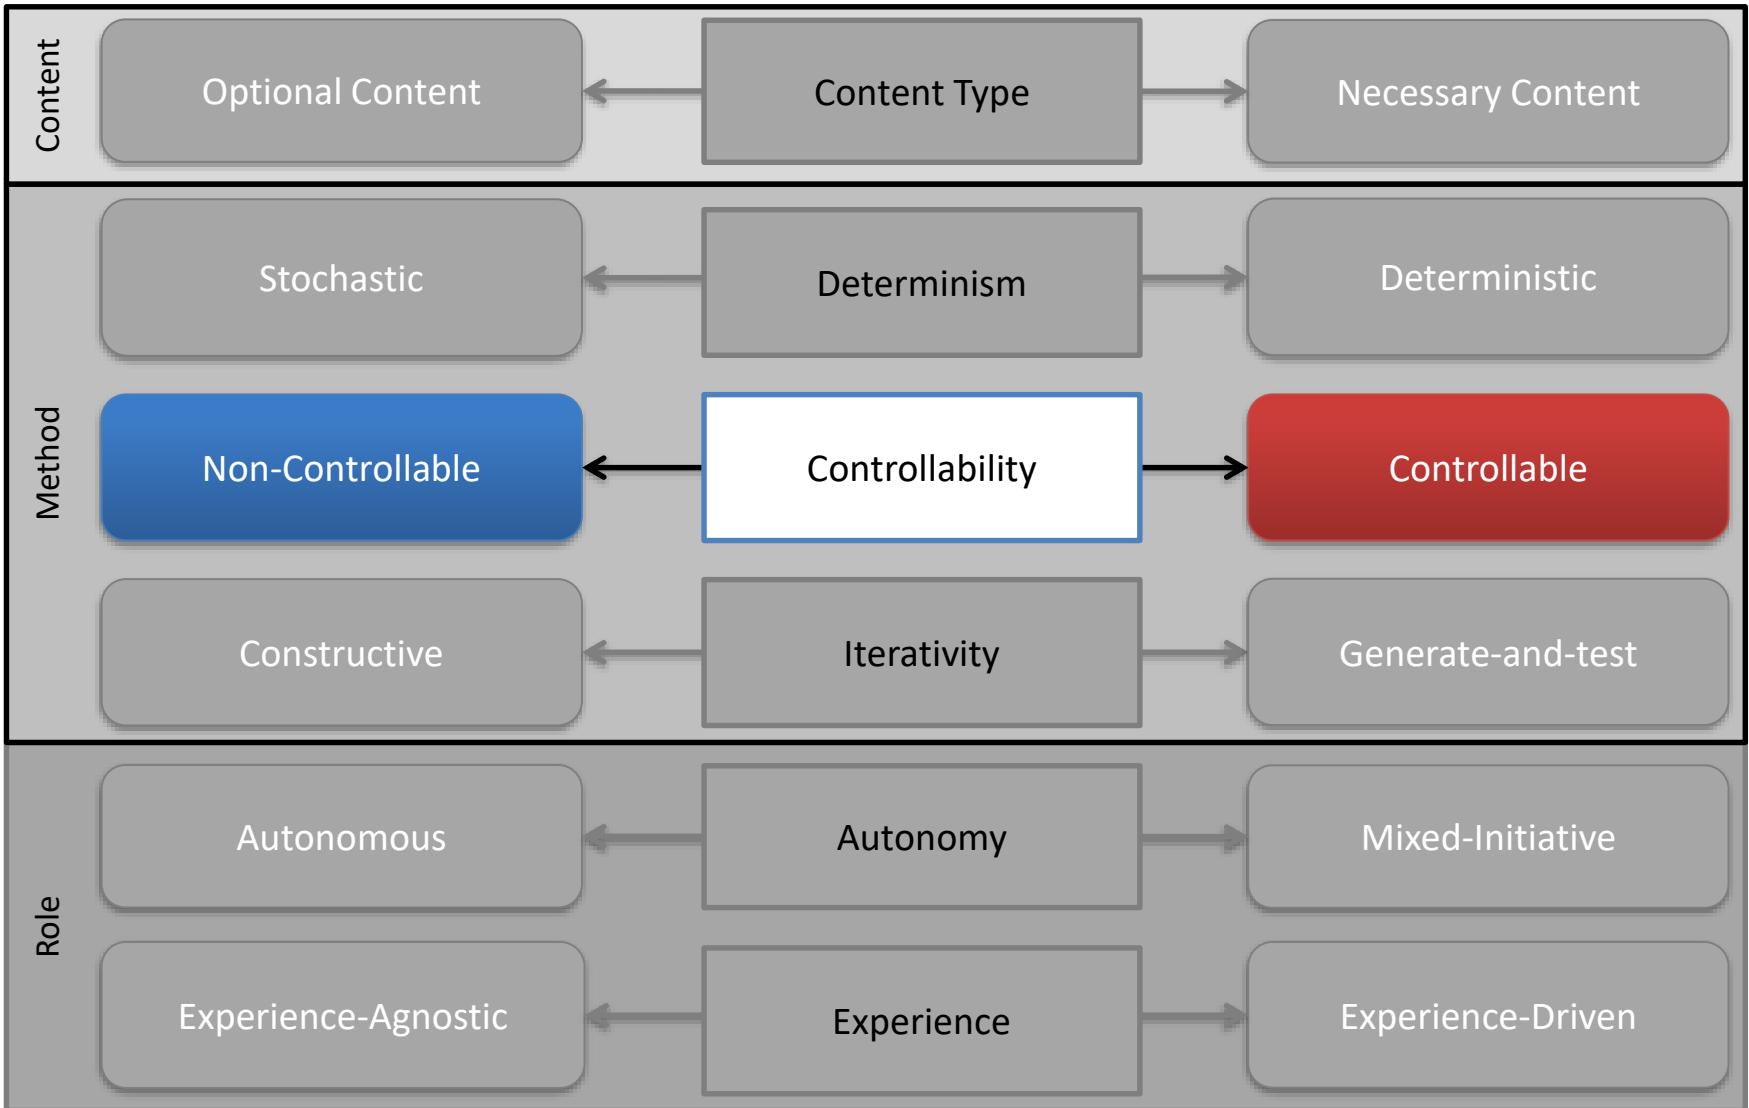
A PCG Taxonomy

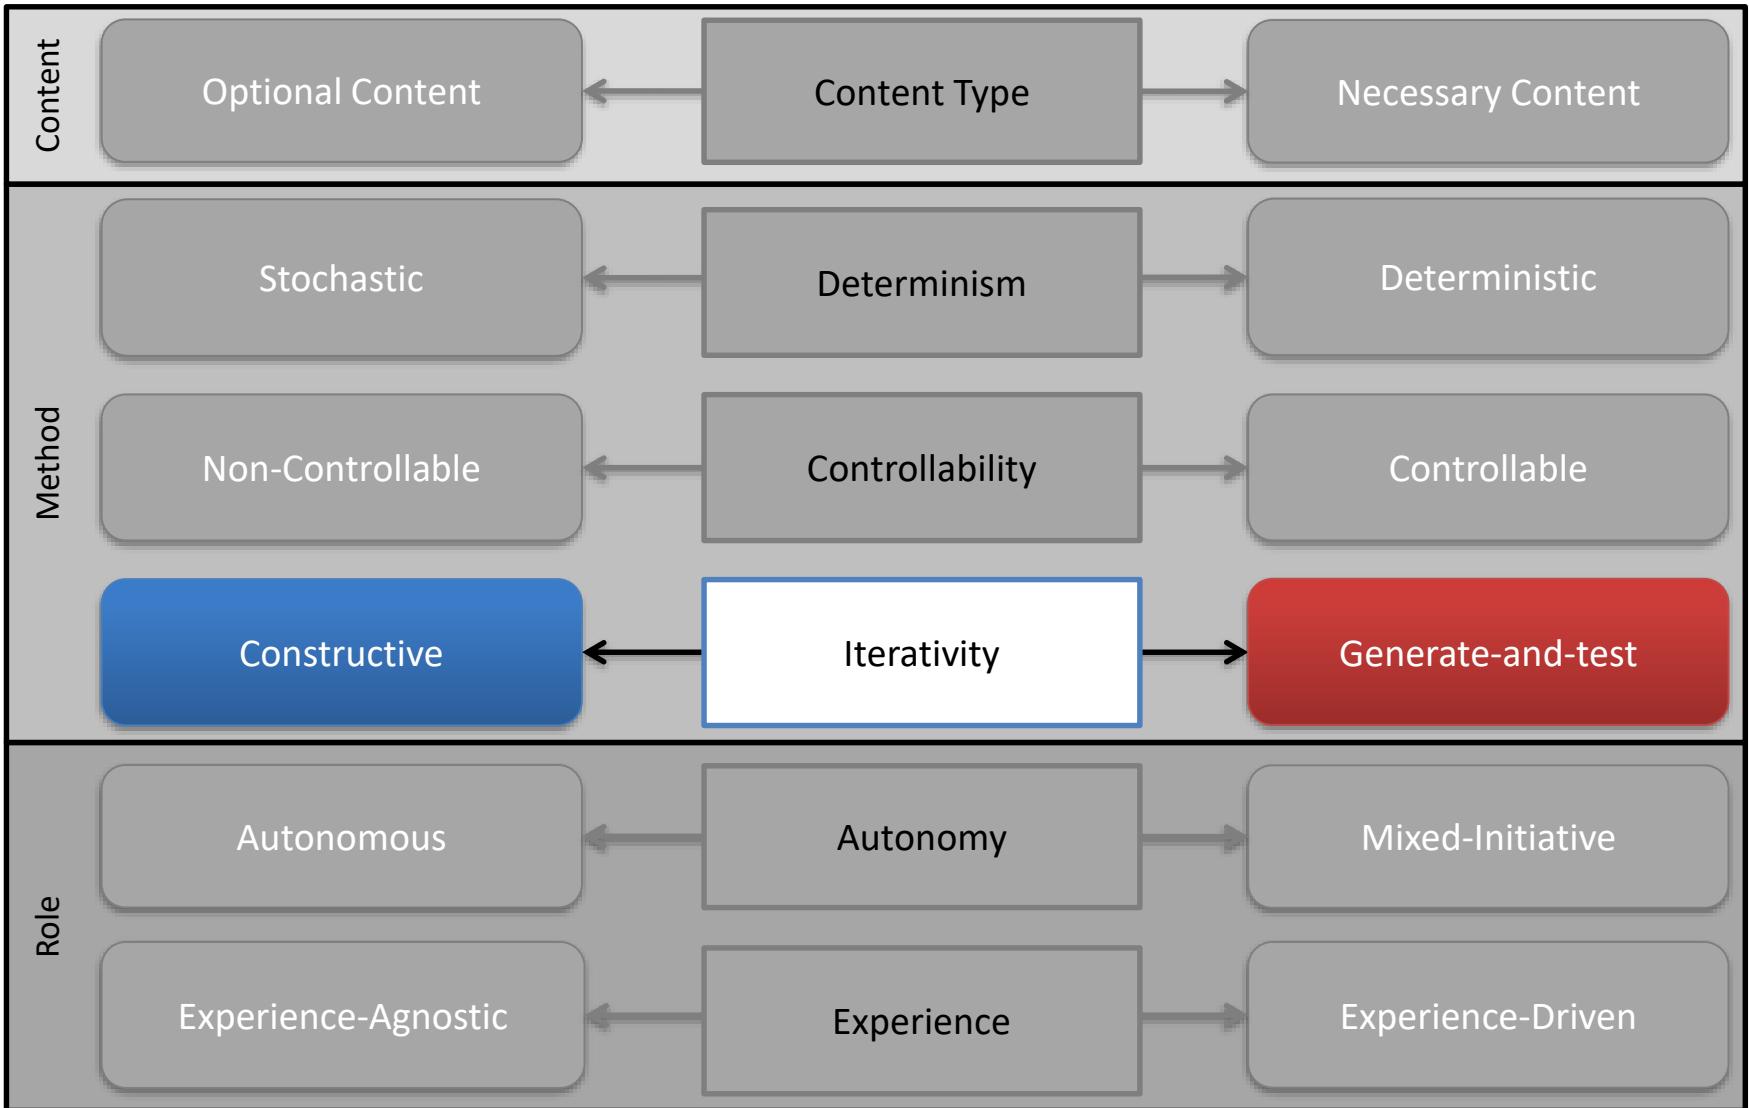




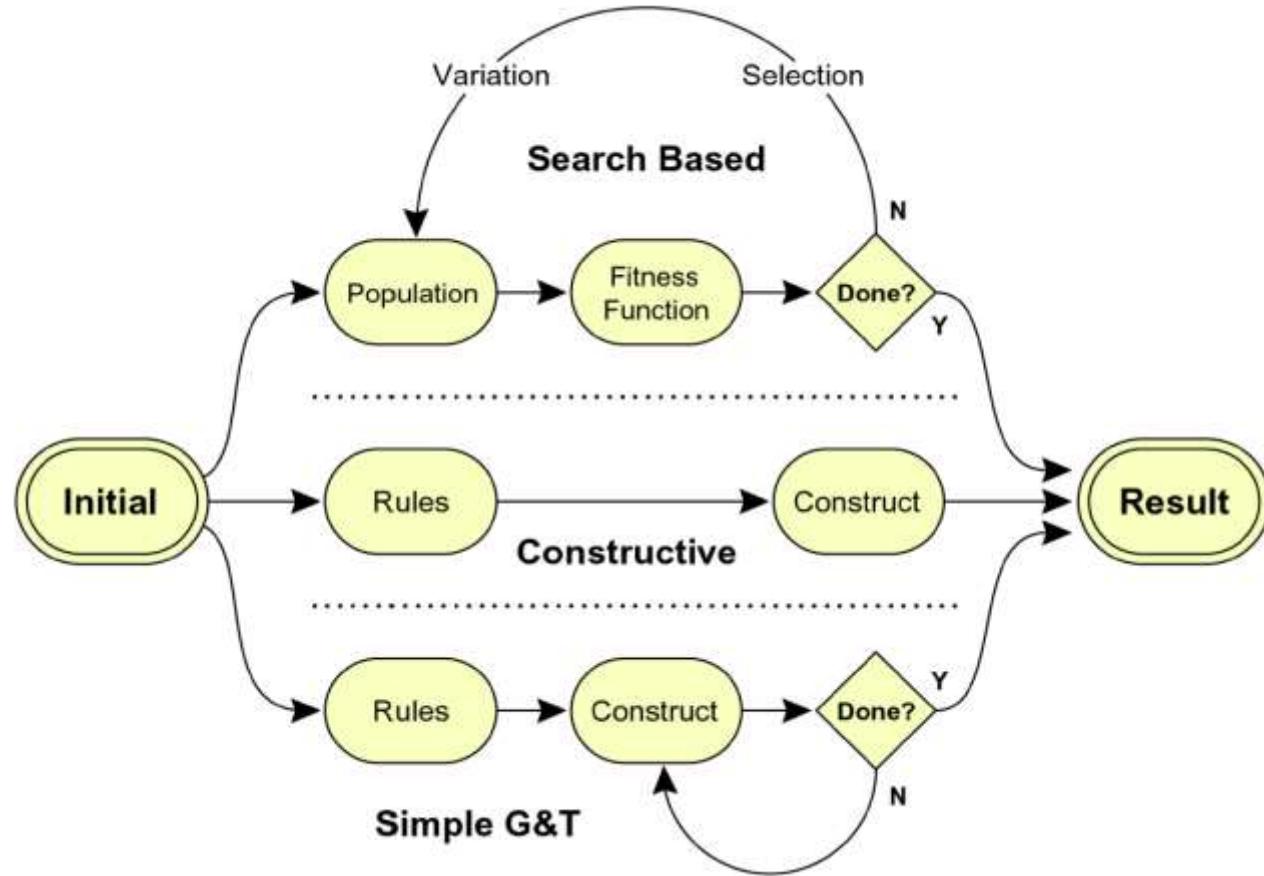








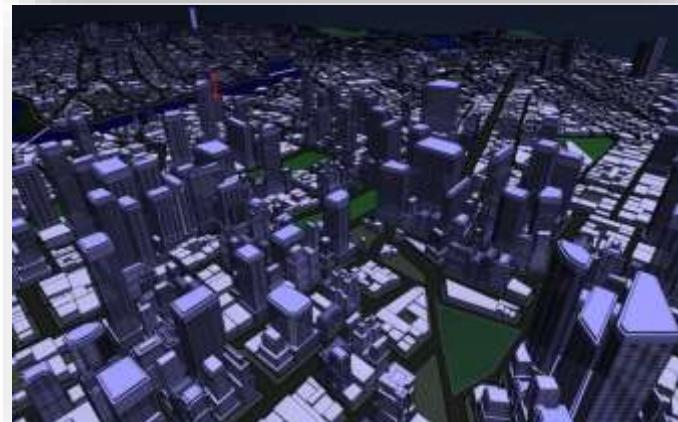
Three Main PCG Approaches



“Classical” PCG Approaches



- Mostly constructive, stochastic, non-controllable, random seed, online
- Mostly optional (non-vital) content such as decoration and redundant items
- Interesting examples from classic games:
 - Elite (deterministic, offline)
 - Rogue (online, crucial content)
- Techniques: L-trees, Fractals, ...



How Could we Generate Content?



PCG Methods

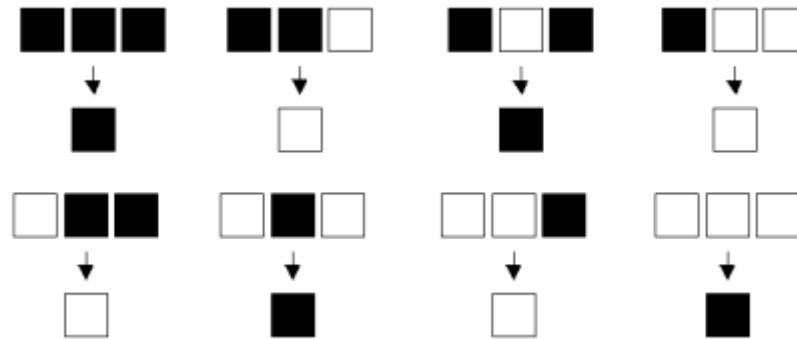


- Cellular automata
- Solver-based
- Grammar-based
- Search-based
- Machine learning
- Noise and fractals
- Ad-hoc constructive methods

Cellular Automata

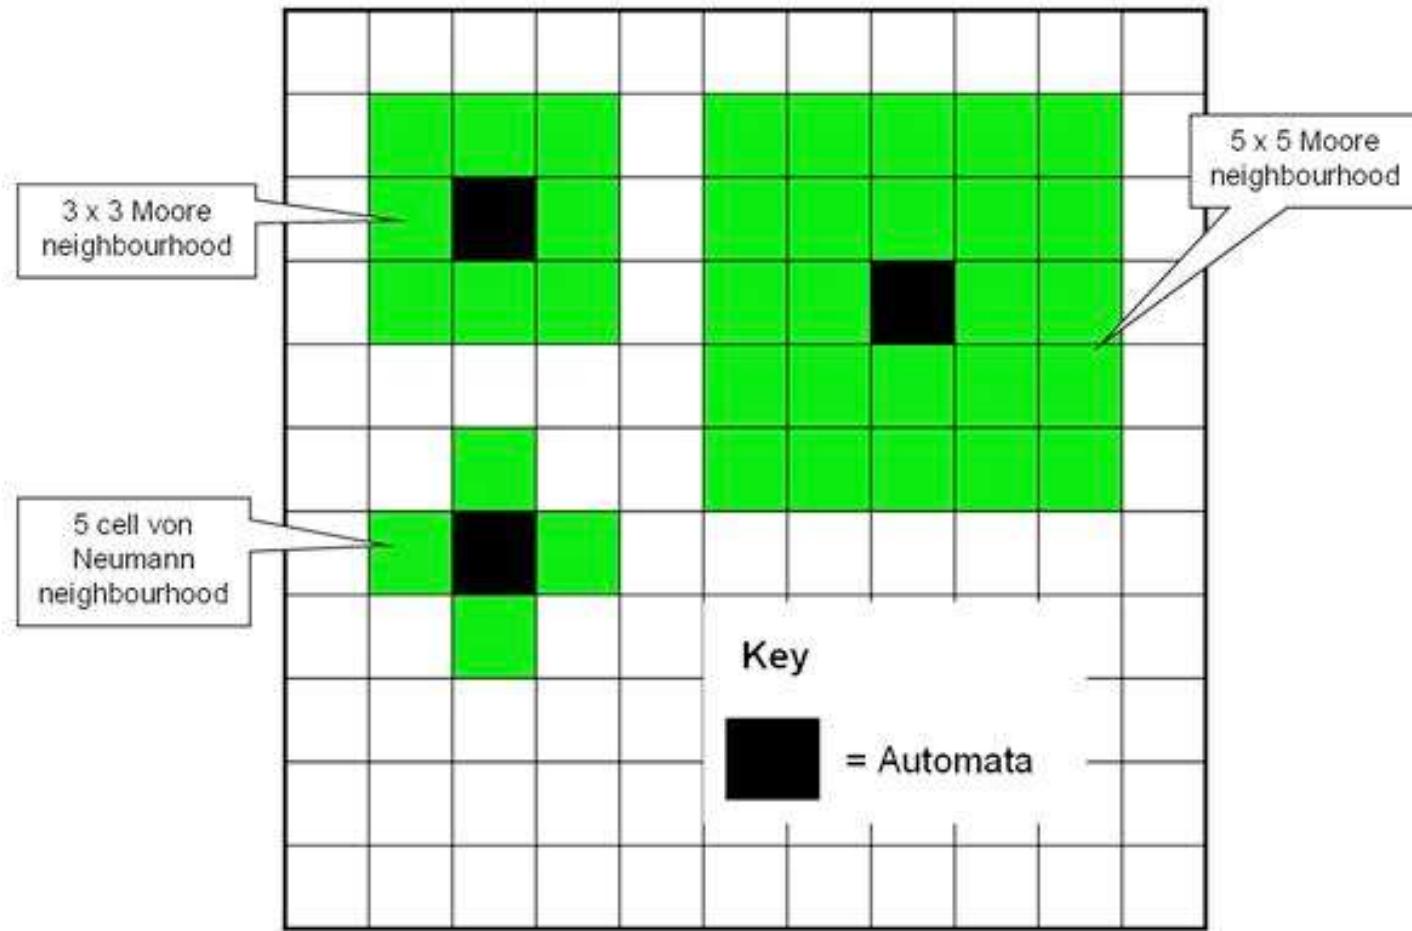


Cellular Automata



- Computational paradigm based on local interaction
- Used in artificial life and complexity studies
- The value of each cell in iteration $n+1$ is based on the value of neighboring cells in iteration n and some rule

2D Cellular Automata



An Example



Cellular automata for real-time generation of infinite cave levels

Lawrence Johnson, Georgios N. Yannakakis and Julian Togelius

FDG PCG Workshop 2010

This...



A CA-based algorithm for generating infinite 2D caves

- simple
- real-time
- looks good
- *somewhat* controllable



The Motivation



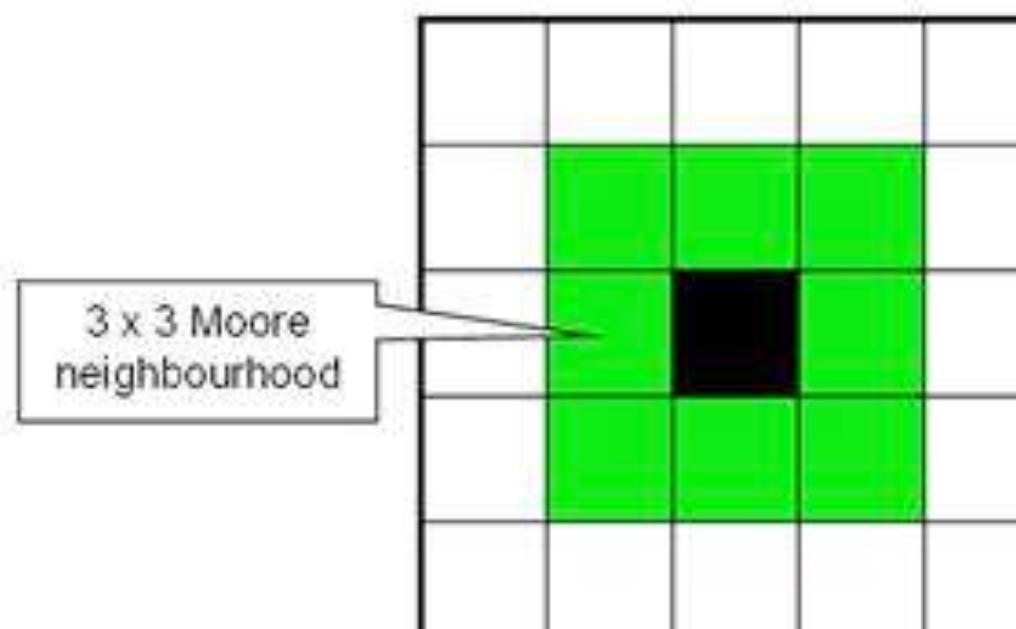
- *Cave Crawler*: a cooperative *abusive* dungeon crawler
- Never ends – therefore needs to produce infinite caves...

CA Cave Generation



- Start with a square grid (e.g. 50×50) – all floor
- Randomly switch a proportion of cells from floor to rock
- Run a CA n times, where each cell is set to:
 - Rock: if at least T neighbors are rock
 - Floor: otherwise
- Fill in the interior of rock formations

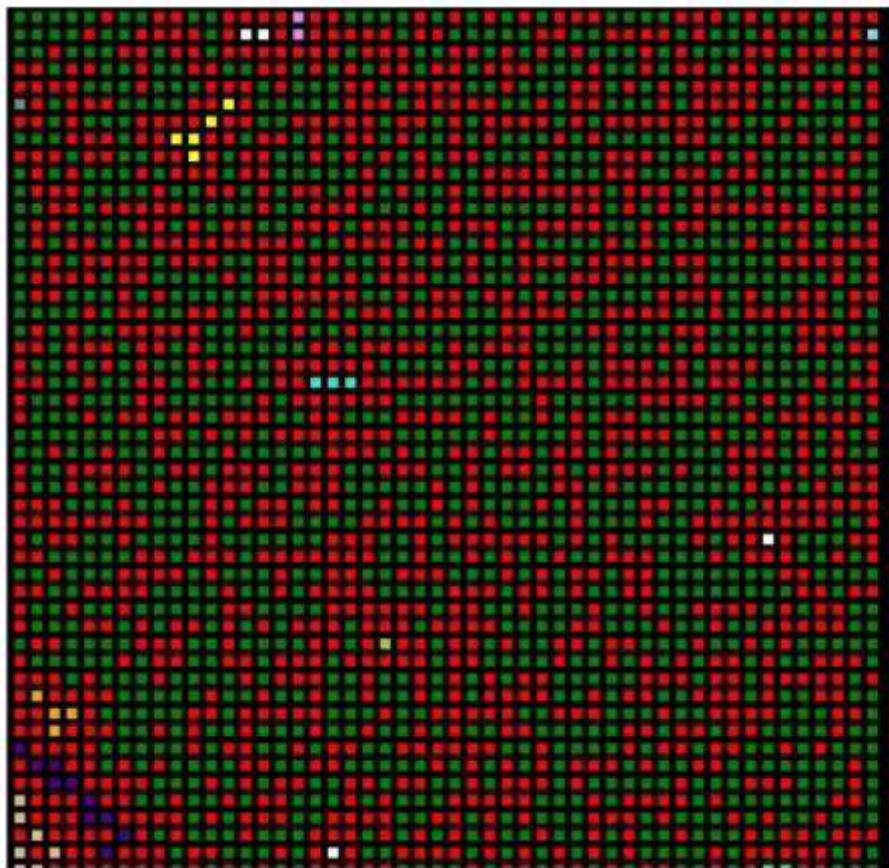
Core CA Mechanic



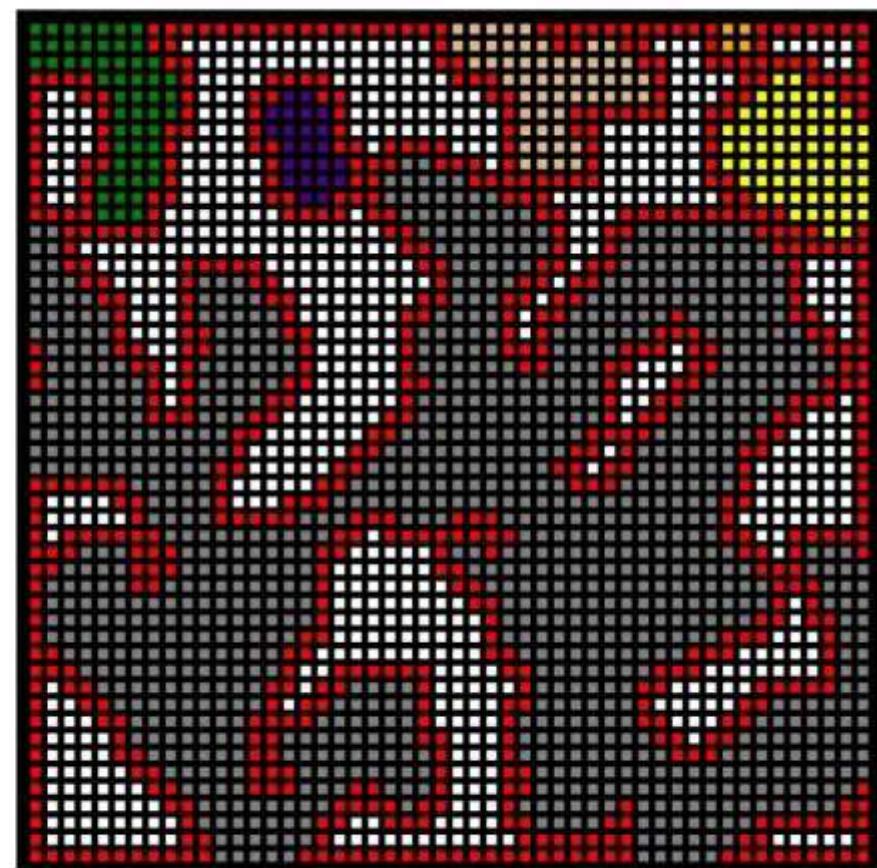
CA Parameters



- r : initial proportion of rock cells (0.5)
- n : CA iterations (4)
- T : neighborhood value threshold that defines a rock (5)
- M : Moore neighborhood size (1)



(a) Random map



(b) CA map

Adjacent Rooms



- The infinite cave needs to be contiguous - and you need to be able to turn back!
(Visited rooms stored as random seeds)
- Generate all four neighbors of a new room
- Dig tunnels from the central room to the new rooms at the shortest points
- Run the CA m times (2) on *all five rooms together* to smooth out edges

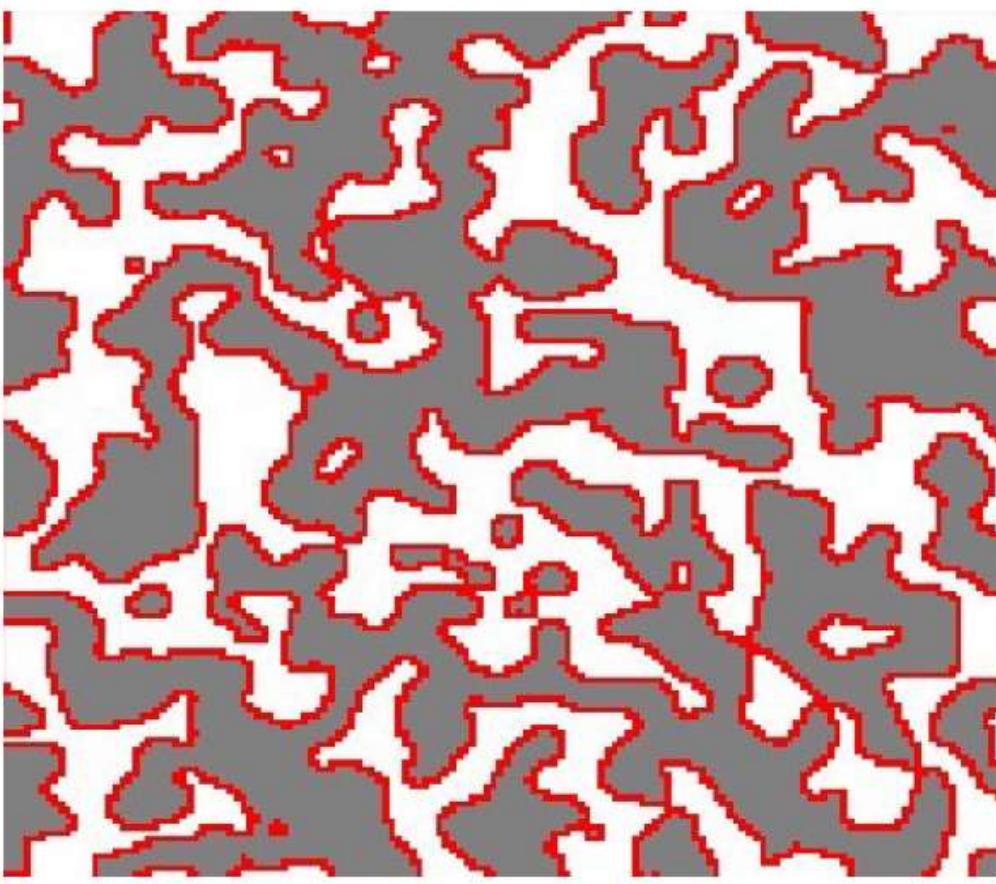
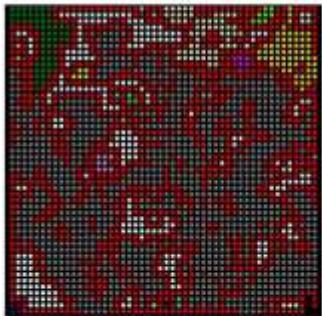


Figure 3: A 3×3 base grid map generated with CA. Rock and wall cells are represented by red and white color respectively. Grey areas represent floor. ($M = 2; T = 13; n = 4; r = 50\%$)

Controllable?



Parameters can be varied...
...but what do they mean?



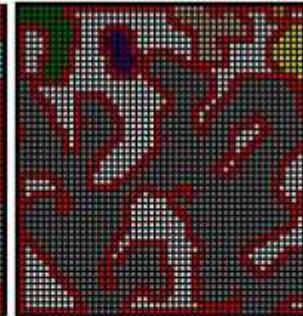
(a) $n = 1, M = 1, T = 5$



(b) $n = 2, M = 1, T = 5$



(c) $n = 3, M = 1, T = 5$



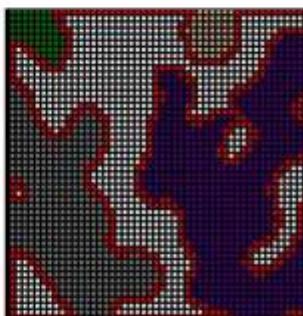
(d) $n = 4, M = 1, T = 5$



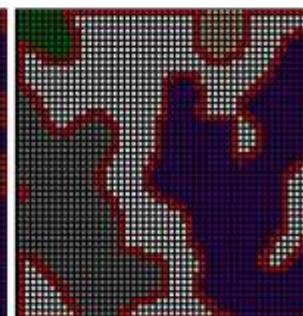
(e) $n = 1, M = 2, T = 13$



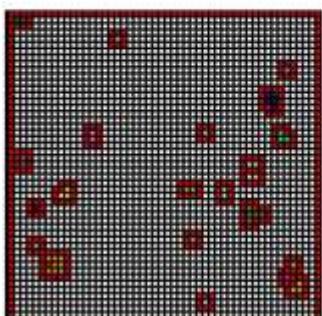
(f) $n = 2, M = 2, T = 13$



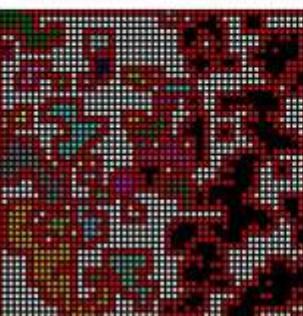
(g) $n = 3, M = 2, T = 13$



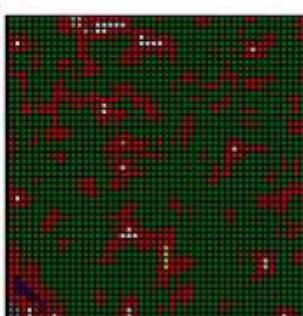
(h) $n = 4, M = 2, T = 13$



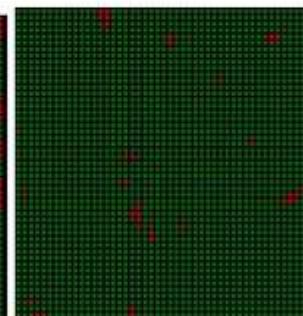
(i) $n = 1, M = 1, T = 2$



(j) $n = 1, M = 1, T = 4$

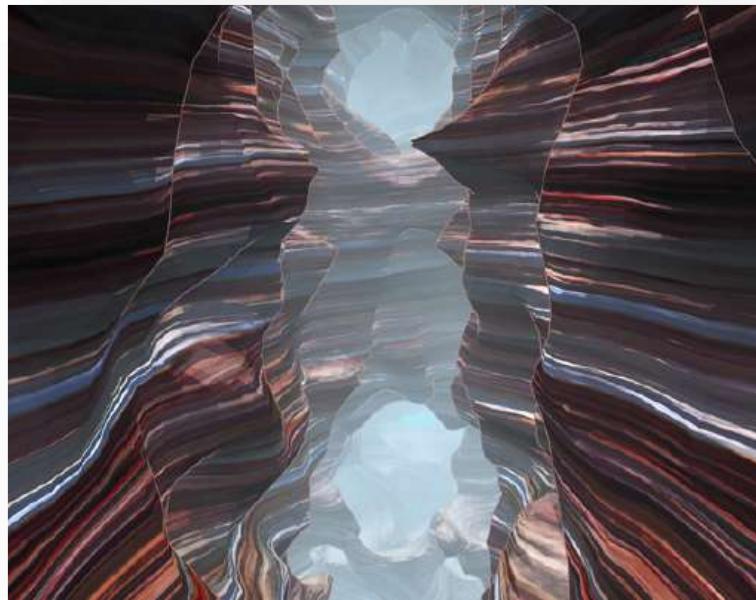
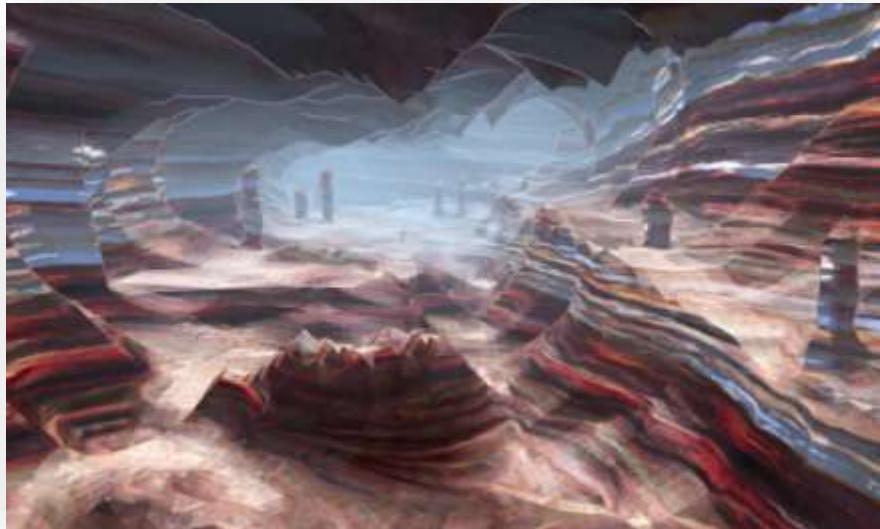


(k) $n = 1, M = 1, T = 6$



(l) $n = 1, M = 1, T = 8$

3D L-systems + CA



Solver-based Methods



Solver-based Methods



- Instead of using evolutionary algorithms, use constraint solvers and specify constraints
- Example: Answer Set Programming (ASP)
 - Declarative programming for search problems
 - Based on logic programming (syntax very similar to Prolog)
 - Finding an answer set is equivalent to solving a satisfiability problem



Grammars



Grammars



- Basic computer science concept, used in theory of computation
- Define a grammar and an alphabet and then watch an axiom unfold into ever more complex strings
- Commonly used for generating e.g. plants



Self-Similarity





Self-Similarity



- Nature has obviously thought out some clever way of representing complex organisms using a compact description...
- ...permitting individual variation...
- ...why is this relevant for us?

L-systems



- Introduced by Aristid Lindenmeyer 1968, to model plant development
- Creates strings (text) from an *alphabet* based on a *grammar* and an *axiom*
- Closely related to Chomsky grammars (but productions carried out in parallel, not sequentially)

An Example L-system



- Alphabet: {a, b}
- Production rules (grammar):
 $a \rightarrow ab$
 $b \rightarrow a$
- Axiom: b

b
|
a
|
L
|
a b
|
|
ab a
|
|
L
|
abaab
| / | L \
abaababa

Example of a derivation in a
DOL-System

Types of L-systems



Context

- **Context-free:** production rules refer only to an individual symbol
- **Context-sensitive:** productions can depend on the symbol's neighbors

Determinism

- **Deterministic:** there is exactly one production for each symbol
- **Non-deterministic:** several productions for a symbol

A Graphical Interpretation of L-systems



- Invented/popularized by Prusinkiewicz in 1986
- Core idea: interpret generated strings as instructions for a turtle in **turtle graphics**
- Read the string from left to right, changing the state of the turtle (x, y, heading)

Example Graphical Systems

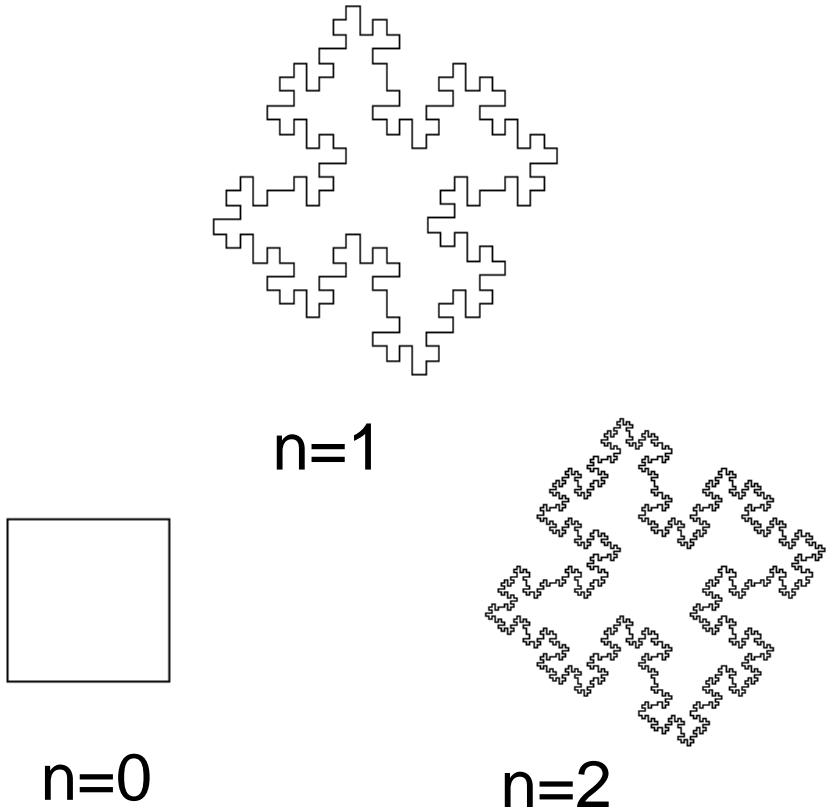


- Alphabet: {F, f, +, -}
- F: move the turtle forward (drawing a line)
- f: move the turtle forward (don't draw)
- +/-: turn right/left (by some angle)

Graphical L-system



- Axiom: $F+F+F+F$
- Grammar:
- $F>F+F-F-FF+F+F-F$
- Turning angle: 90°



Graphical L-systems



What's the limit of these systems?

Bracketed L-systems



- Alphabet: {F, f, +, -, [,]}
- [: push the current state (x, y, heading of the turtle) onto a pushdown stack
-]: pop the current state of the turtle and *move* the turtle there *without drawing*
- Enables branching structures!

Bracketed L-systems

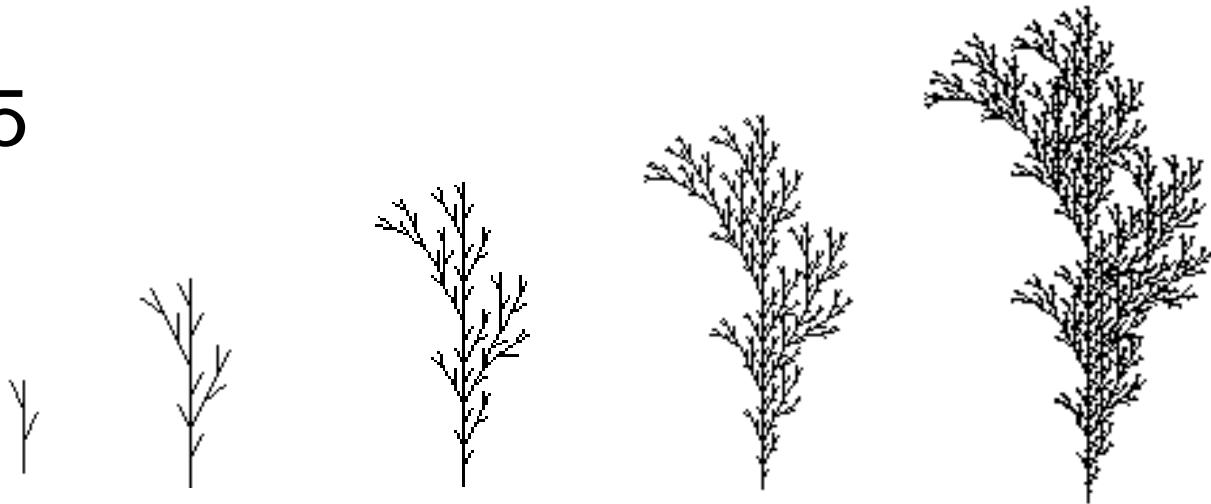


Axiom: F

Grammar: F>F[-F]F[+F][F]

Turning angle: 30°

n=1,...,5

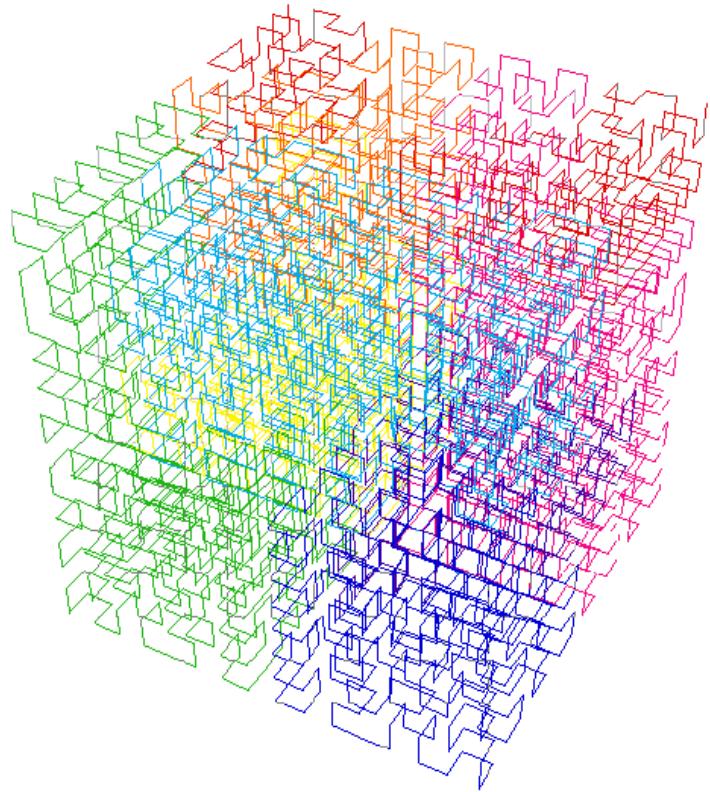
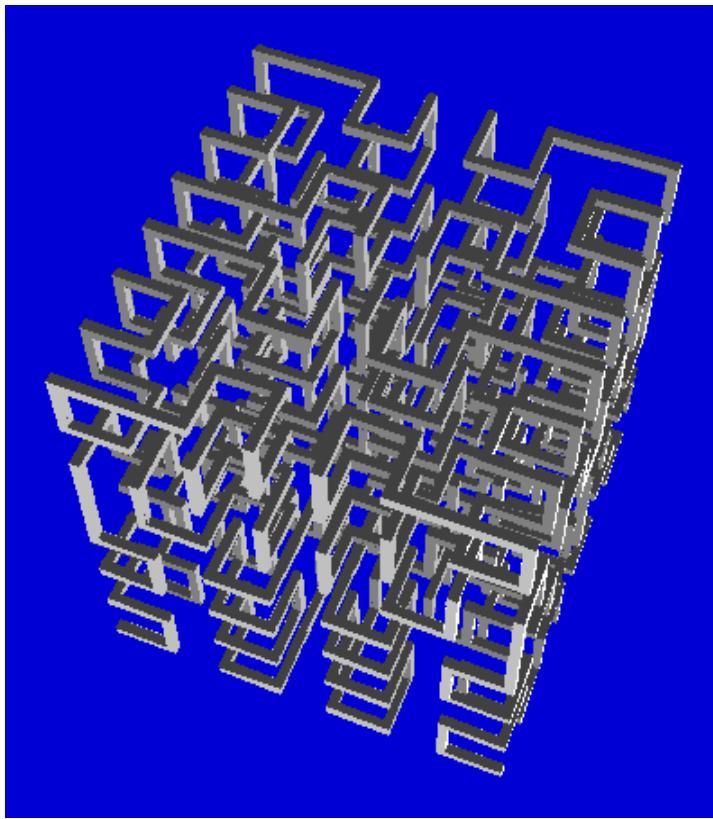


3D Graphics

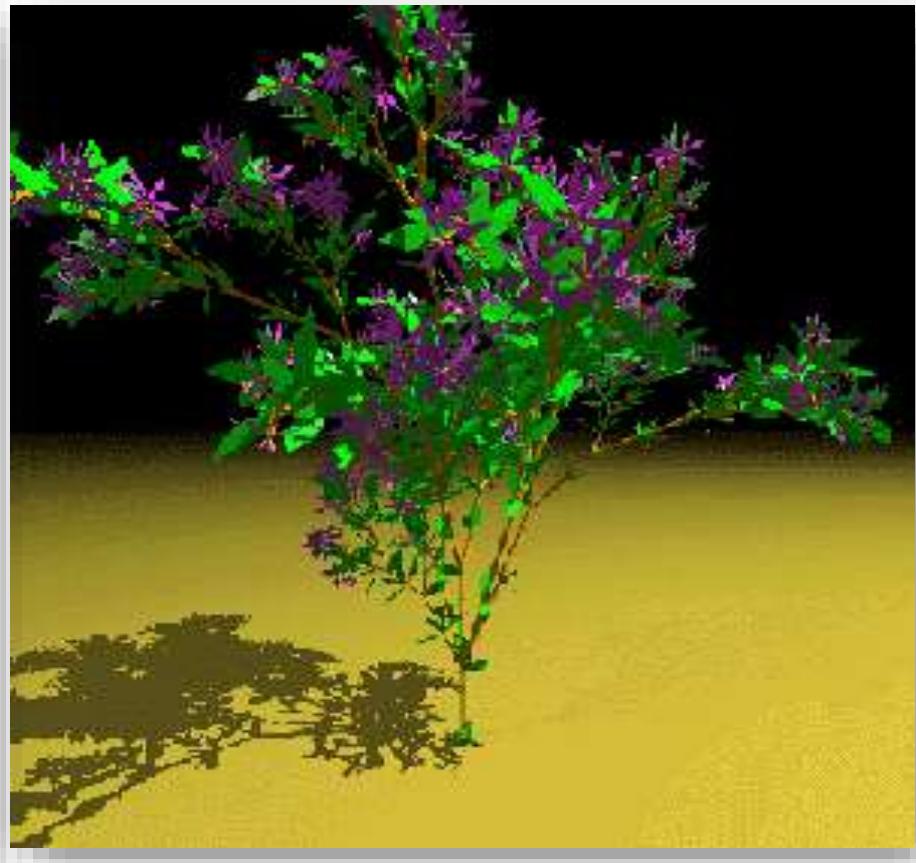
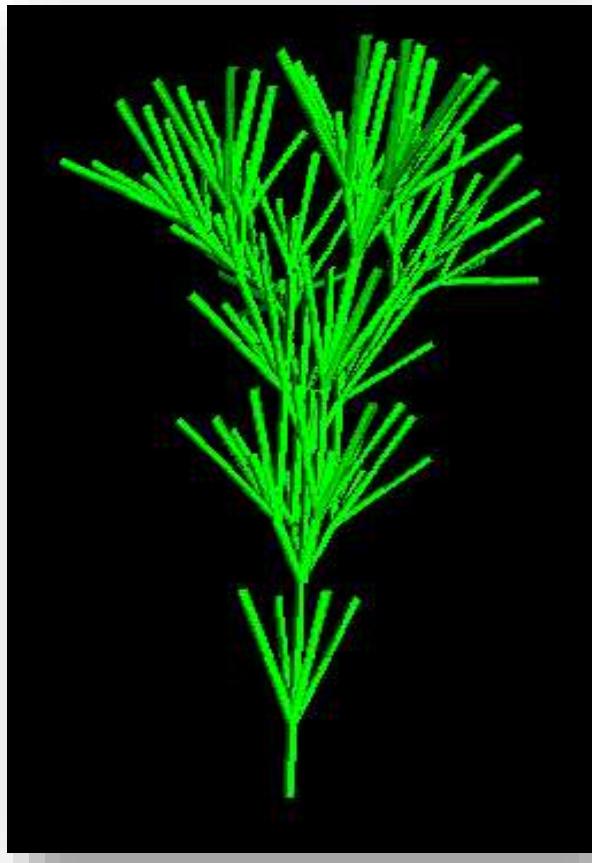


- Turtle graphics L-system interpretation can be extended to 3D space:
- Represent state as x, y, z and pitch, roll, yaw
- +, -: turn (yaw) left/right
- &, ^: pitch down/up
- \, /: roll left/right (counterclockwise/clockwise)

3D Interpretation of L-systems

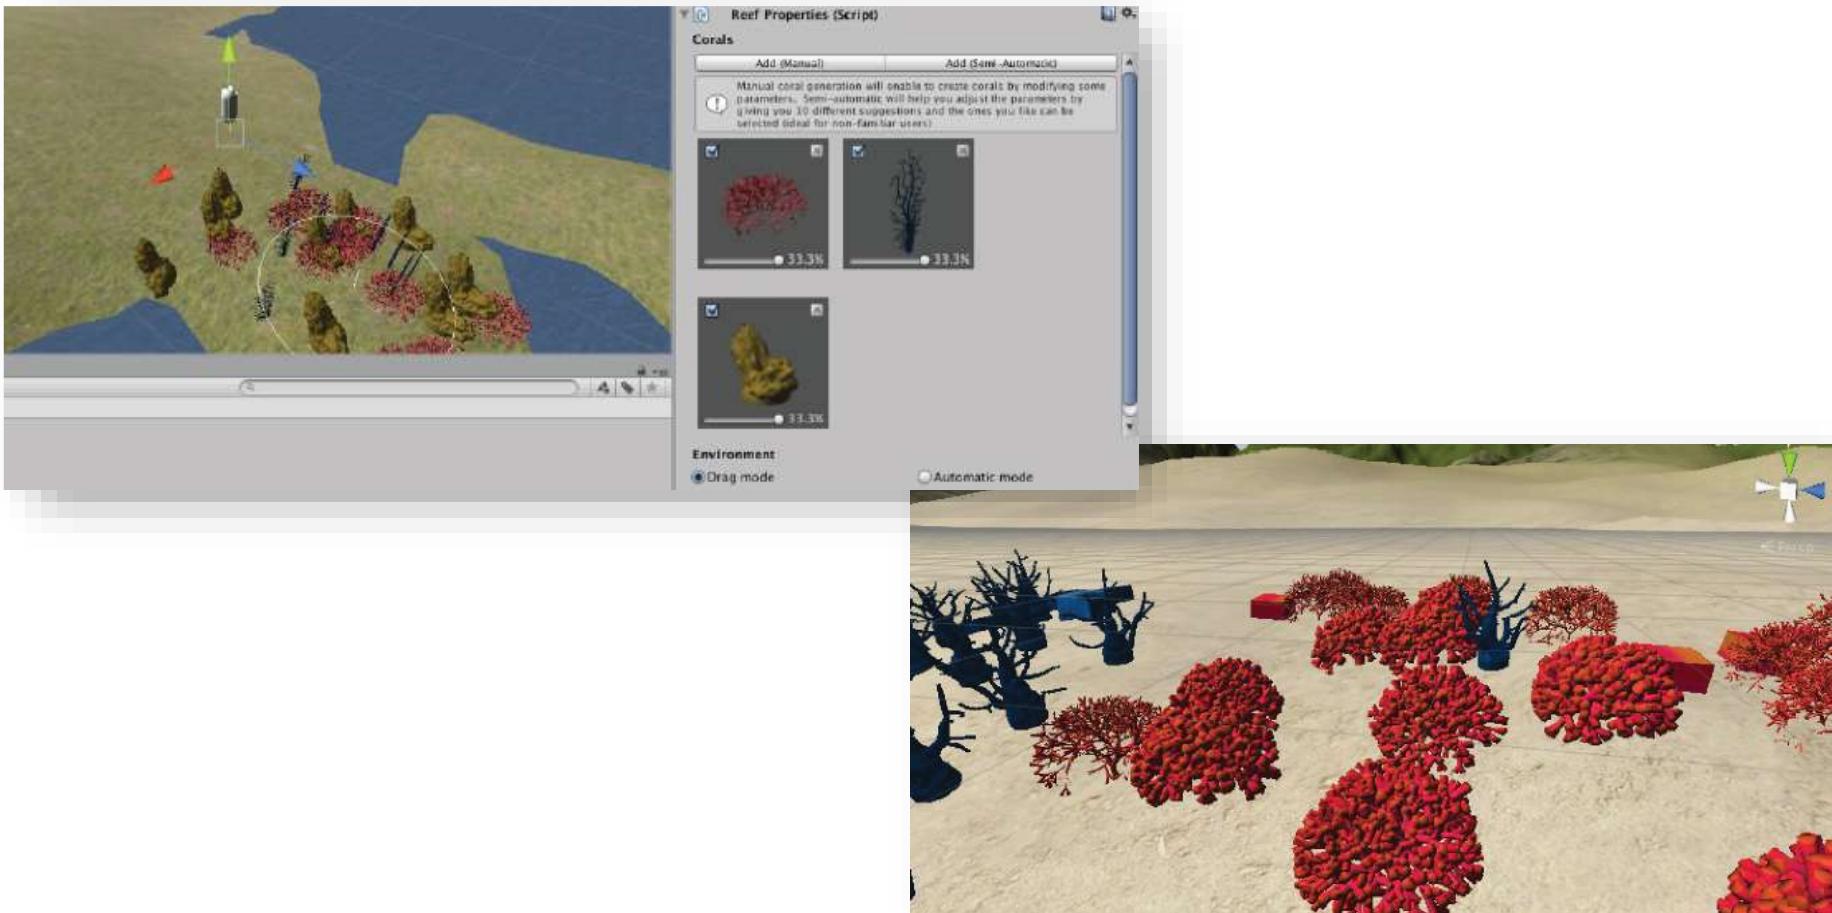


3D Interpretation of Bracketed L-systems



Coralize: 3D Corals in Unity

Abela, R., Liapis, A. and Yannakakis, G.N., 2015. **A constructive approach for the generation of underwater environments.** In Proceedings of the FDG workshop on Procedural Content Generation in Games.



2D L-systems



Axiom: \boxed{A}

Rules:

$\boxed{A} \rightarrow \begin{array}{|c|c|} \hline A & B \\ \hline B & A \\ \hline \end{array}$

$\boxed{B} \rightarrow \begin{array}{|c|c|} \hline A & A \\ \hline B & B \\ \hline \end{array}$

Two Expansions:

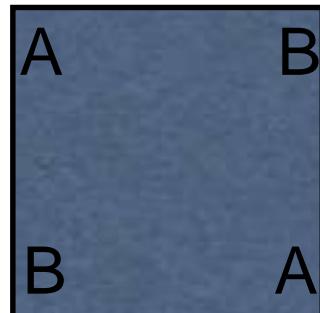
$\boxed{A} \rightarrow \begin{array}{|c|c|} \hline A & B \\ \hline B & A \\ \hline \end{array}$

$\begin{array}{|c|c|c|c|} \hline A & B & A & A \\ \hline B & A & B & B \\ \hline A & A & A & B \\ \hline B & B & B & A \\ \hline \end{array}$

Terrain Interpretation of 2D L-systems



- Each group of four letters is interpreted as instructions for lowering or raising the corners of a square
- E.g. A=+0.5, B=-0.5



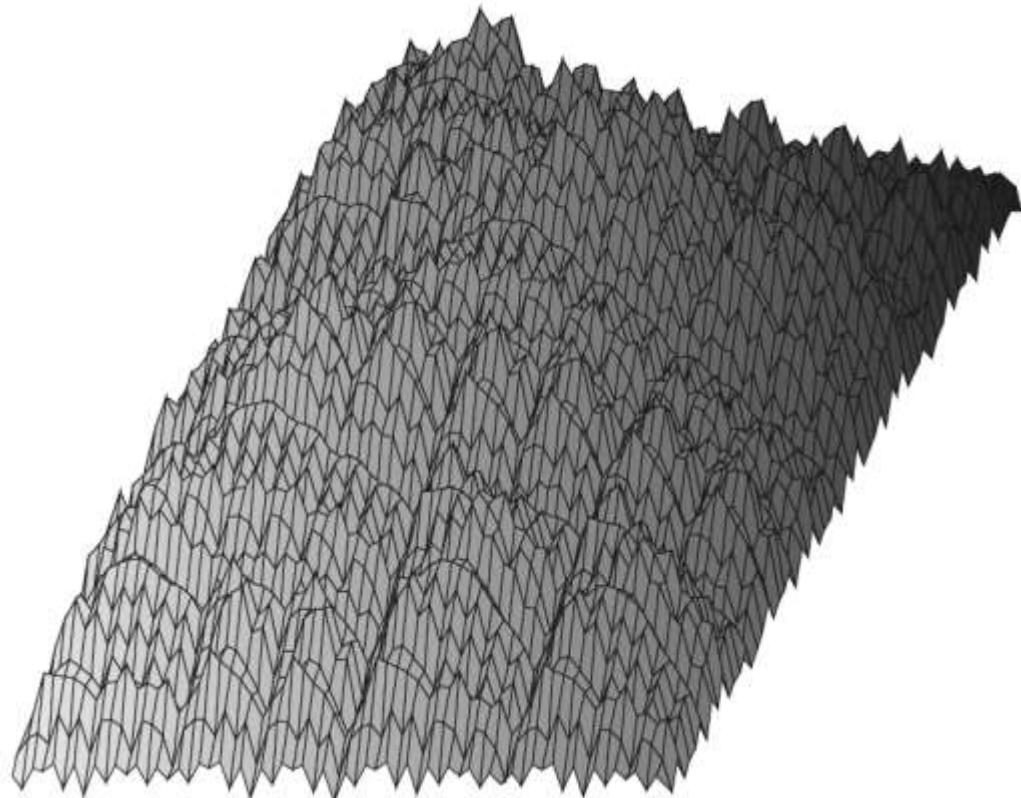
Terrain Interpretation of 2D L-systems



- In next iteration, the 2D L-system is rewritten once, and each square is divided into two
- “Doubling the resolution”

A	BA	B
B	AB	A
A	BA	B
B	AB	A

Terrain Interpretation of 2D L-systems



Six rewritings of $A \rightarrow ABBA$, $B \rightarrow AABB$

Grammars for Adventure Level Design



1. Dungeon → Obstacle + treasure
2. Obstacle → key + Obstacle + lock + Obstacle
3. Obstacle → monster + Obstacle
4. Obstacle → room

could give...

key + monster + room + lock + monster + room + treasure
key + monster + key + room + lock + monster + room + lock +
room + treasure
room + treasure
monster + monster + monster + monster + room + treasure

Search-Based PCG

Togelius, J., Yannakakis, G.N., Stanley, K.O. and Browne, C., 2011. Search-based procedural content generation: A taxonomy and survey. *IEEE Transactions on Computational Intelligence and AI in Games*, 3(3), pp.172-186.



Search-Based PCG



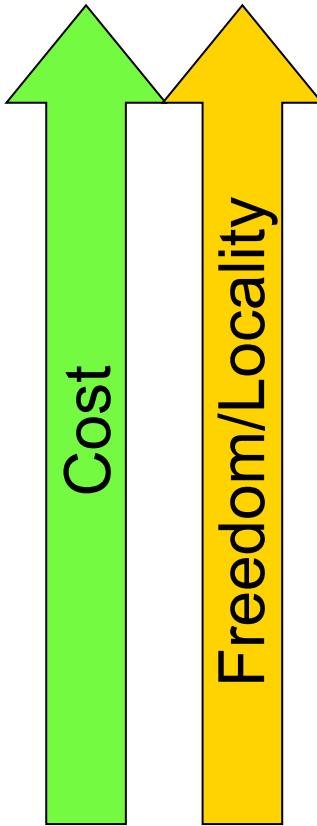
- Use evolutionary computation to search the design space for good artifacts (e.g. levels)
 - Technically, we could use other stochastic search / optimization algorithms
- Major issues:
 - Representing the content
 - Devising a good evaluation / fitness function

The Algorithm



- Lots of different types of evolutionary algorithms: Genetic Algorithms, Evolution Strategies, Evolutionary Programming
- And evolution-like algorithms: Particle Swarm Optimization, Differential Evolution
- Keep It Simple, Stupid!
 - Often, simple $\mu+\lambda$ ES with no crossover and no self-adaptation works well enough

Representing Content (e.g. a dungeon)



- *Directly*: grid
- *More indirectly*: position and orientation of walls
- *Even more indirectly*: patterns of walls and floor
- *Very Indirectly*: number of rooms and doors
- *Indirectly*: random seed

How to Evaluate Content Quality



- **Directly**
 - A direct mapping between content and quality; e.g. number of jumps in a platform game
- **Simulation-based**
 - An AI (maybe a human imitator) plays the game for a while and content is evaluated
- **Interactively**
 - Real-time evaluation via a player (or players)

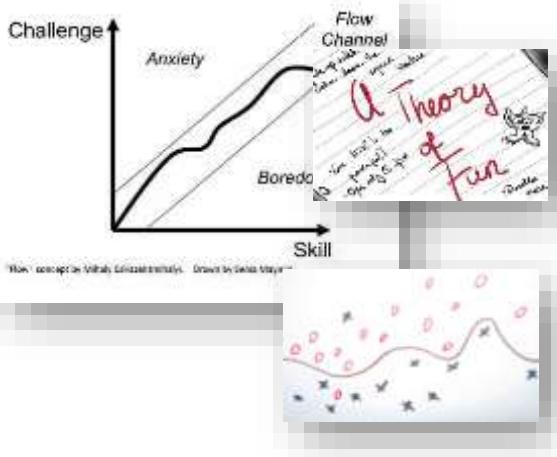
How to Evaluate Content Quality



- **Directly**
 - **Theory-driven**: evaluation function is based on a theoretical model – e.g. Koster's theory of fun
 - **Data-driven**: evaluation function is derived via gameplay (or other modalities of) data
- **Simulation-based**
 - **Static**: evaluation function does not change over time
 - **Dynamic**: evaluation function is affected as time goes by
- **Interactively**
 - **Implicit**: game behavior gives value to content (e.g. preference over a weapon)
 - **Explicit**: ask players to score content

Content Quality

Direct



Simulation-based



Interactive

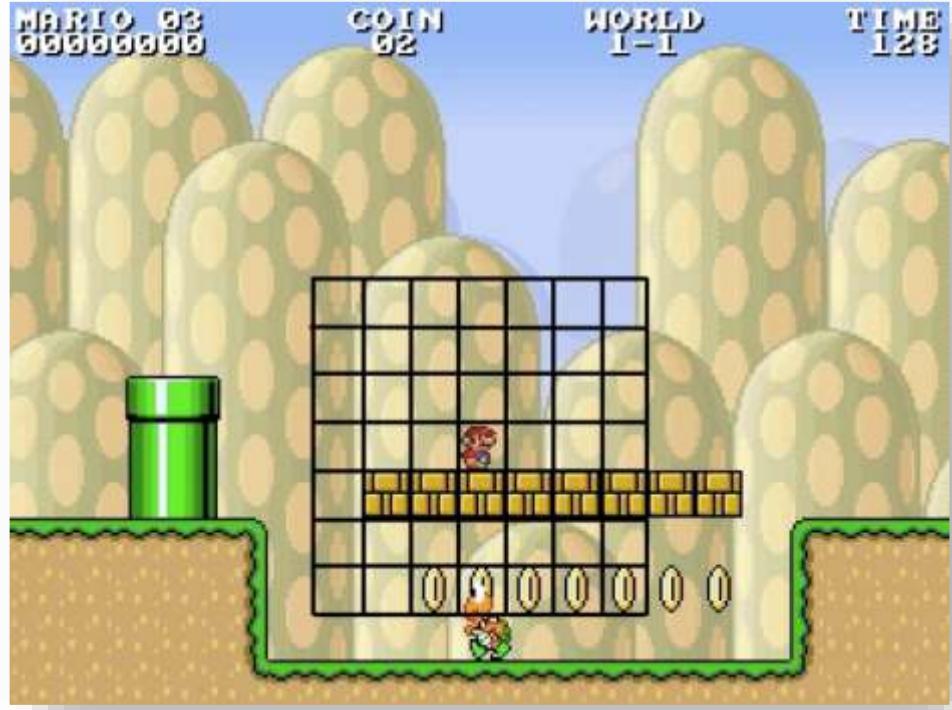


Search-Based PCG Example #1

How would we generate levels for Super Mario Bros?



The Mario AI Benchmark

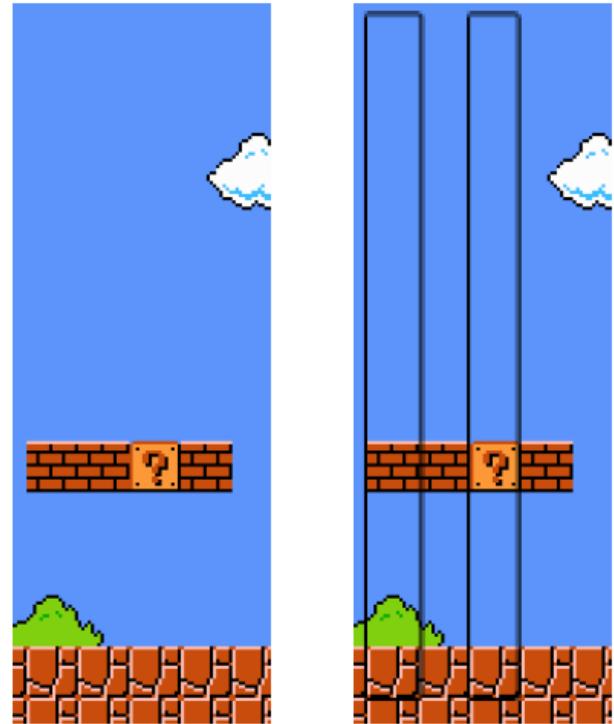


- Reasonably faithful clone of SMB 1/3
- APIs for level generators and AI controllers

Representation



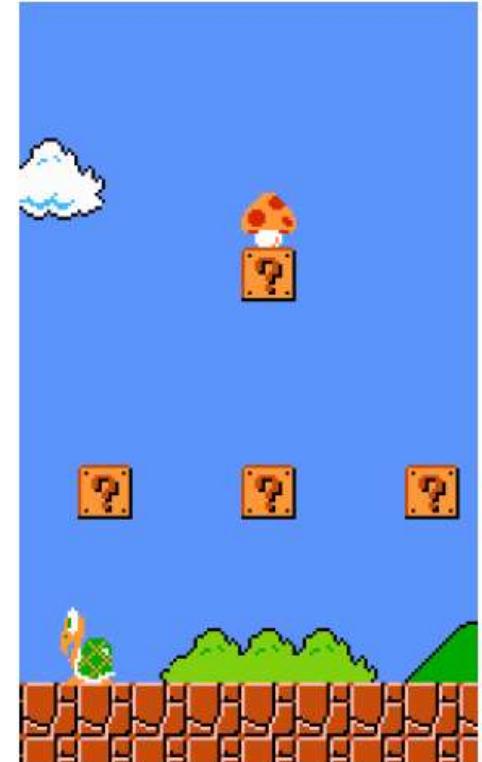
- A number of “vertical slices” are identified from the original SMB levels
- Levels are represented as strings, where each character correspond to a pattern



Evaluation



- 25 patterns are identified in the original SMB levels
- e.g. enemy hordes, pipe valleys, 3-paths...
- The fitness function counts the number of patterns found in the level



MARIO 03

COIN
15

WORLD
none

TIME
029

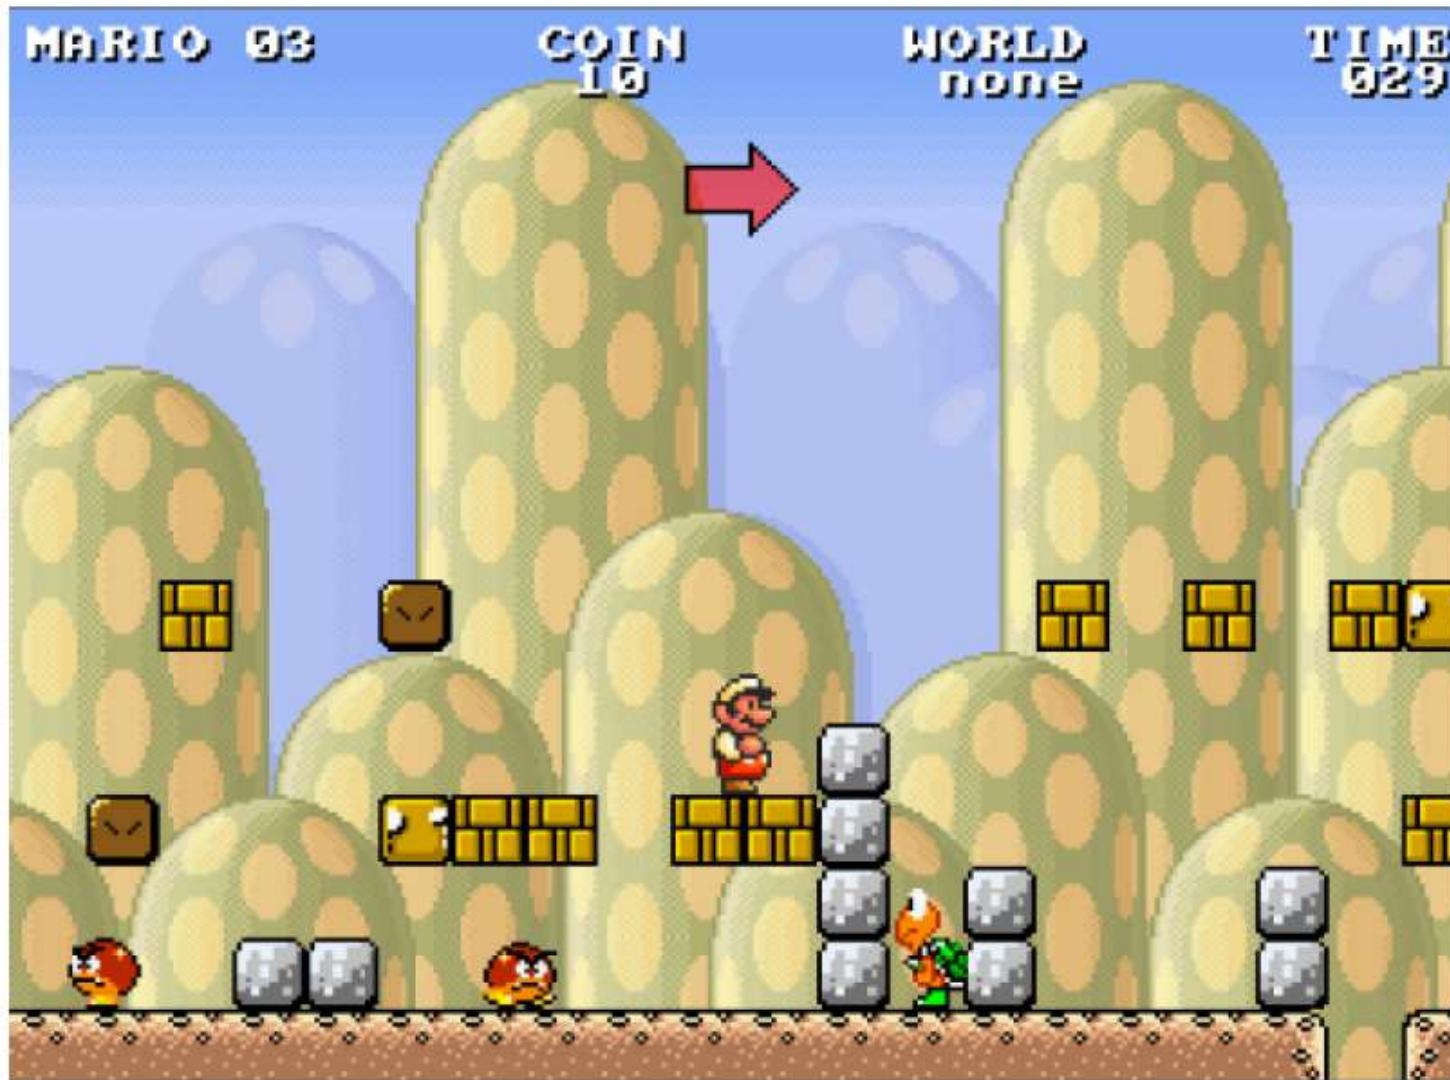




Fig. 15. FFMacro #98 MC: 6, fitness value: 2332 (highest).



Fig. 16. FFMesoB #6 MC: 0, fitness value: 545.



Fig. 17. FFMesoB #30 MC: 3, fitness value: 409 (lowest).



Fig. 18. FFMesoB #64 MC: 6, fitness value: 2065 (highest).

Search-Based PCG Example #2

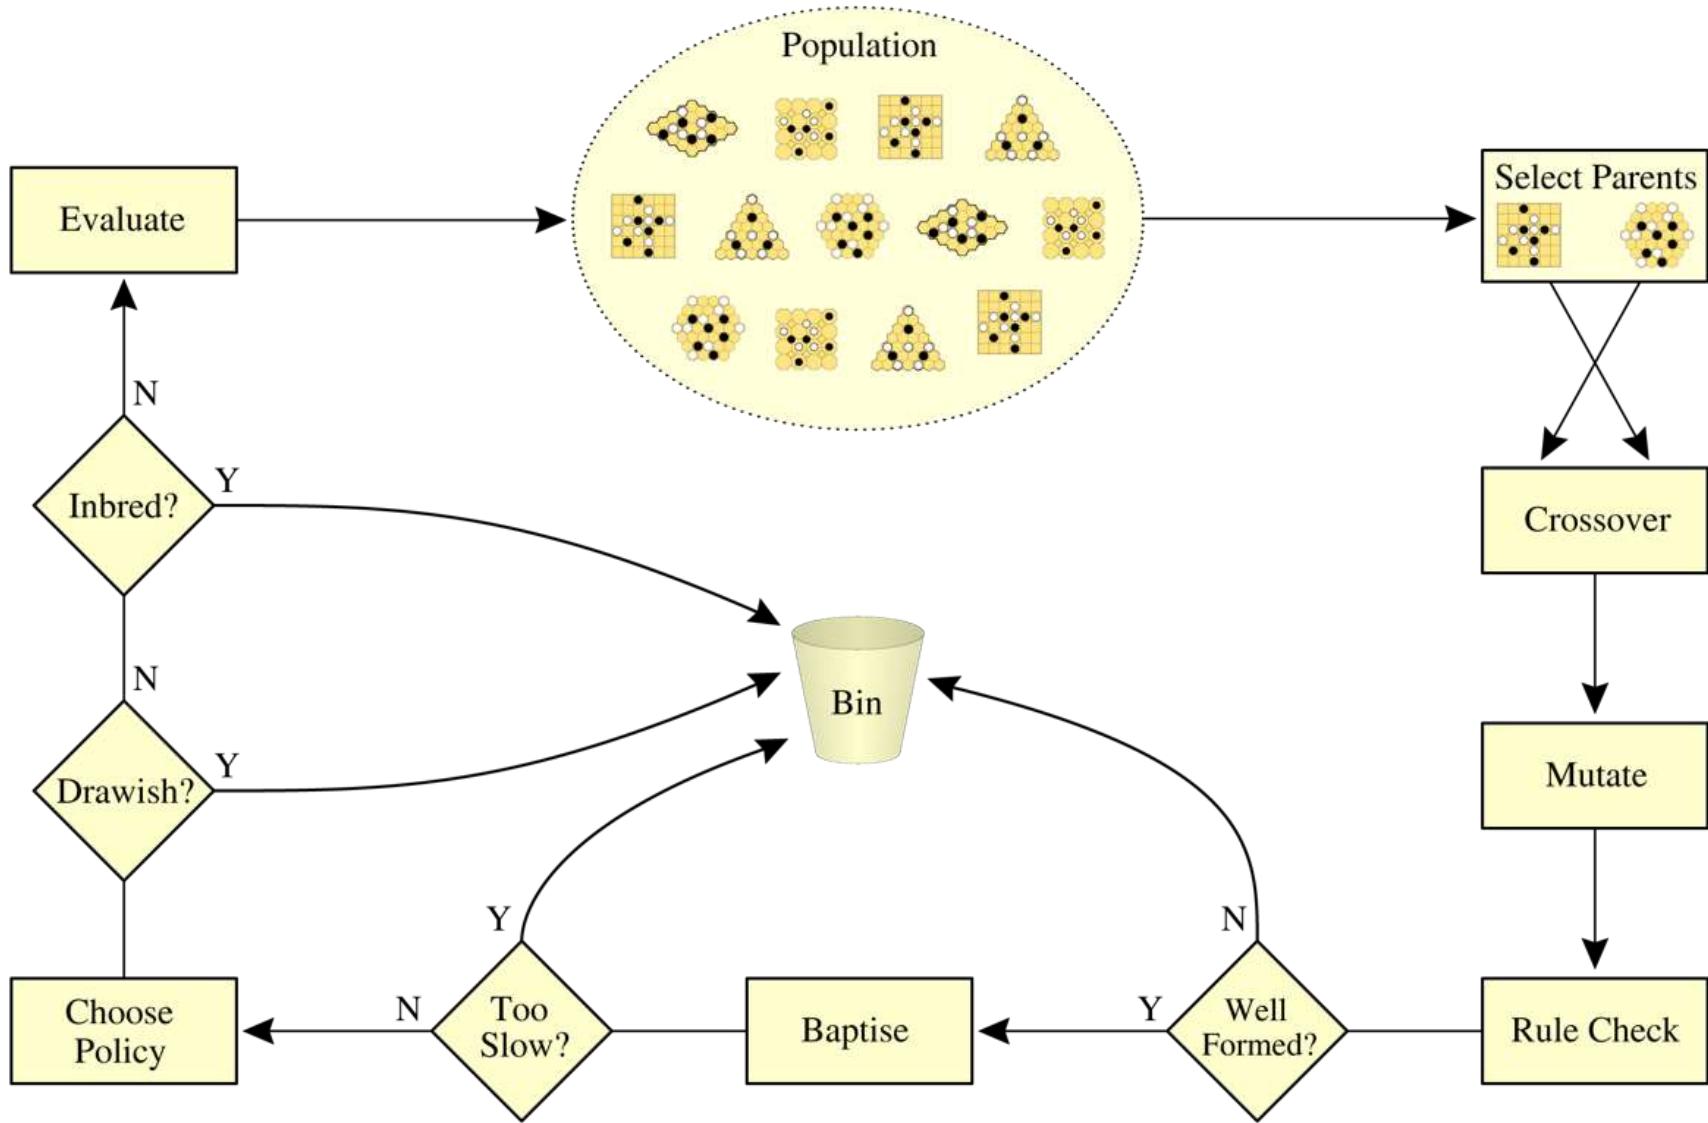
How would we create new game rules?



Creating Game Rules



- Rules are also content...
- Will need simulation-based evaluation - you can only judge game rules by playing the game
- Has been attempted for simple Pac-Man-like games (Togelius 2008), GVGAI games (Nielsen et al 2015)
- Perhaps most convincingly for board games (Browne 2008)



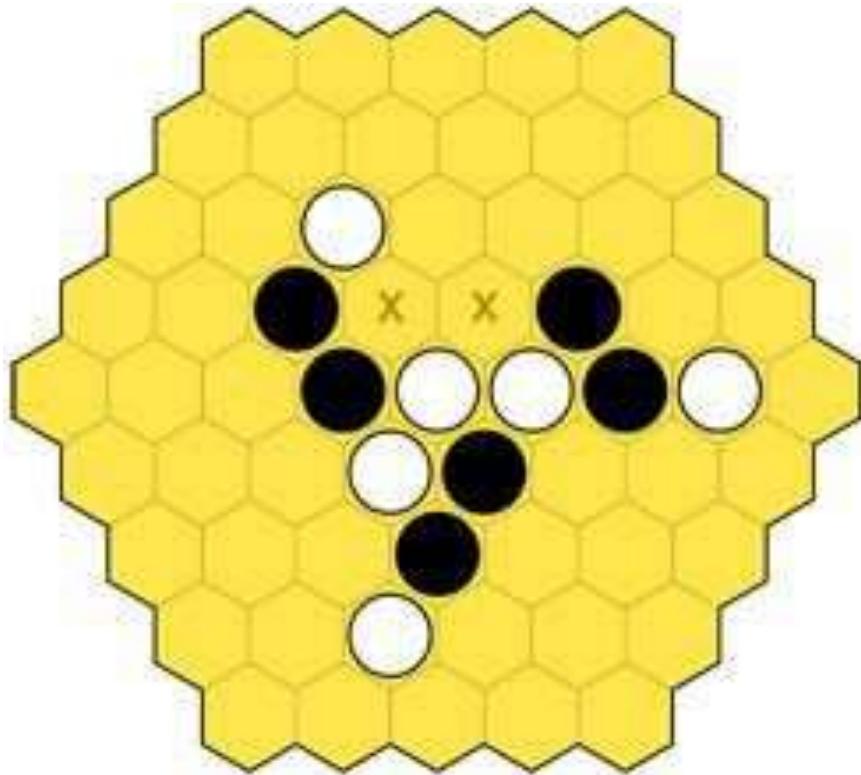
Yavalath

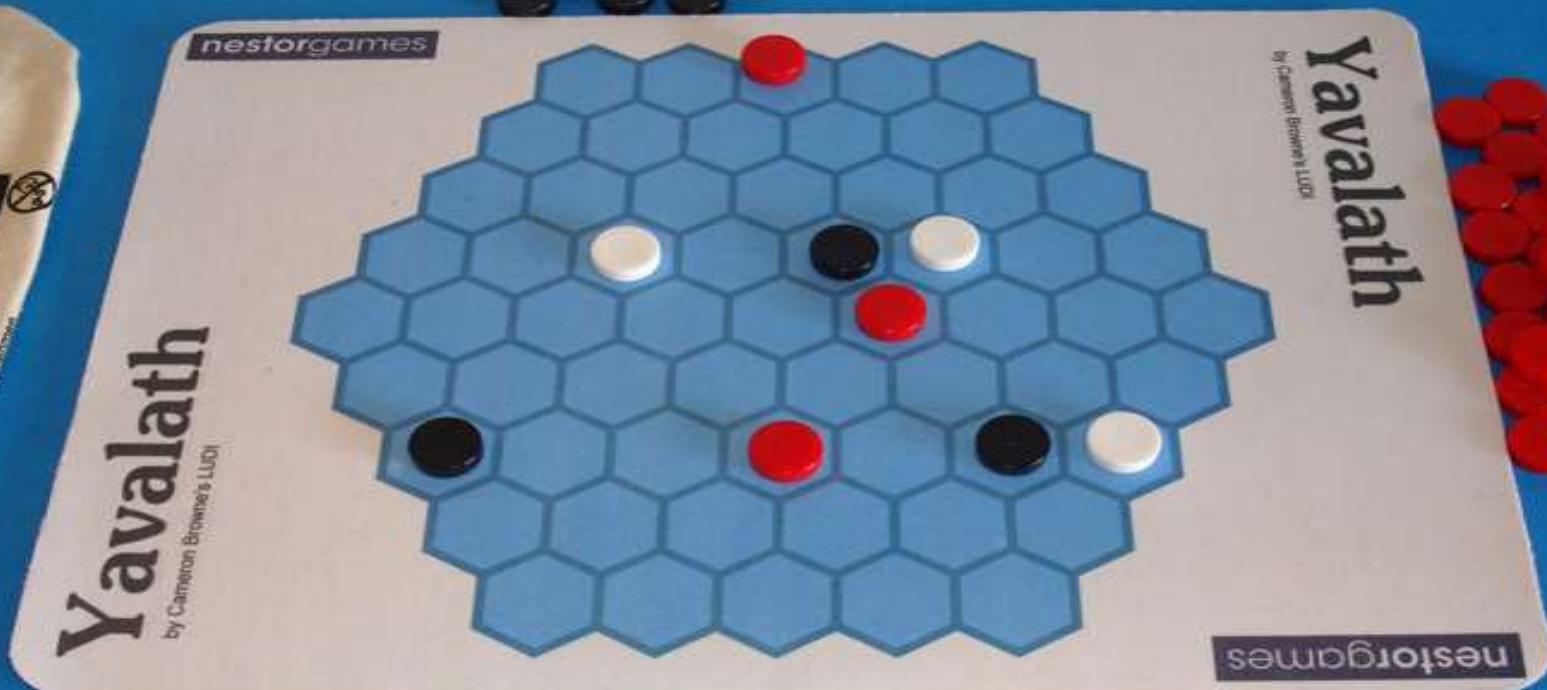
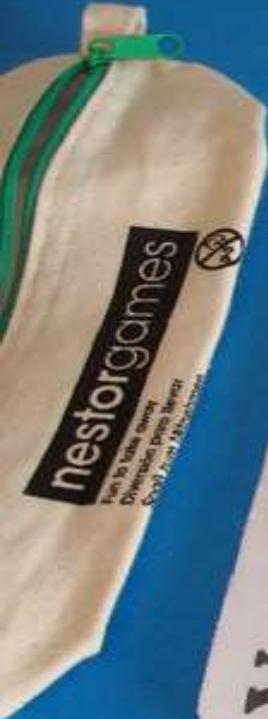


YAVALATH (#2)

```
(game Yavalath
  (players White Black)
  (board (tiling hex) (shape hex) (size 5))
  (end
    (All win (in-a-row 4))
    (All lose (and (in-a-row 3) (not (in-a-row 4))))))
  )
```

Yavalath





Search-Based PCG Example #2

How would we design L-systems?



Evolving L-systems



How can we combine L-systems with evolutionary computation?

Evolving L-systems



- Evolving the axiom
- Evolving the grammar:
 - Change the shape of one or more production rules, or
 - Add/remove/replace productions
- Evolving the interpretation:
 - Evolve production probabilities
 - Evolve other aspects (e.g. turning angles)

Evolving L-systems



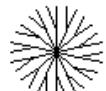
- One example: Ochoa evolved the consequent of a single production rule
 - starting from $F>F[-F]F[+F]F$
- Mutation: replace single symbols, or blocks of a few symbols
- Crossover: swap complete “sub-trees”
(like in genetic programming)

Fitness Functions



- Phototropism
- Bilateral symmetry
- Proportion of branching points

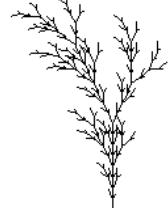
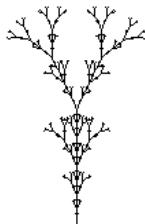
Evolved Systems



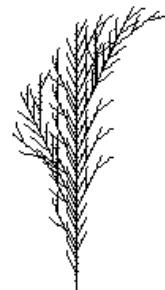
Symmetry



Branching
points

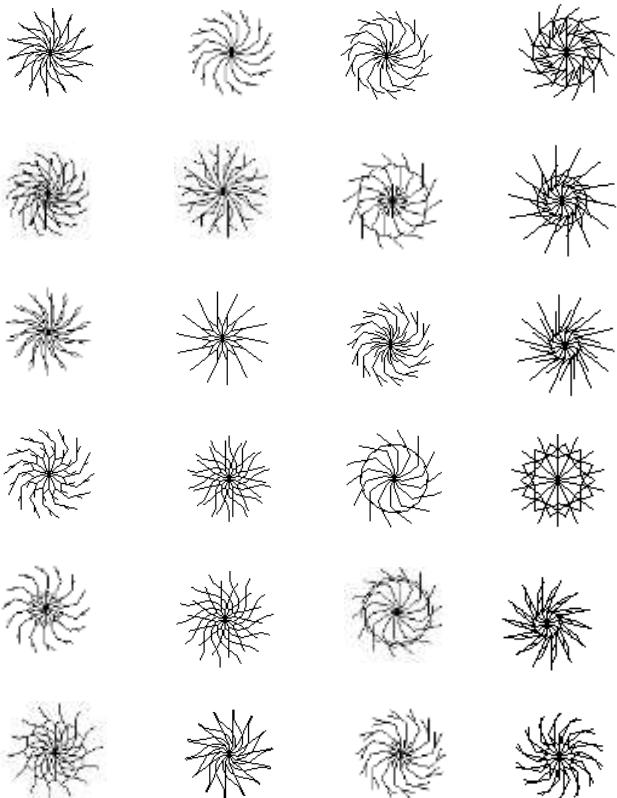


All 3

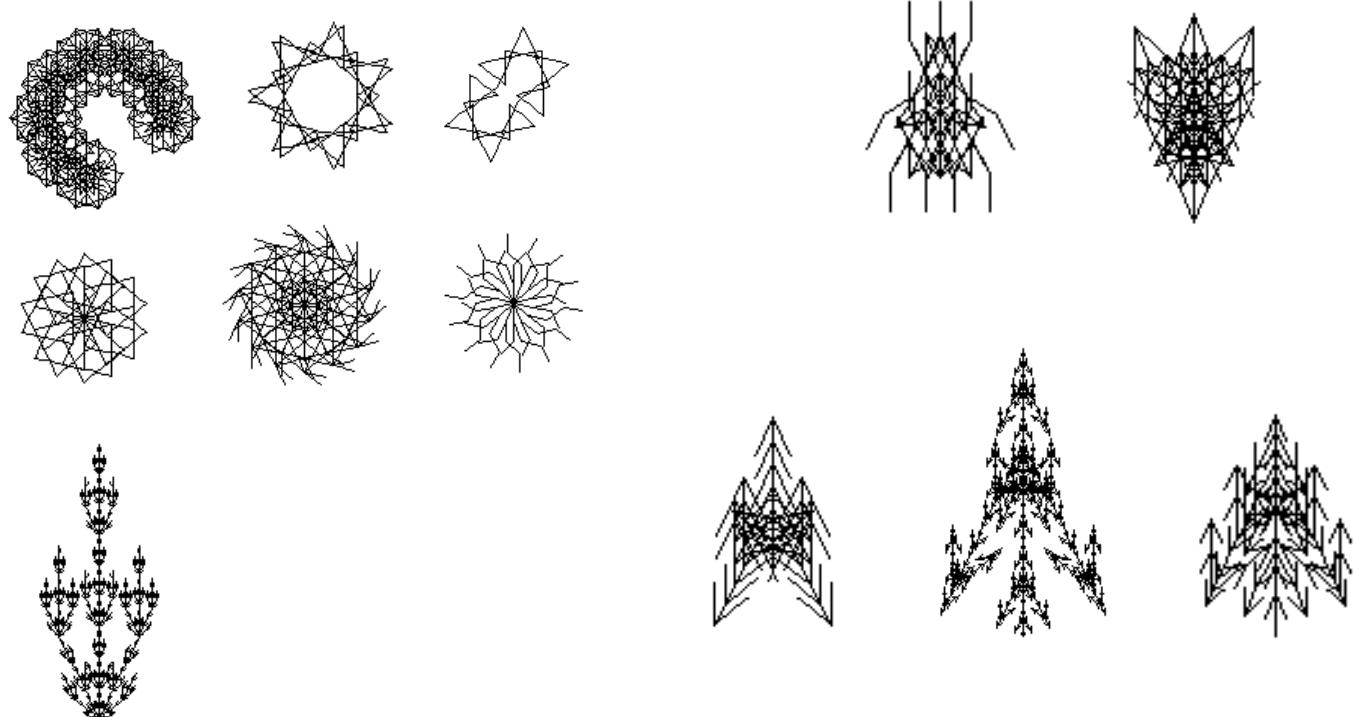


Phototropism +
Symmetry

Evolved Systems

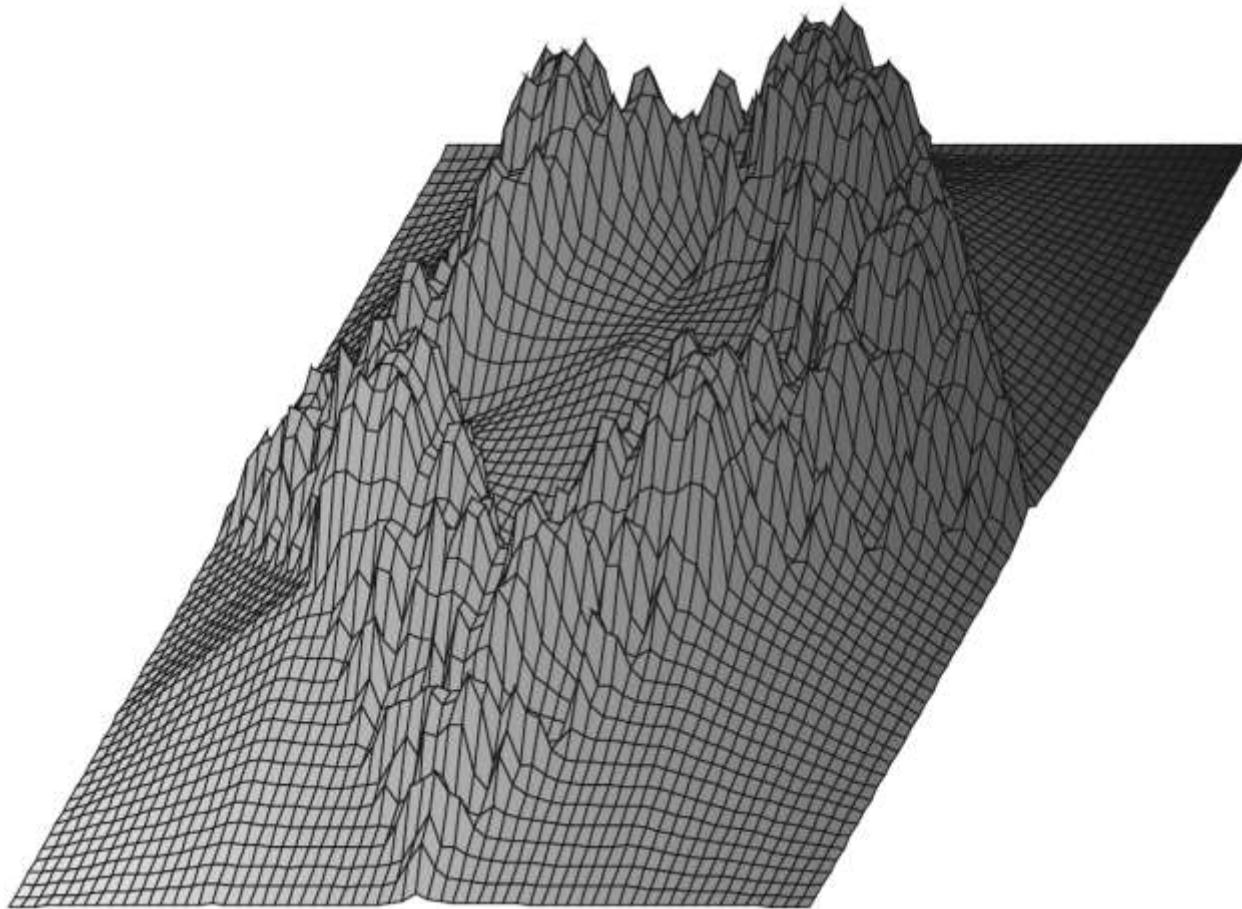


Evolved Systems

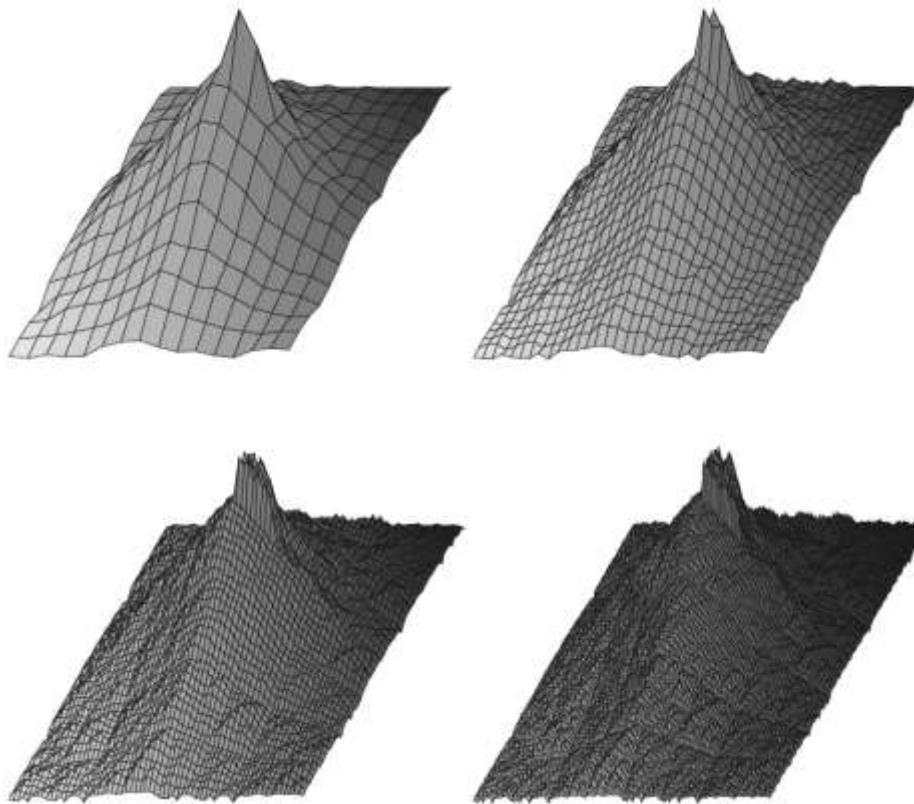


...and this was an extremely simple L-system!

Evolved 2D L-system Terrains



Evolved 2D L-system Terrains



Very short specification, yet infinite resolution!

PCG via Machine Learning



PCG via Machine Learning



- Basic idea of Procedural Content Generation via Machine Learning (PCGML): train machine learning models on corpuses of existing content, then generate new content
- Many methods useful, including n-grams, Markov chains, and... yes, even deep learning
- Reference: Summerville et al. (2018): Procedural Content Generation via Machine Learning (PCGML)

N-Grams



- Capture the statistic co-occurrence of sequential characters
- n: number of previous characters a new character depends on
- Commonly used for predictive text
- Can also be used for linear game content

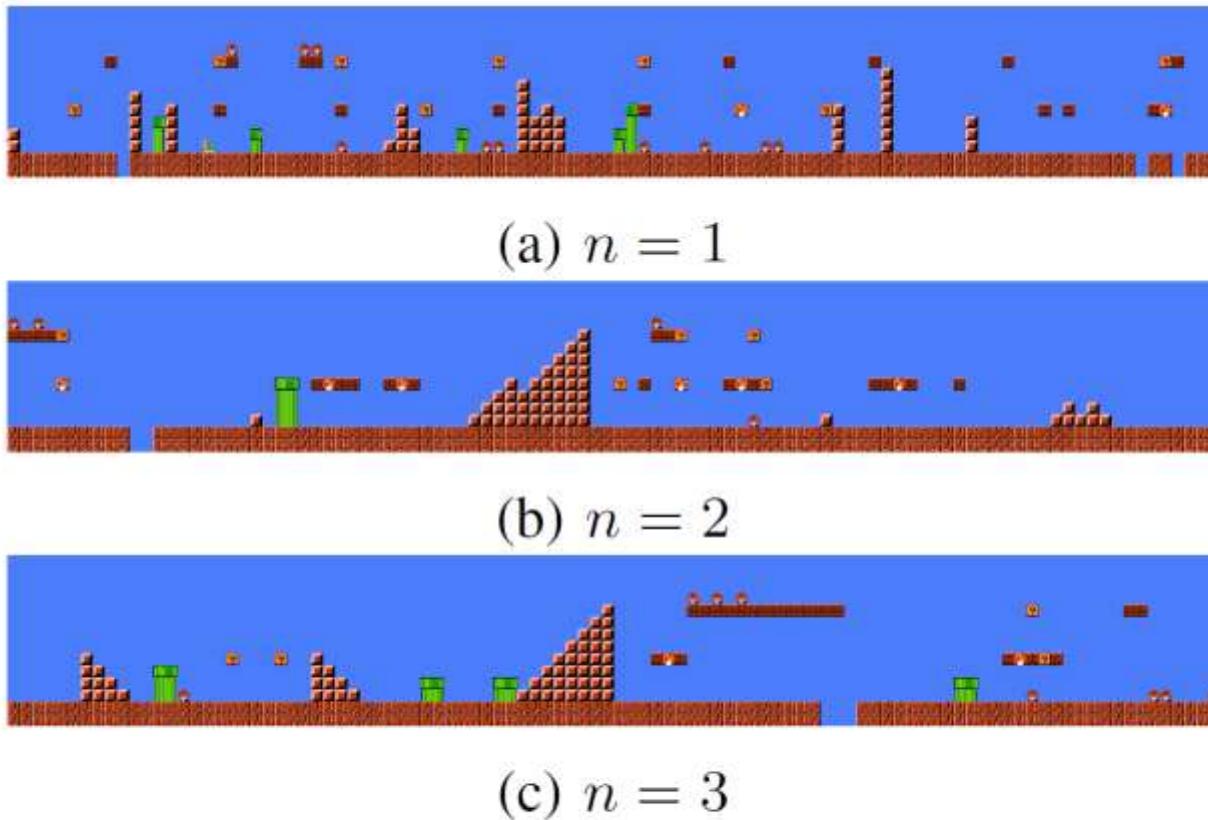


Figure 6: Mario levels reconstructed by n -grams with n set to 1, 2, and 3 respectively.

Neuroevolution



- Evolving the weights of neural networks, either interactively or according to some automatic fitness function
- Can be used for PCG: predicting some aspect of a level (e.g. enemies in Super Mario Bros) from other aspects (e.g. bricks and question mark blocks)
- Analogically: The different layers of a level are akin to different instruments in a song

Neuroevolution

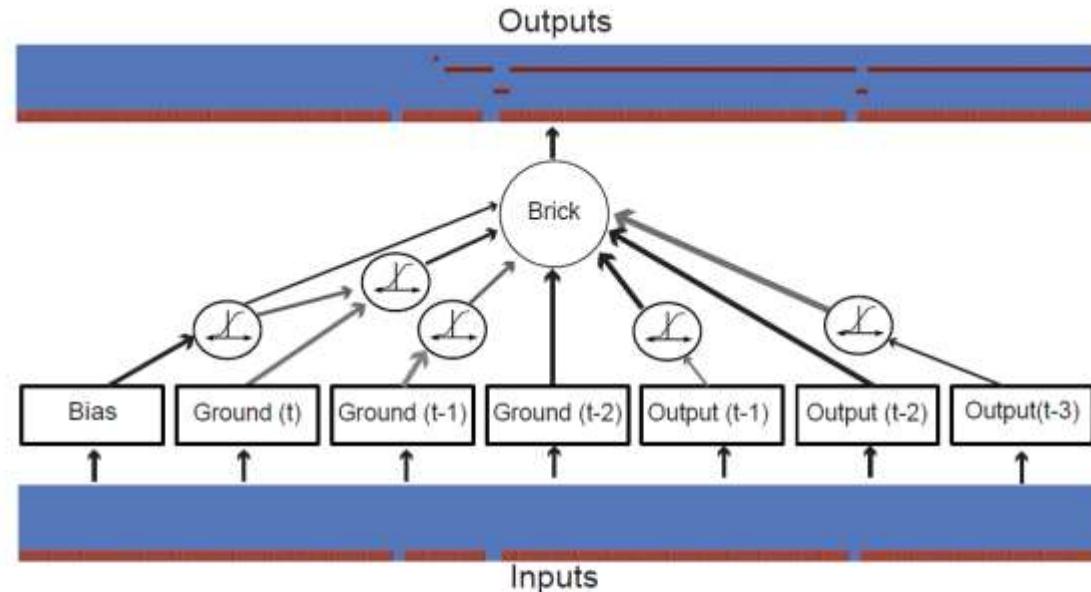


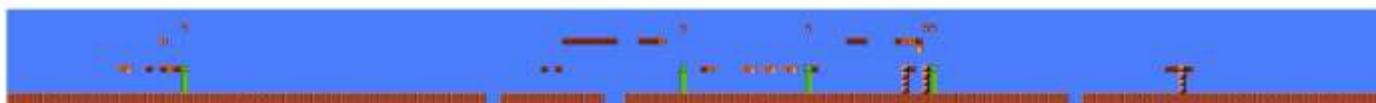
Figure 1: **The ANN Representation.** At each column in a given level, an ANN inputs visual assets (i.e., musical voices) from at least one tile type of *Super Mario Bros.* and outputs a different type of visual asset. This example figure shows ground tiles being input to the ANN while the brick tile placements are output from predictions made through NeuroEvolution of Augmenting Topologies [39]. To best capture the regularities in a level, each column is also provided information about the most recently encountered inputs (i.e., input values for the three previous columns). The inputs and outputs are then fed back into the ANN at each subsequent tick. Once trained, ANNs can potentially suggest reasonable placements for new human composed in-game tiles.



(a) Original World One Level One in Super Mario Bros.



(b) The Ground, Brick, and Question Mark Voices in World One Level One



(c) The remaining voices regenerated from several trained networks, using the three-voice level in (b) as input



(d) Author-created Level with Ground, Brick, and Question Mark Voices



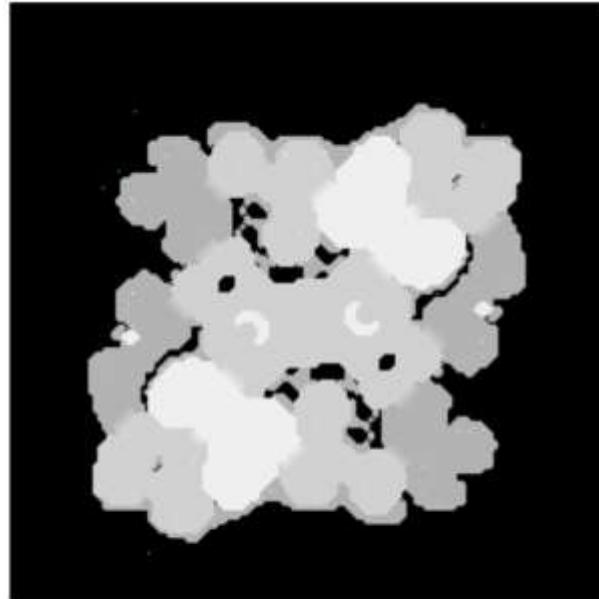
(e) Additional voices generated by trained networks for the author-created level.

Convolutional Neural Nets

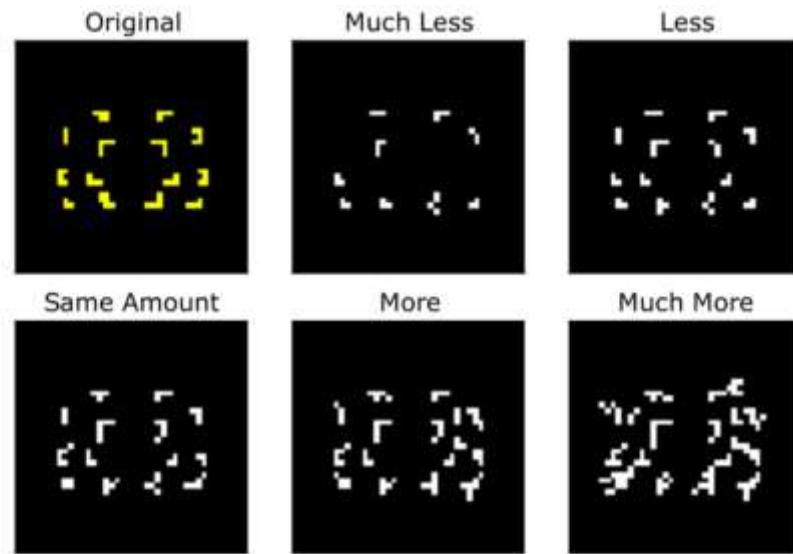


- Commonly used for classification and prediction from high dimensional matrices, such as images
- Makes use of the convolution operation to find recurring features in images while minimizing the number of parameters
- Can be used for PCG: predicting level features from other features (e.g. resources from base locations)

CNNs for *StarCraft* Resources



(a) A *StarCraft II* heatmap



(b) Varying resource amounts.

Figure 4: A heatmap and generated resource locations for *StarCraft II*

Long Short-Term Memory



- Commonly used for sequence recognition and prediction, for example for audio or text
- Often used for generative text, through predicting the next word or letter based on some training set
- Can be used for PCG: predicting the next element in a sequence, e.g. a level segment or a word in a description

LSTM for Mario



Figure 2: Example output of the LSTM approach, including generated exemplar player path.

LSTM for Magic Cards



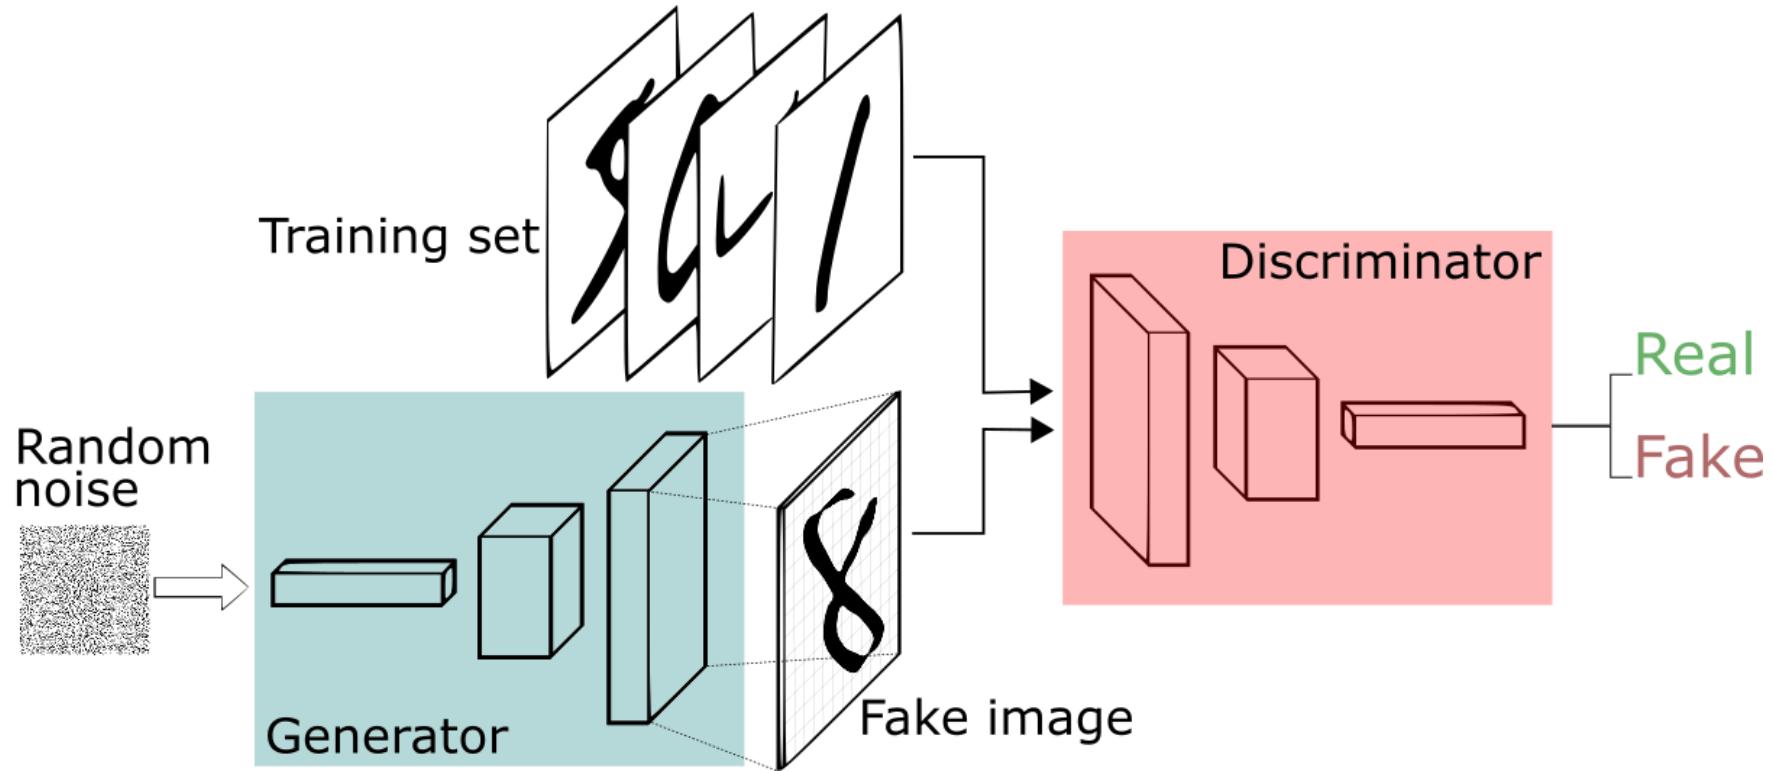
Figure 5: Example partial card specification and its resulting output.

Generative Adversarial Networks



- Train two networks intermittently: a generator and a discriminator
- The discriminator is trained to distinguish real artifacts from fake (generated) ones
- The generator is trained to generate artifacts that fool the discriminator
- Essentially the same idea as competitive coevolution

Generative Adversarial Networks



Evolving Mario Levels in the Latent Space of a Deep Convolutional Generative Adversarial Network

Vanessa Volz

TU Dortmund University

Dortmund, Germany

vanessa.volz@tu-dortmund.de

Jacob Schrum

Southwestern University

Georgetown, TX USA

schrumb2@southwestern.edu

Jialin Liu

Queen Mary University of London

London, UK

jialin.liu@qmul.ac.uk

Simon M. Lucas

Queen Mary University of London

London, UK

simon.lucas@qmul.ac.uk

Adam Smith

University of California

Santa Cruz, CA USA

amsmith@ucsc.edu

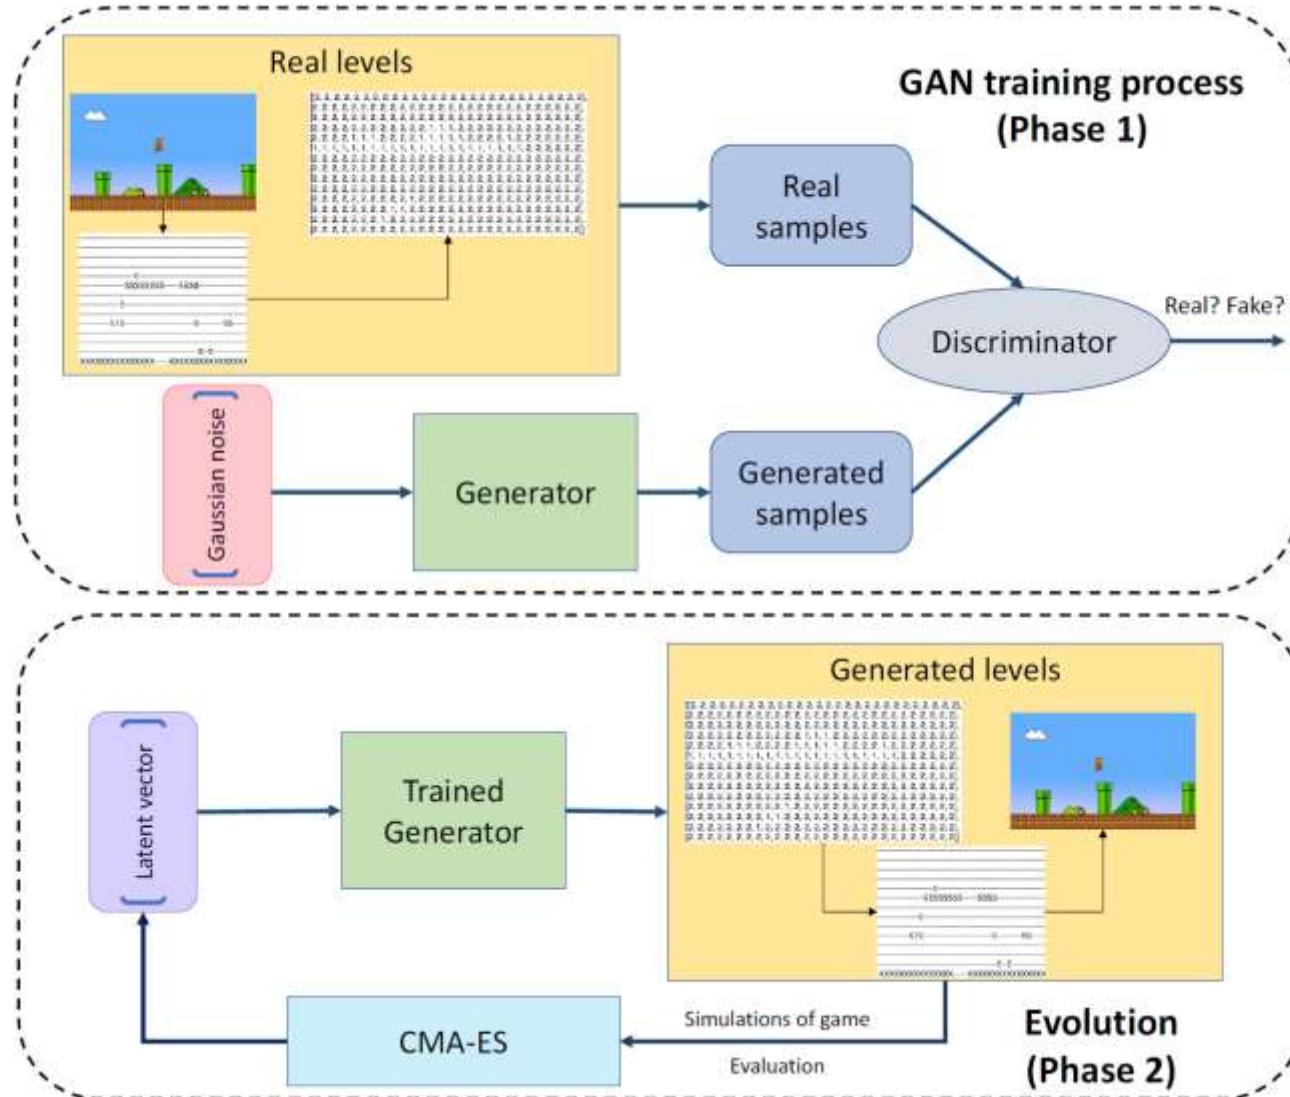
Sebastian Risi

IT University of Copenhagen

Copenhagen, Denmark

sebr@itu.dk

- Train a GAN on a Mario level, so that it can produce level segments
- Search the latent space of the trained GAN for level segments with particular properties
- Approach called *Latent Variable Evolution*



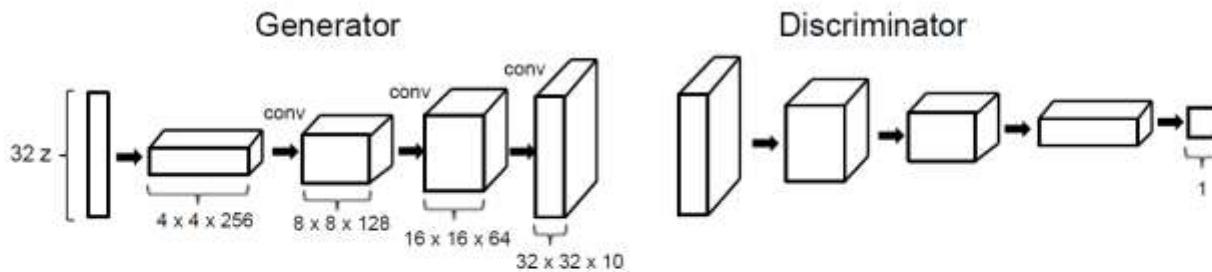


Figure 3: The Mario DCGAN architecture.

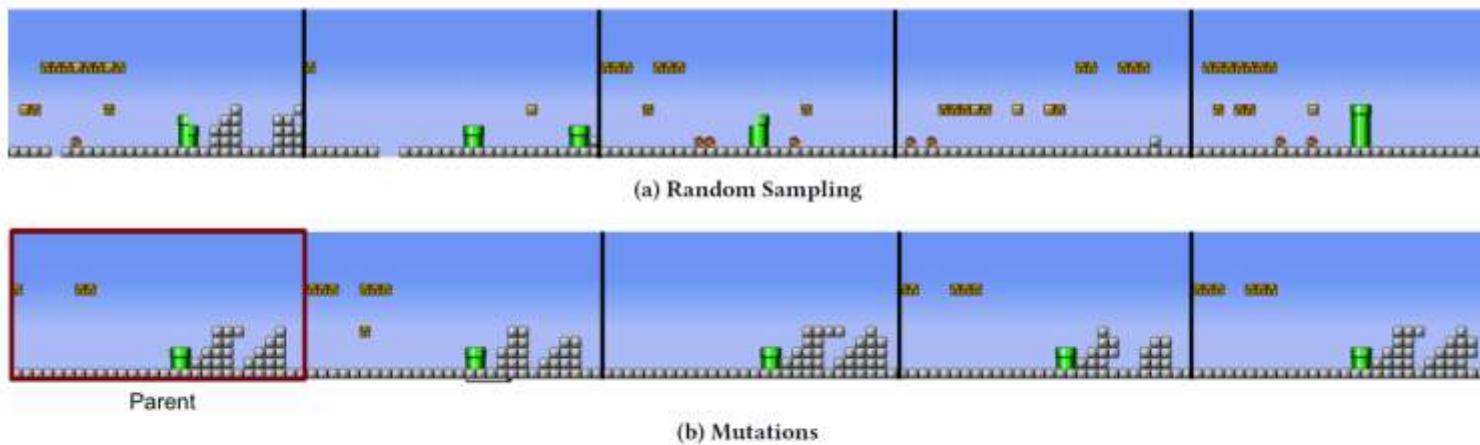


Figure 4: Generated Examples. Shown are samples produced by the GAN by (a) sampling random latent vectors, and (b) randomly mutating a specific latent vector. The main result is that the generator is able to produce a wide variety of different level layouts, but varied offspring still resemble their parent.

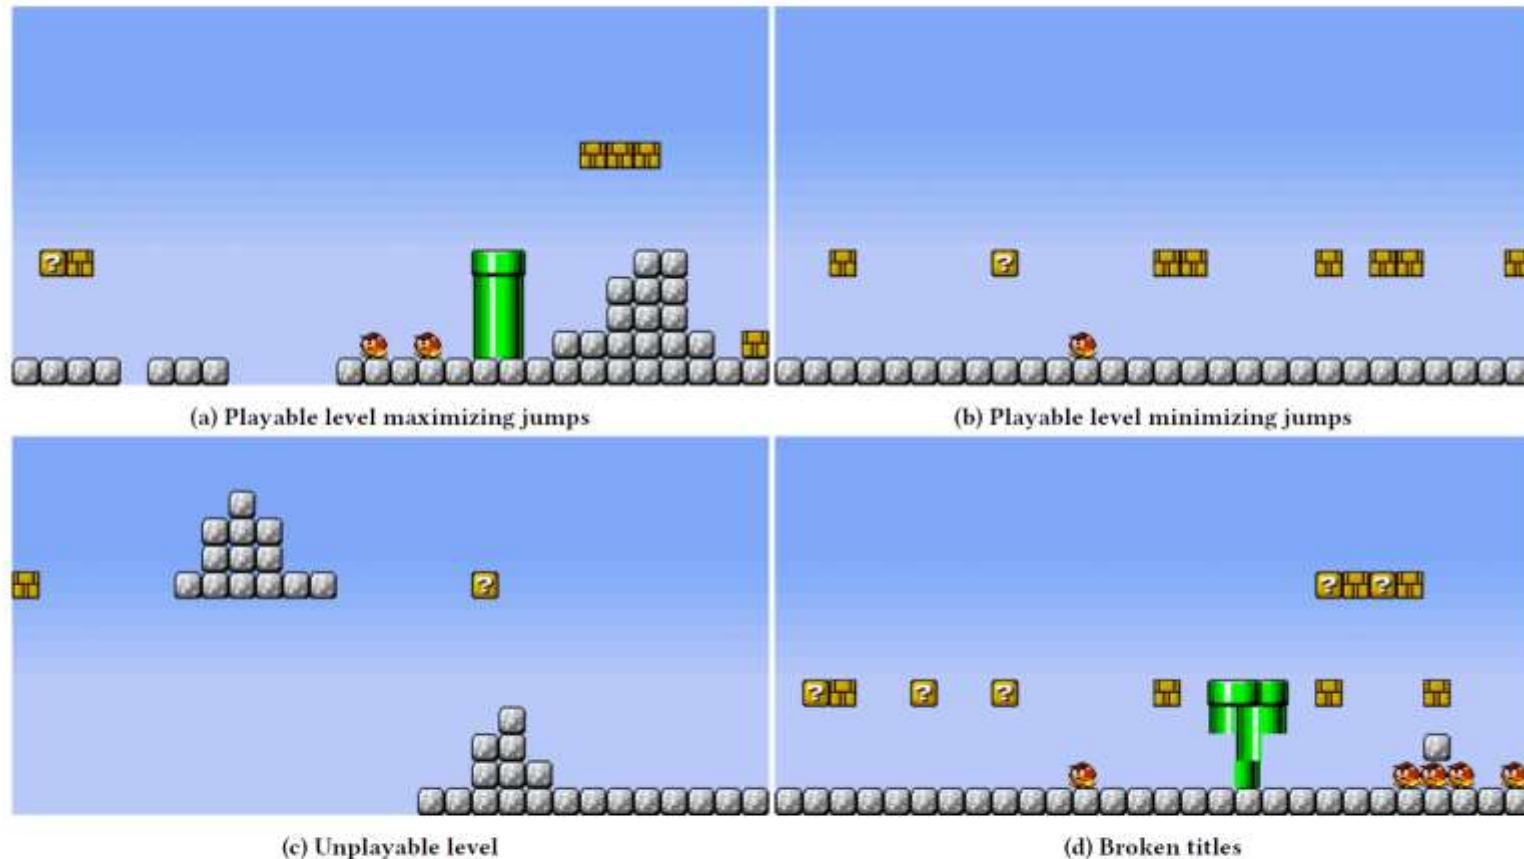
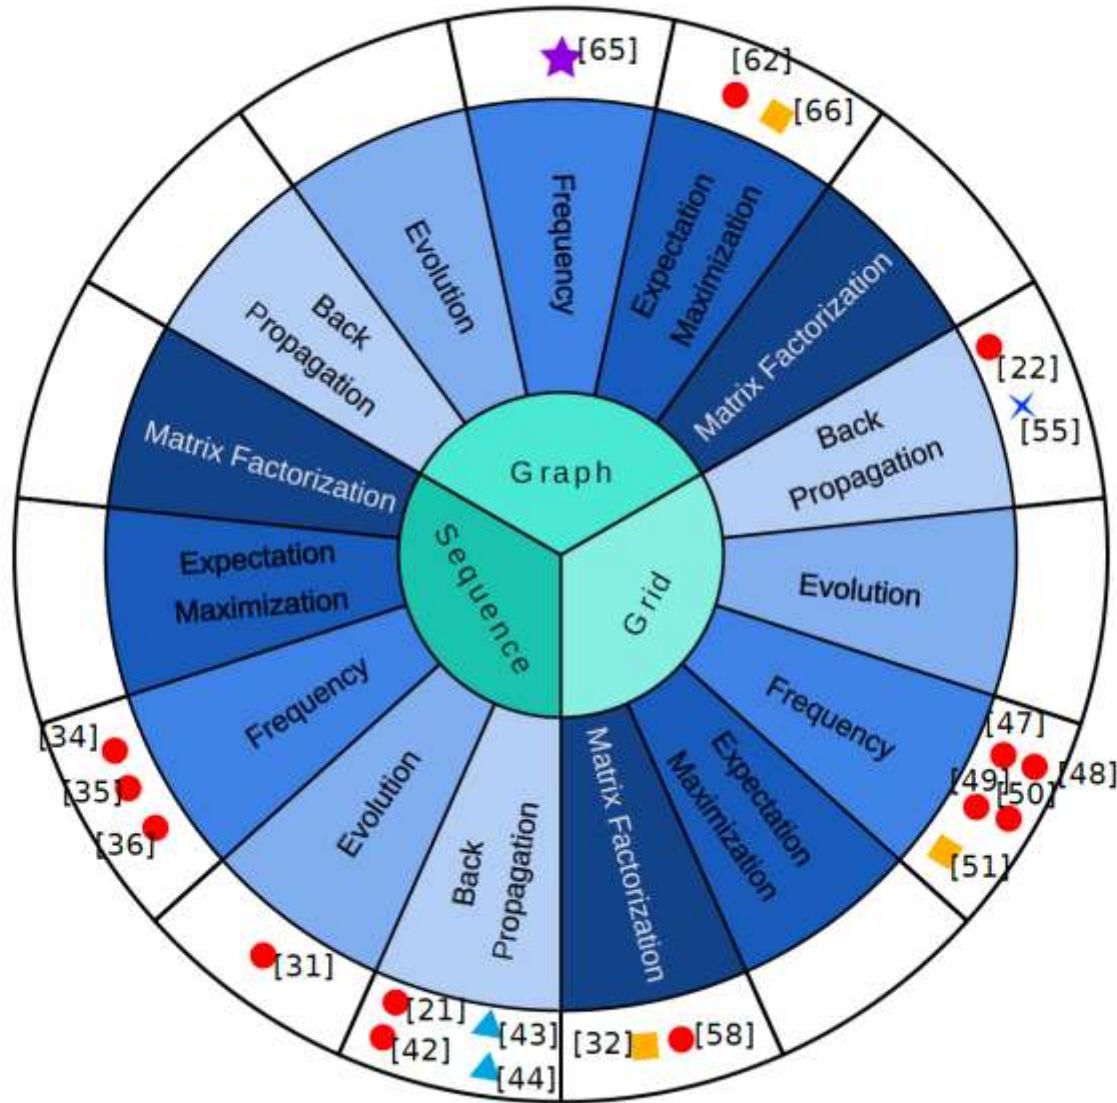
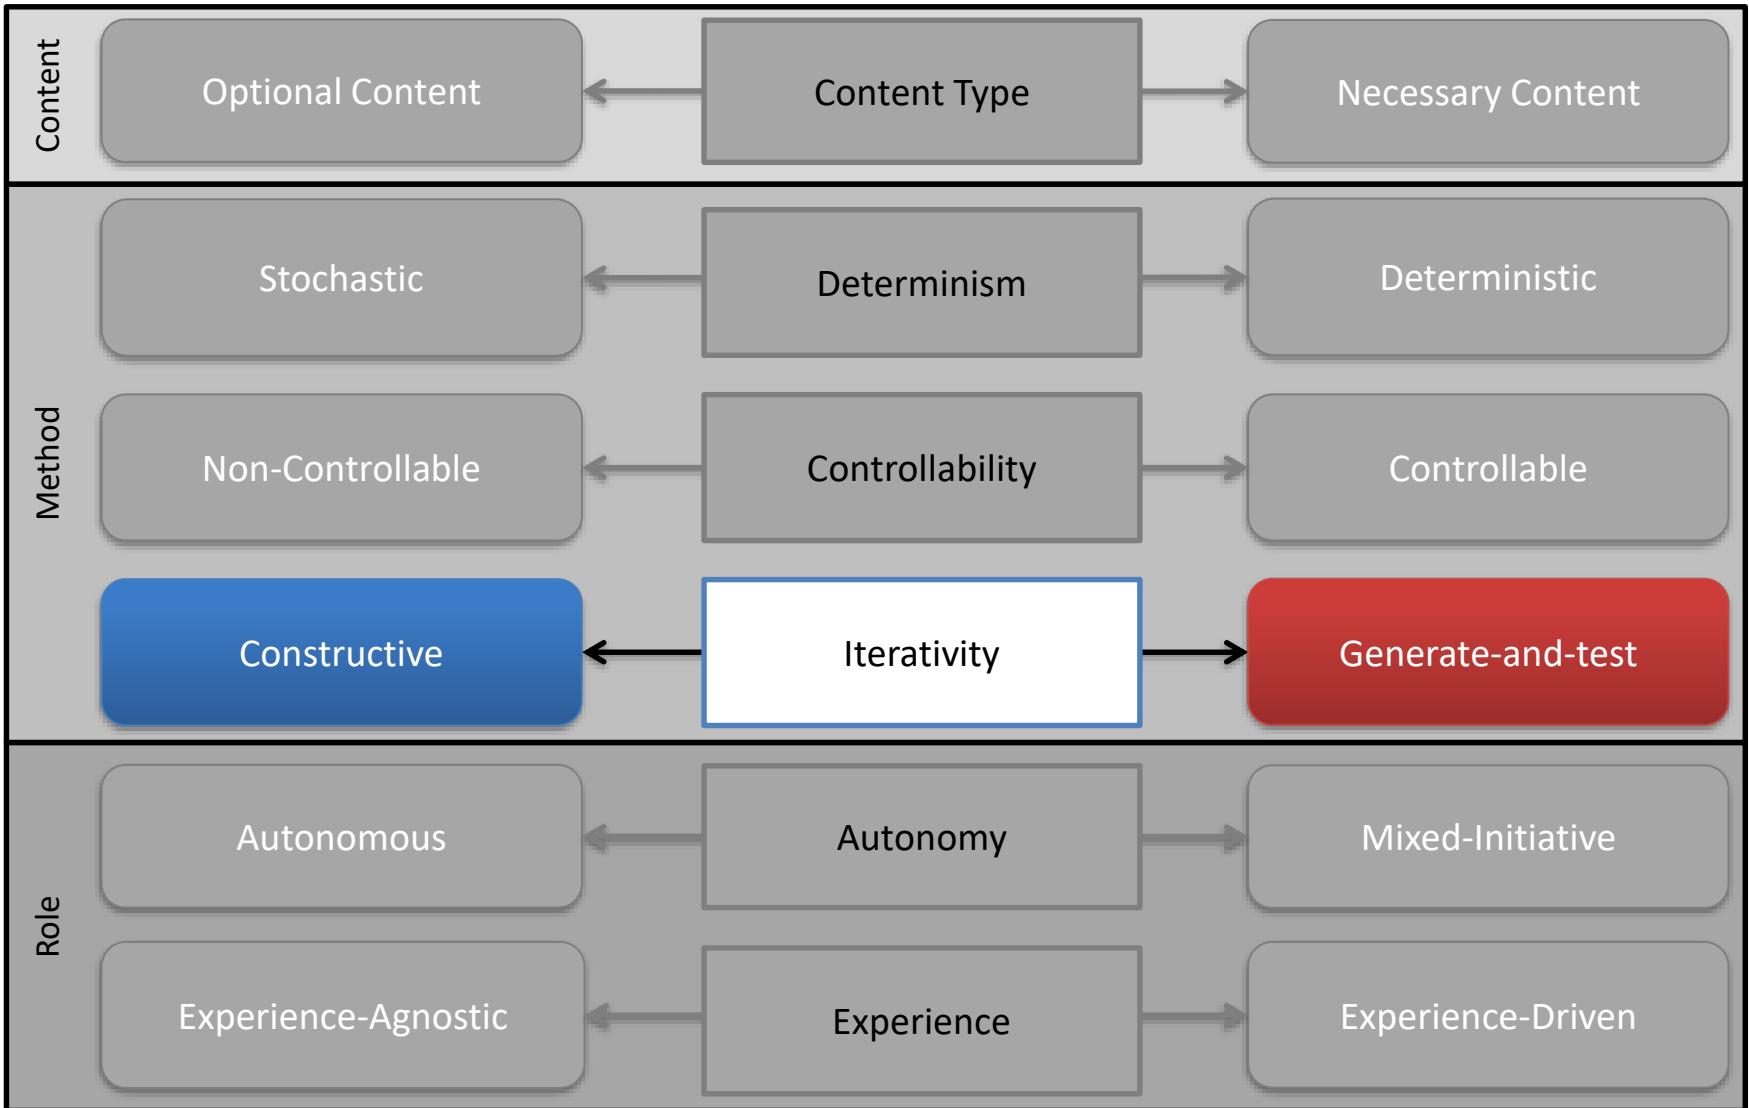


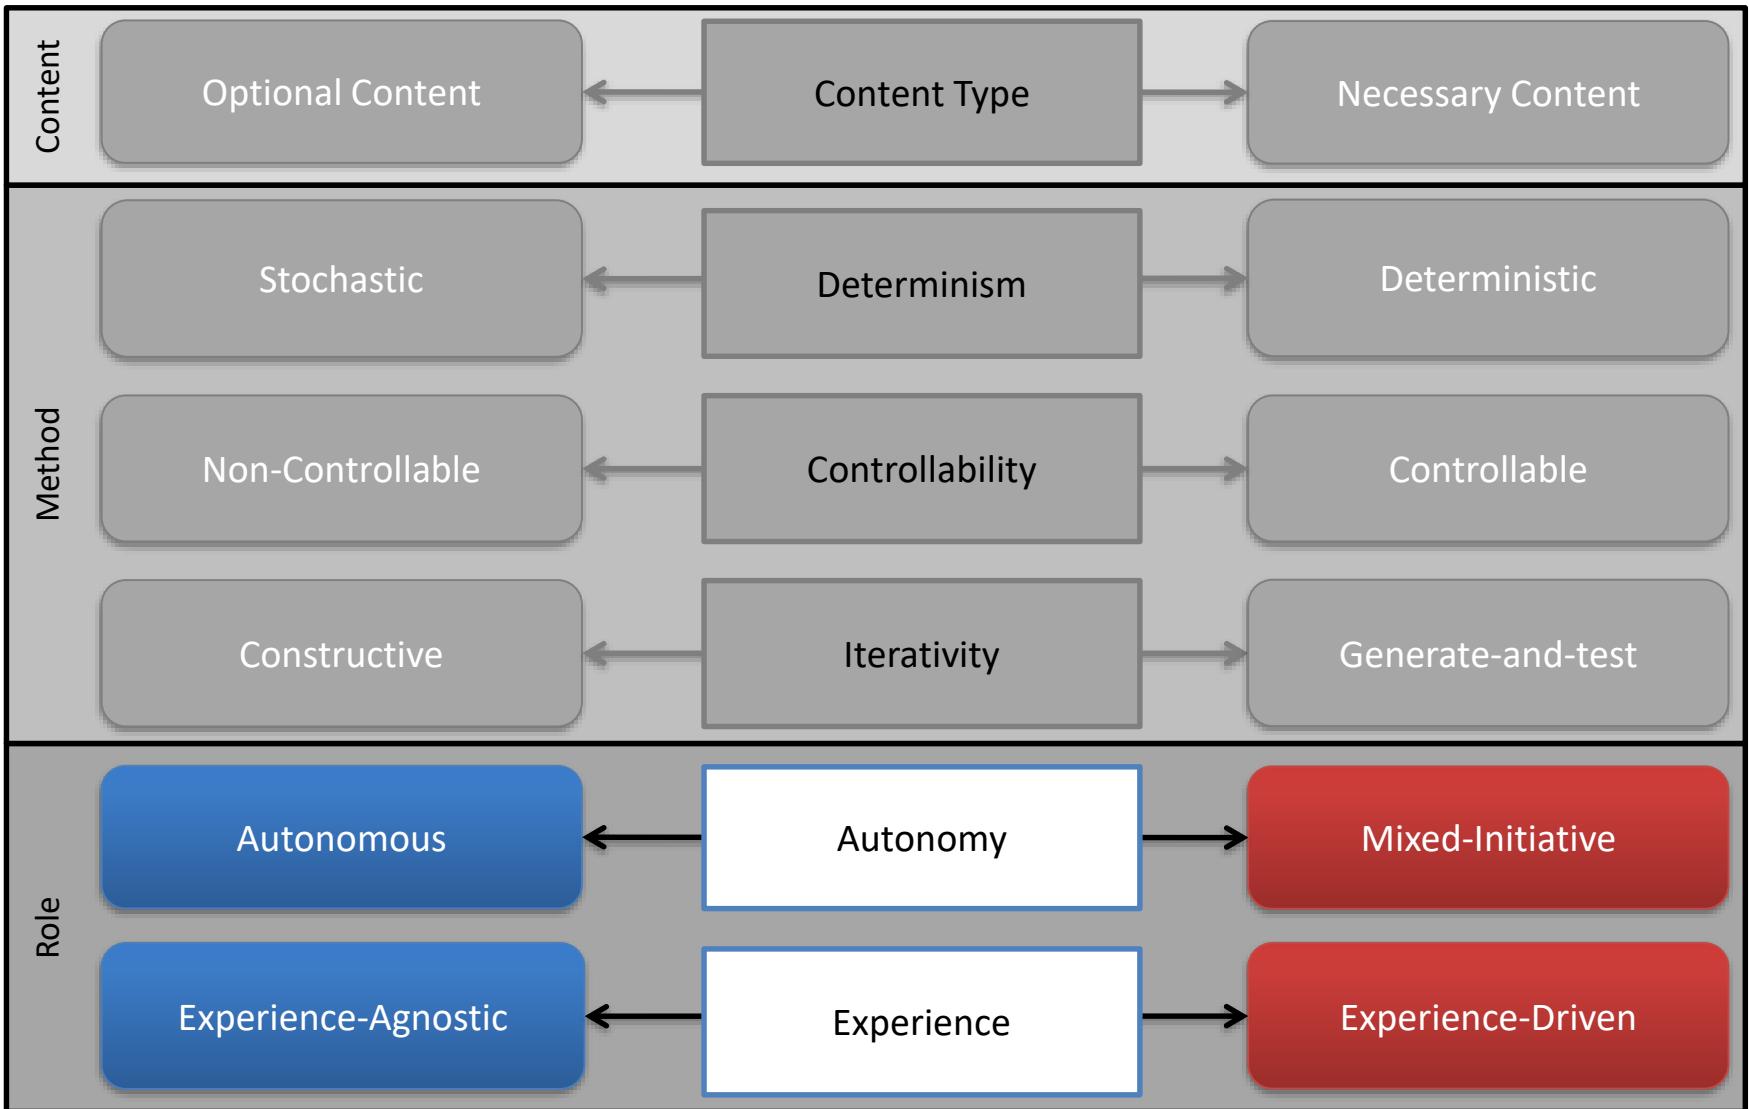
Figure 7: Agent-based optimization examples. (a) and (b) show good examples of levels in which the number of jumps is maximized (F_1), and minimized (F_2), respectively. (c) shows an example of one of the worst individuals found (not playable, F_1). An example of an individual that reaches high fitness (maximizing jumps, F_1) but has broken titles is shown in (d).

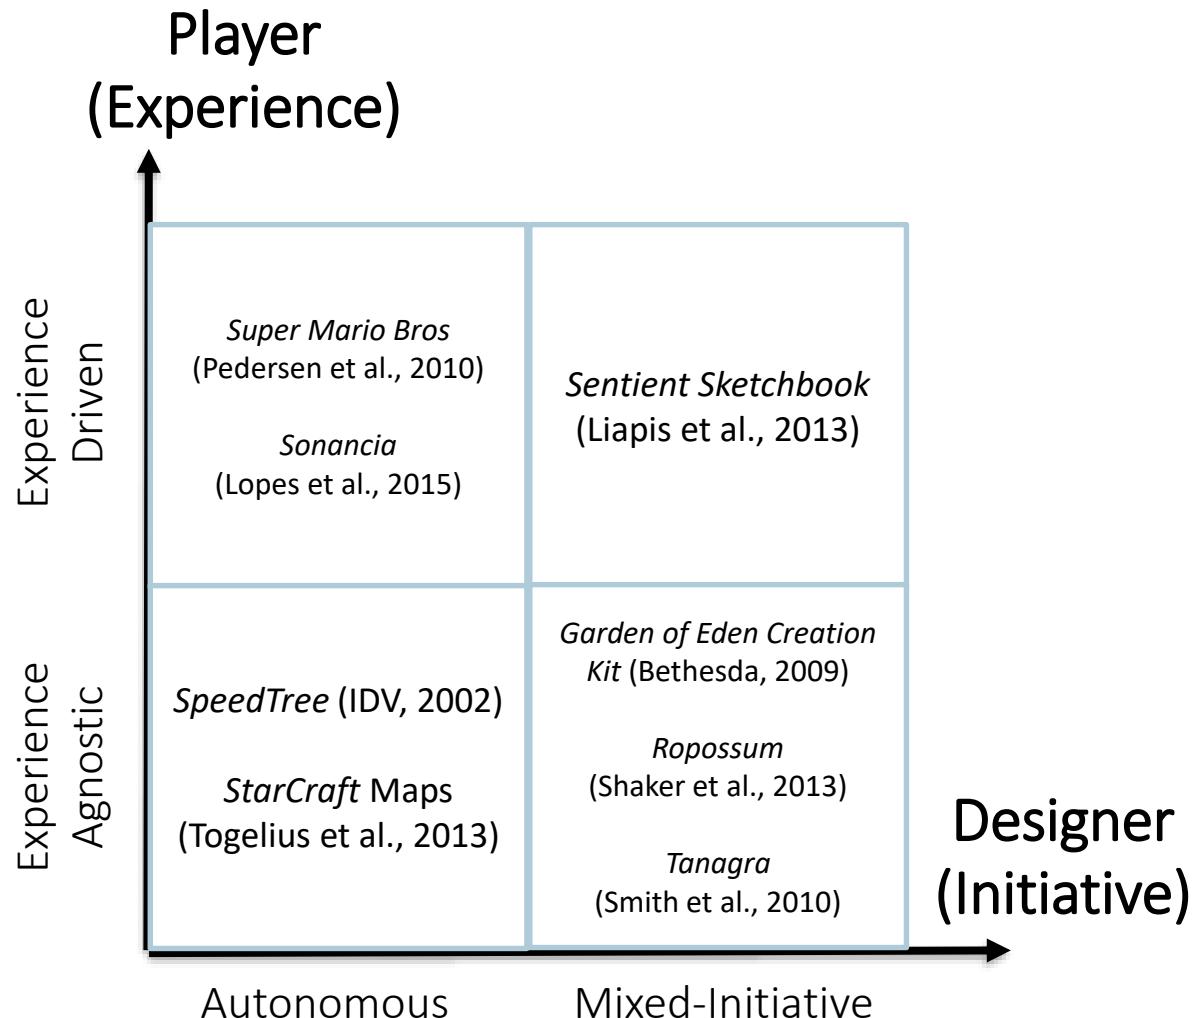


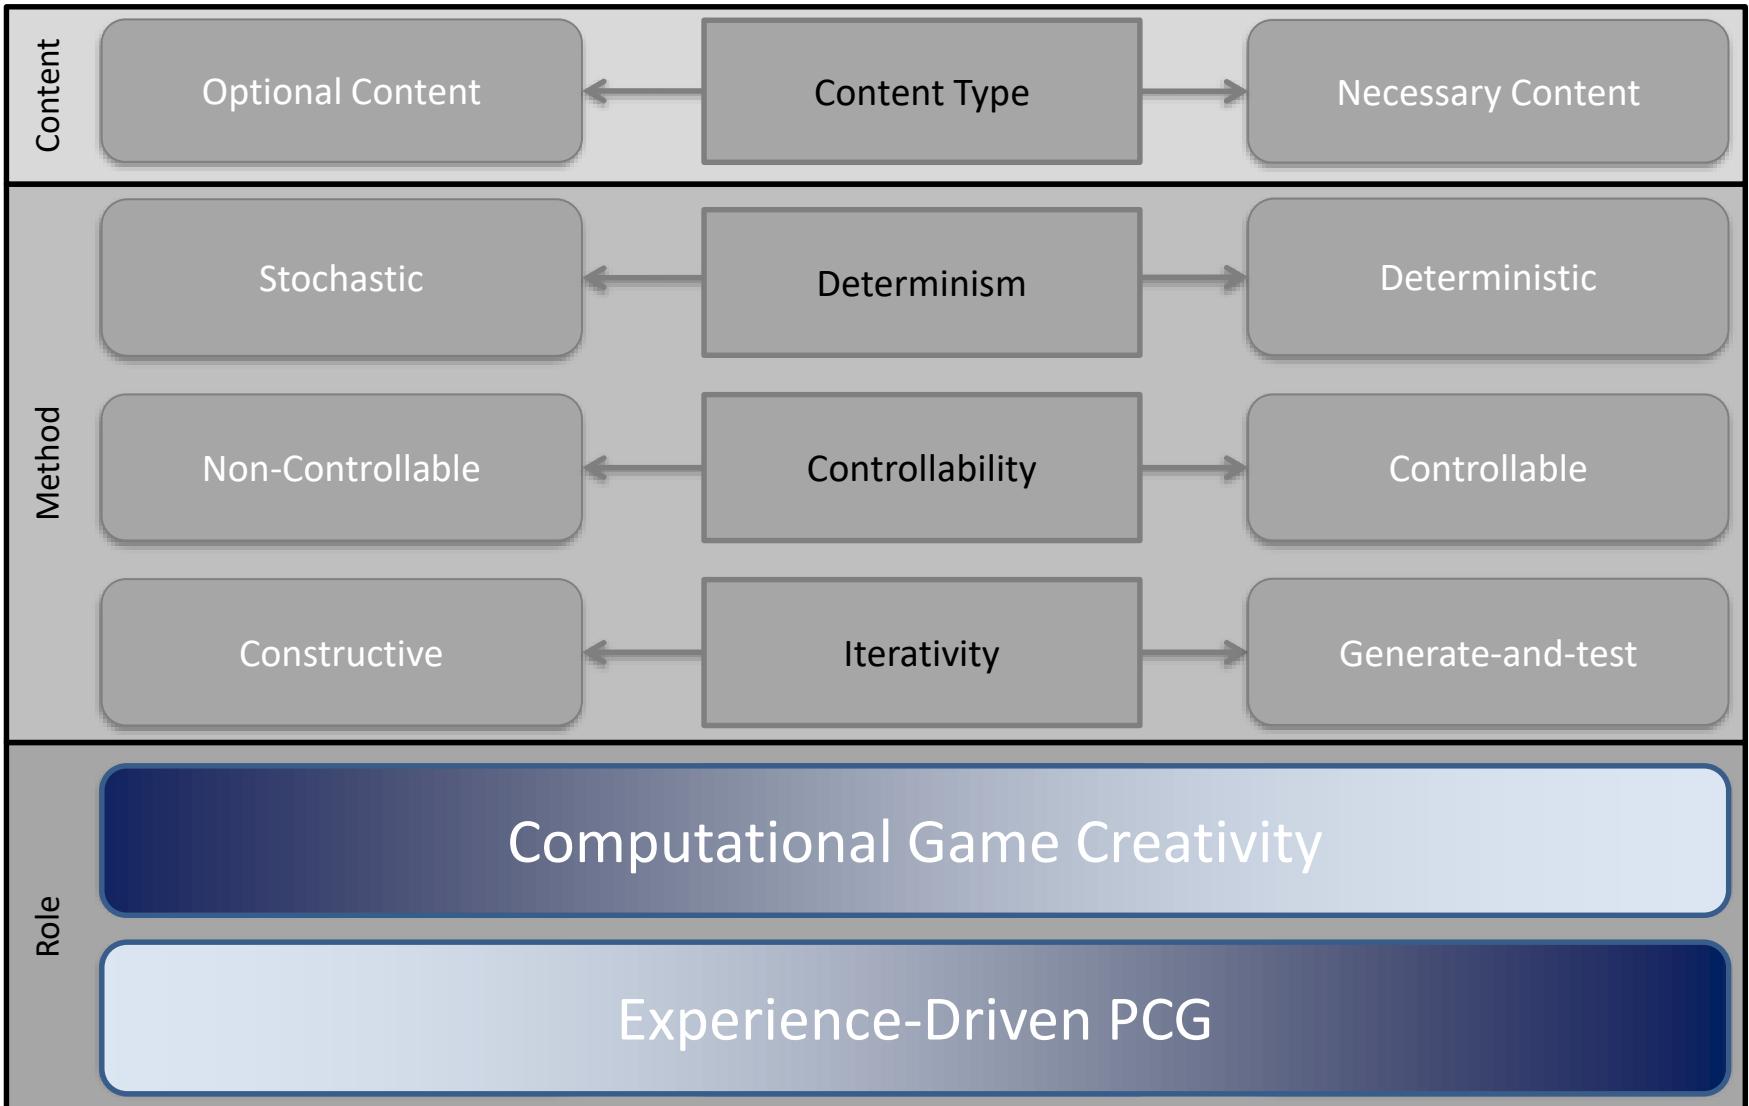
Figure 6: Level with increasing difficulty. Our LVE approach can create levels composed of multiple parts that gradually increase in difficulty (less ground tiles, more enemies). In the future this approach could be used to create a level in real-time that is tailored to the particular skill of the player (dynamic difficulty adaptation).











Computational Creativity **in** and **for** Games!

Liapis, Yannakakis, Togelius: "Computational Game Creativity," in *Proceedings of the Fifth International Conference on Computational Creativity*, 2014.

Computational Creativity



“Computational Creativity is a recent area of creativity research that brings together academics and practitioners from diverse disciplines, genres and modalities, to explore **the potential of computers to be autonomously creative** or to **collaborate as co-creators** with humans.”

-- PROSECCO Network of Excellence



PROMOTING THE
SCIENTIFIC EXPLORATION
OF COMPUTATIONAL
CREATIVITY

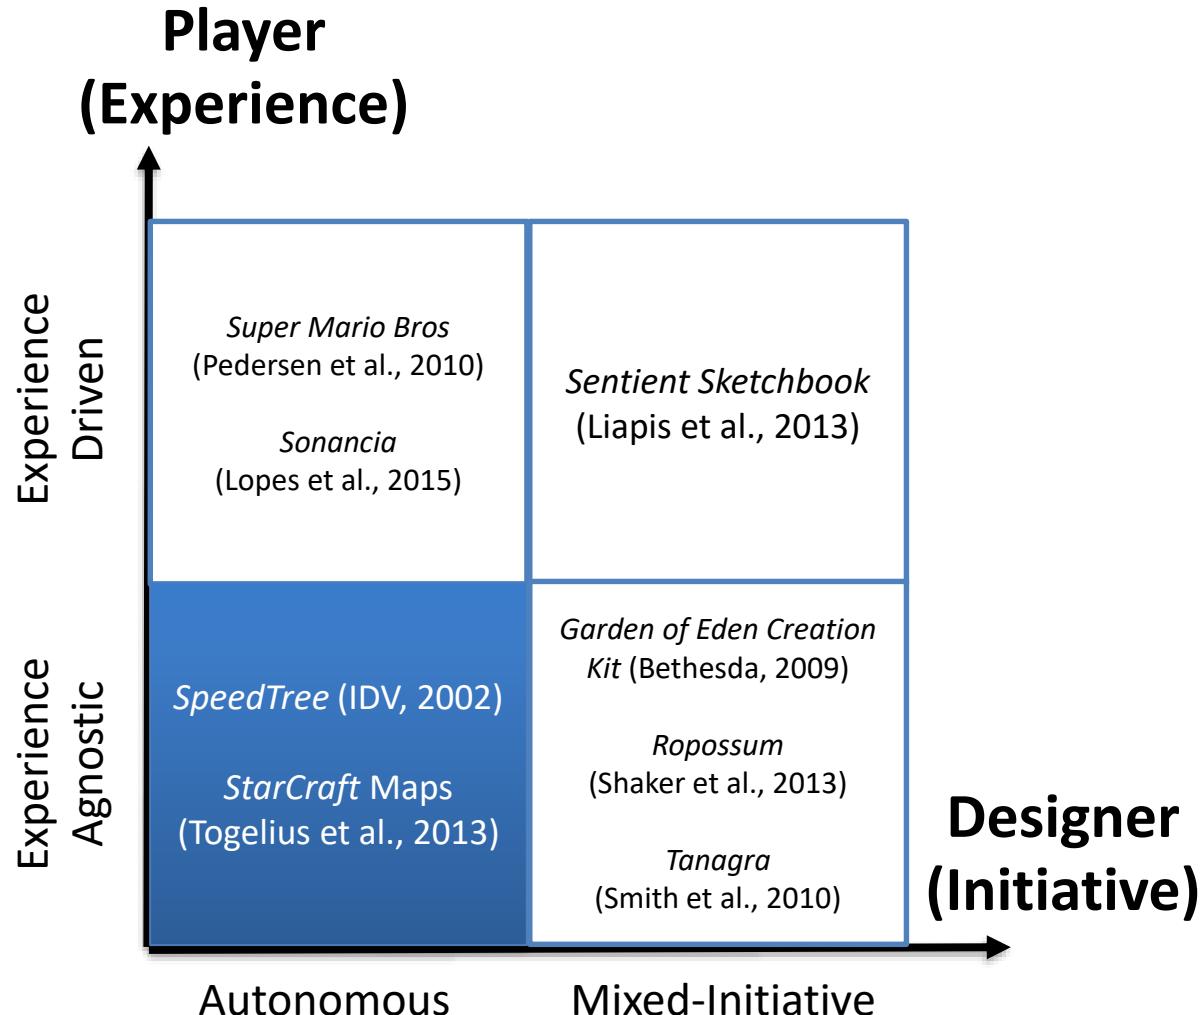
Computational Game Creativity



“The study of Computational Creativity **within**
and **for** digital games”

Within: Games as an ideal canvas for studying CC

For: Games benefit as products from artifacts of CC





SpeedTree Modeler 10.0.0.1000000000000000

File Edit View Select Tools Simulation Render Help

Complete mode Drawing mode

Generators

- Plants
- Groundcover Plants
- Random Plants
- Wind Driven
- Spirals
- Spines
- Whorls
- Branches
- Search
- Texture Generators
- Materials
- Dynamism

Branch

Source: Mesh

Moving: Mesh Vert Position

Amount: 1

Jagged: 1

Jagged: 1

Offset: 1

U Off: 1

V Off: 1

U Off: 1

V Off: 1

Cap

Source: Mesh Vert Position

Amount: 0

Vertical: 0

Vertical: 0

Vertical: 1

Angle: 0

Advanced Simulation

- Level of Detail
- Wind
- Physics

Spine Generator

This Spine Generator is designed for generating spine points. A "spine" is defined as the center axis of a branch or trunk, so both branches and trunks are part of the same generator.

Quake

00:00:00.000000 00:00:00.000000 00:00:00.000000 00:00:00.000000

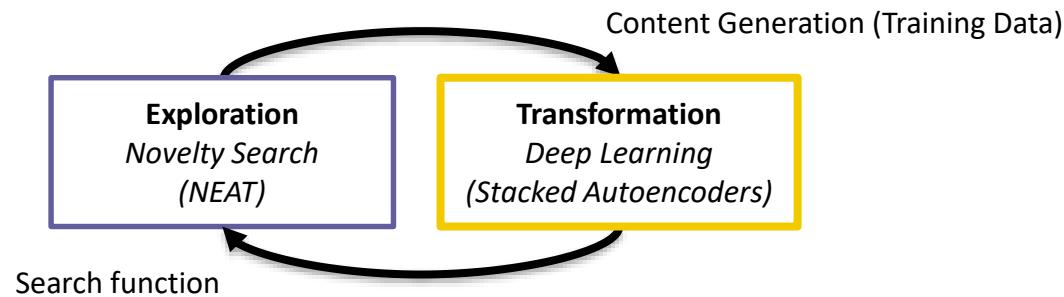
Copyright © 2010 SpeedTree Inc. All rights reserved.
This application must be registered to use.



Togelius, Julian, Mike Preuss, Nicola Beume, Simon Wessing, J. Hagelback, and Georgios N. Yannakakis. "**Multiobjective exploration of the Starcraft map space.**" In *Computational Intelligence and Games (CIG)*, IEEE Conference on, pp. 265-272, 2010.

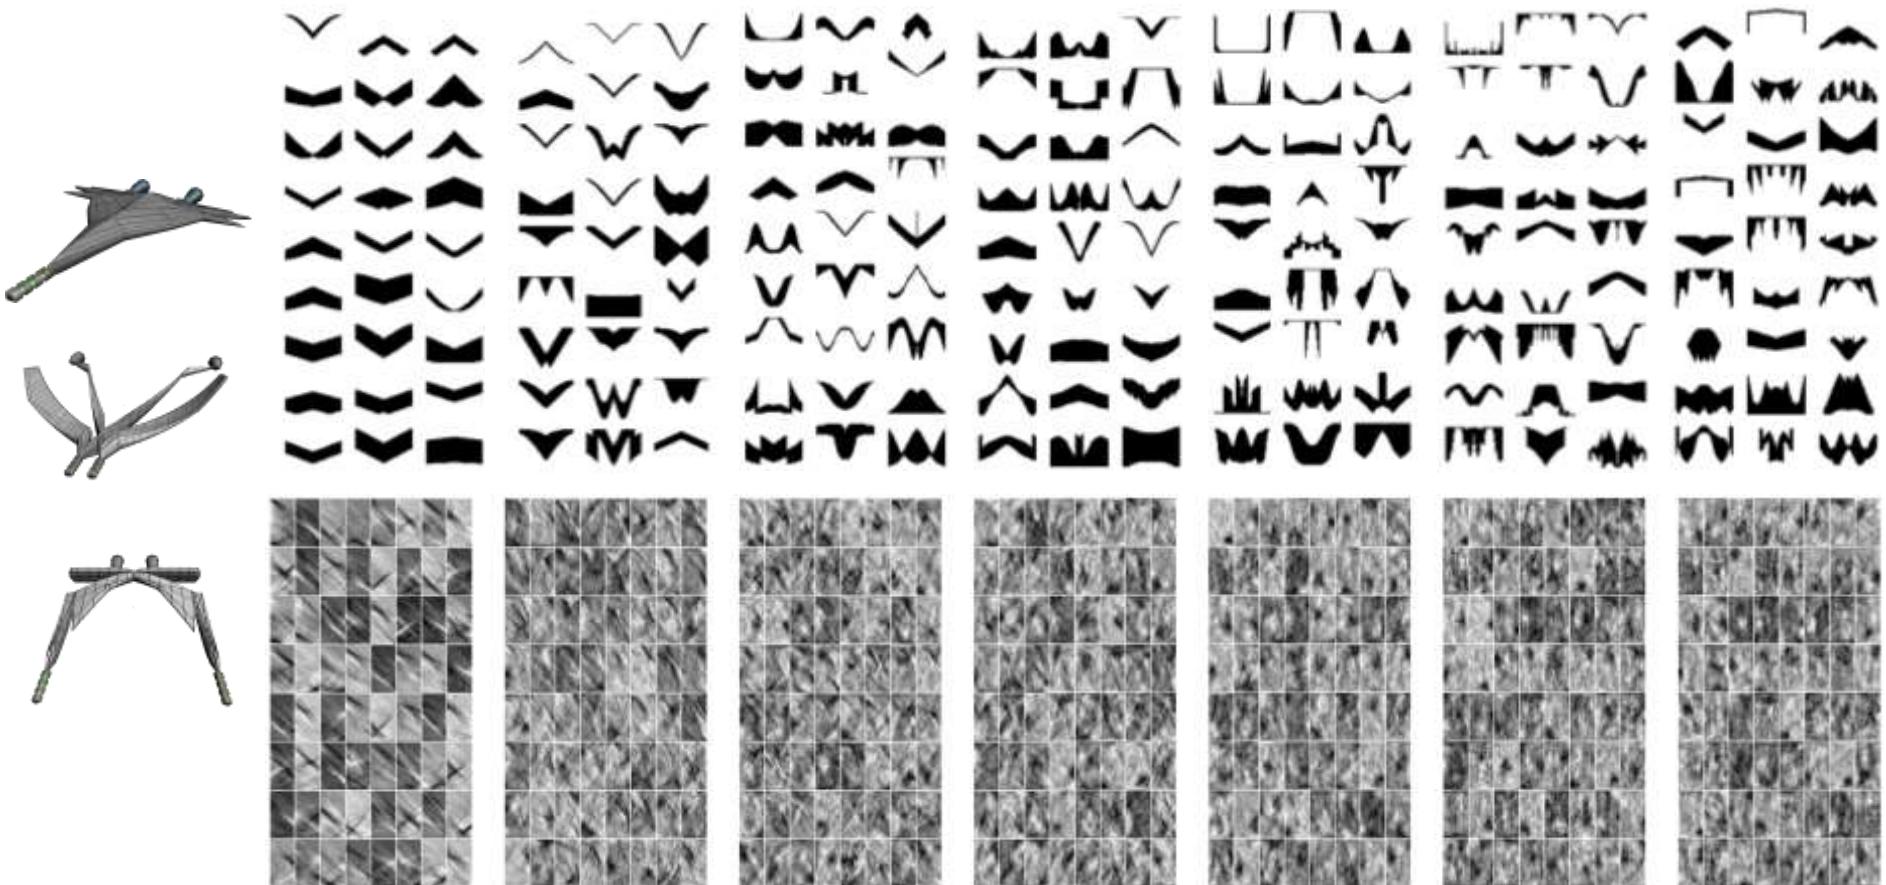
Deep Learning Meets Novelty Search

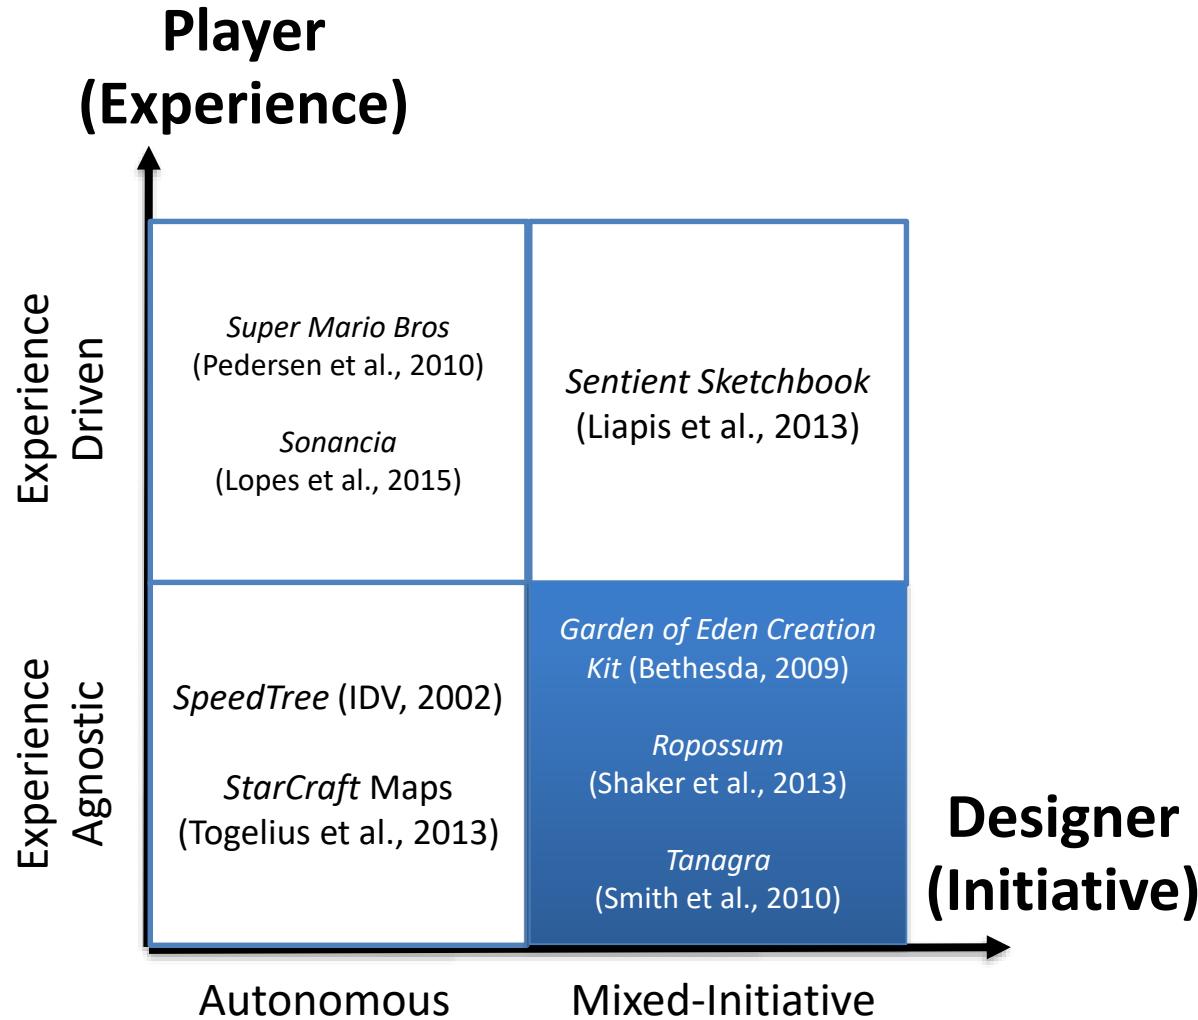
Liapis, Martínez, Togelius, and Yannakakis: "Transforming Exploratory Creativity with DeLeNoX," in *Proceedings of the Fourth International Conference on Computational Creativity*, 2013.



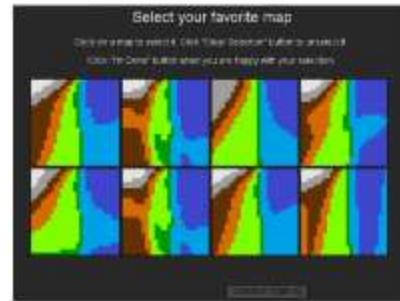
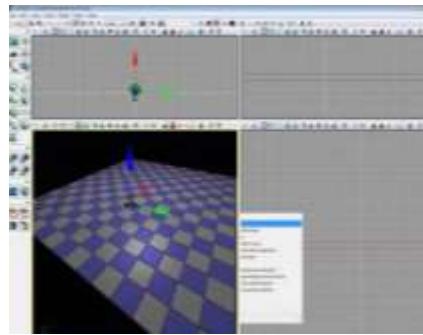
Deep Learning Meets Novelty Search

Liapis, Martínez, Togelius, and Yannakakis: "Transforming Exploratory Creativity with DeLeNoX," in *Proceedings of the Fourth International Conference on Computational Creativity*, 2013.





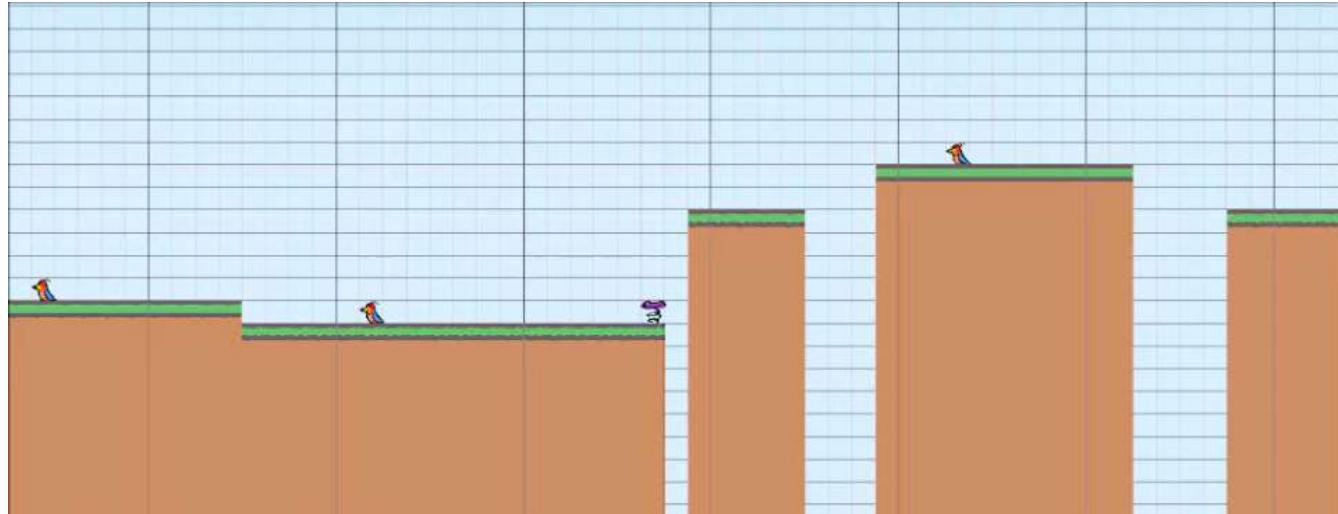
Human
Creativity



Computational
Creativity



Tanagra: Constraint Solver for MI-PCG



Tanagra: An AI-Supported Level Design Tool

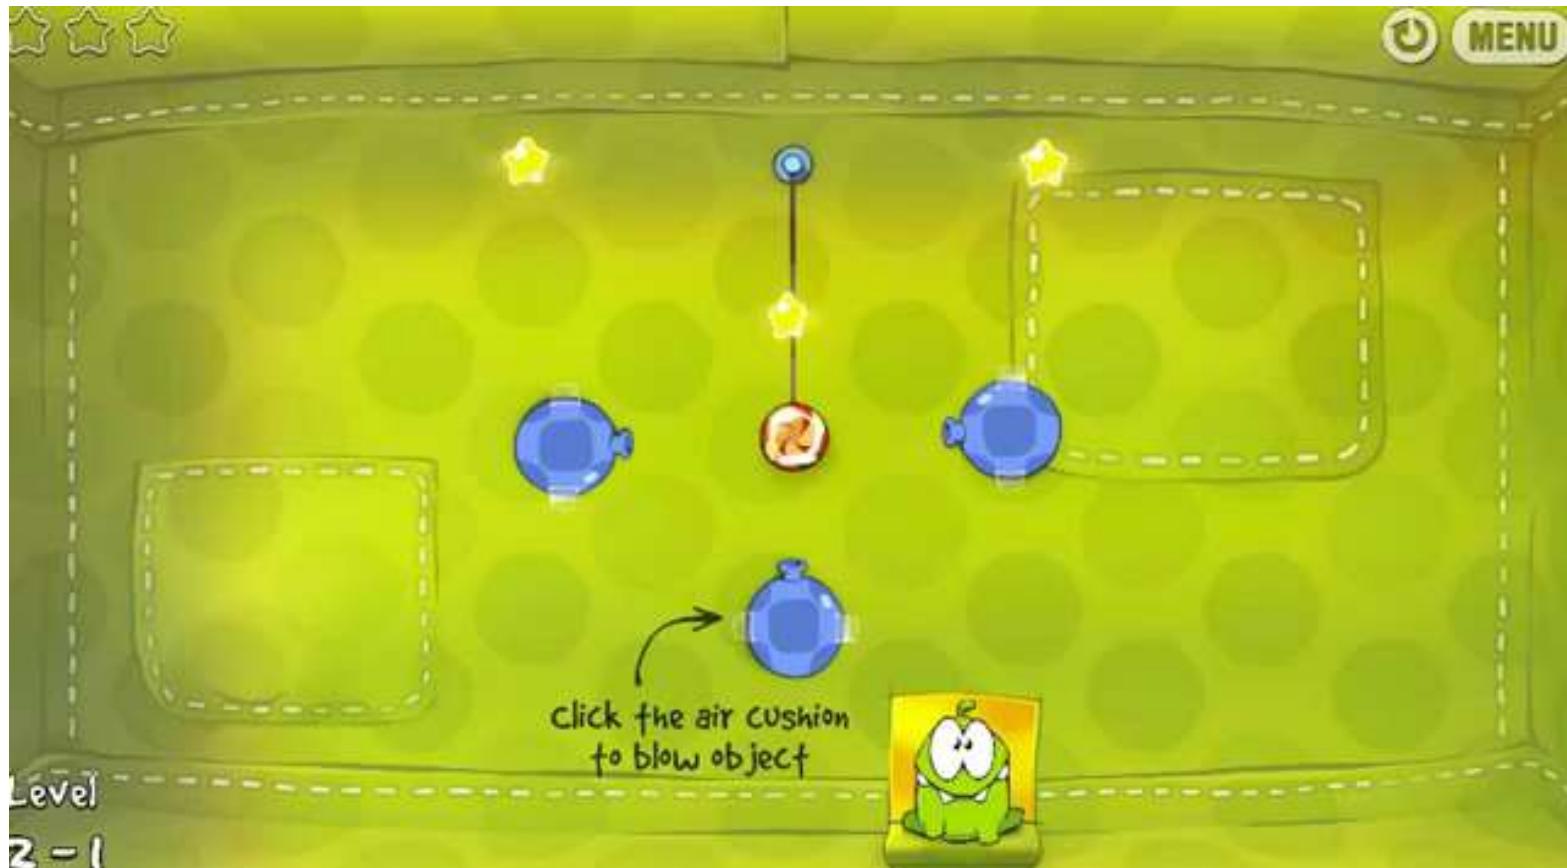
Gillian Smith, Jim Whitehead, Michael Mateas



Physics-based puzzle “cut the rope”

- evolutionary grammars for creating new puzzles
- playability module for testing how (if?) to solve a puzzle
- using the designer’s input in complete or partial designs

Ropossum





Generate Level Samples

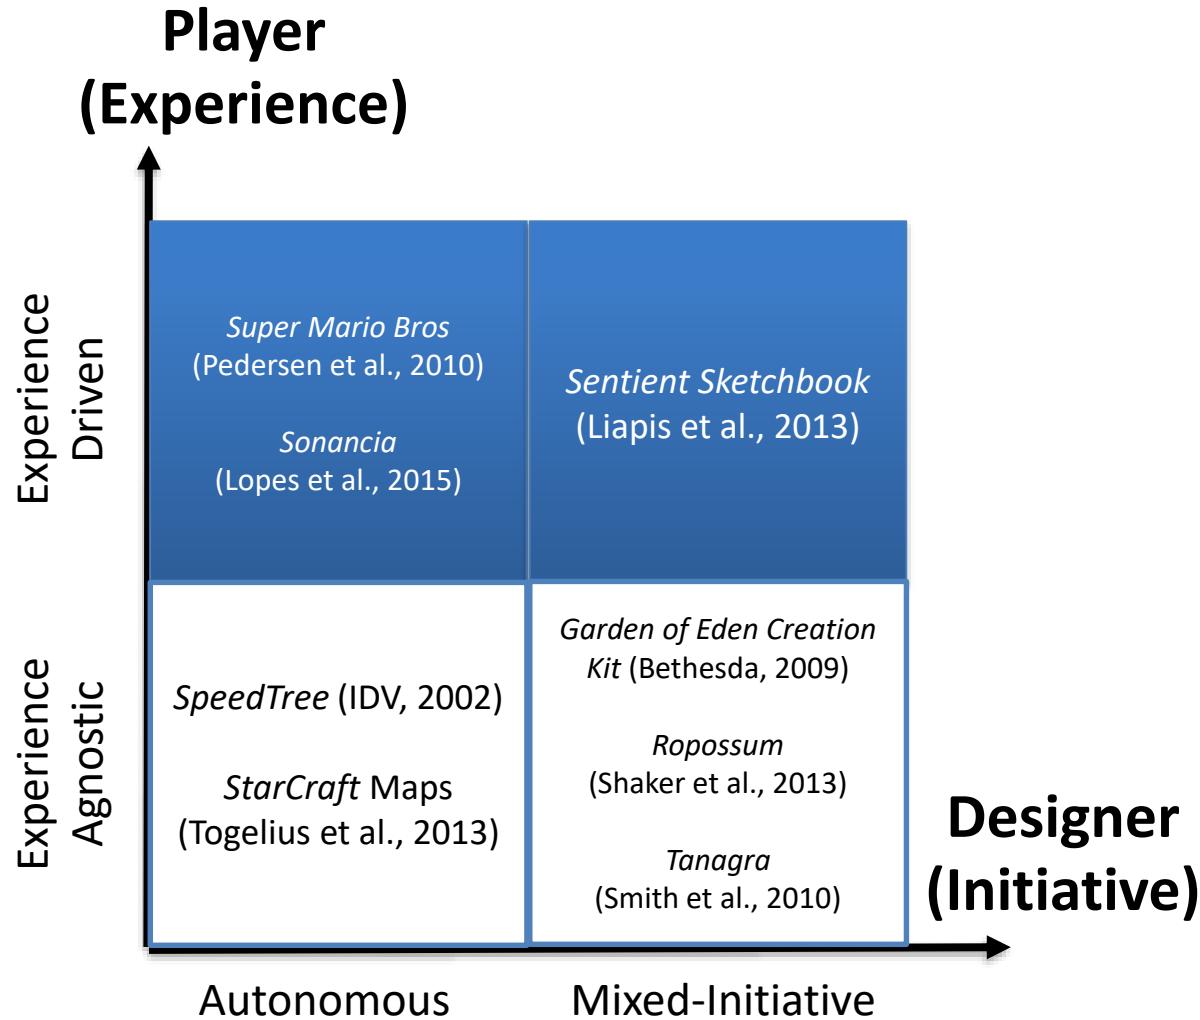


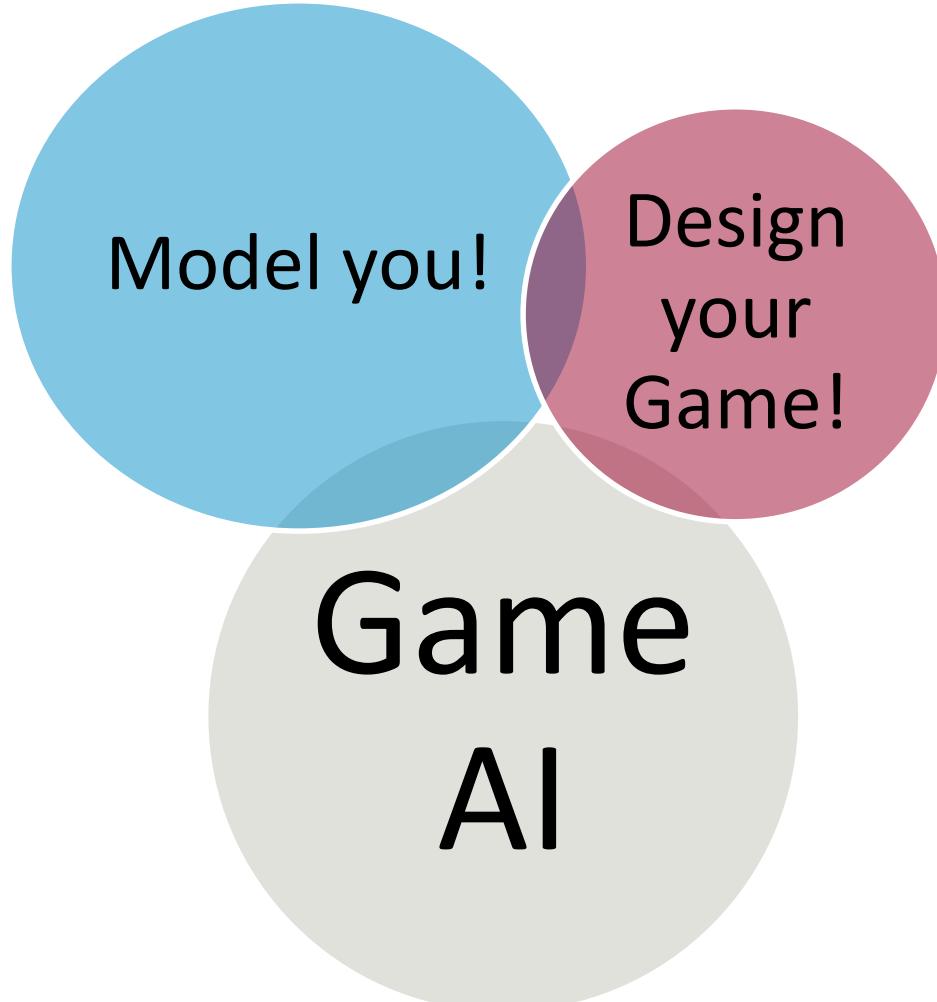
coralize

A reef / coral PCG tool for Unity



R. Abela, A. Liapis, G. N. Yannakakis: "**A Constructive Approach for the Generation of Underwater Environments**," in *Proceedings of FDG*, 2015.



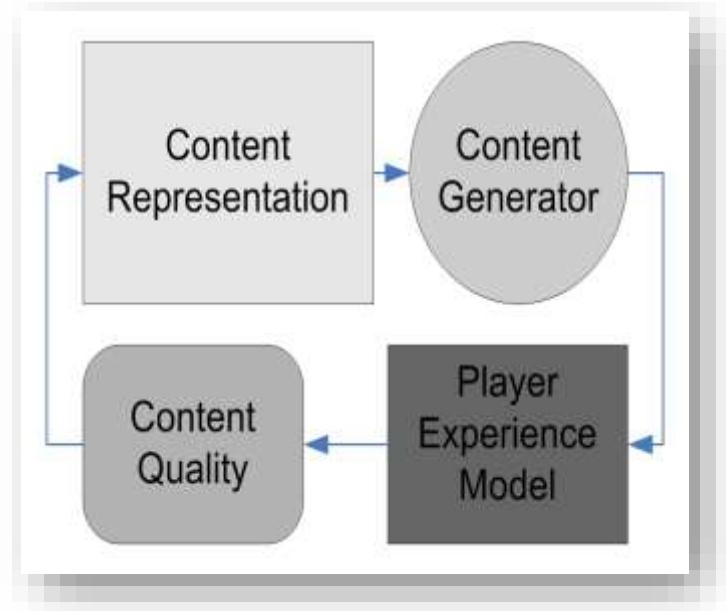


Yannakakis and Togelius, **Experience-driven Procedural Content Generation**, *IEEE Transactions on Affective Computing*, 2011.

A General PCG Framework



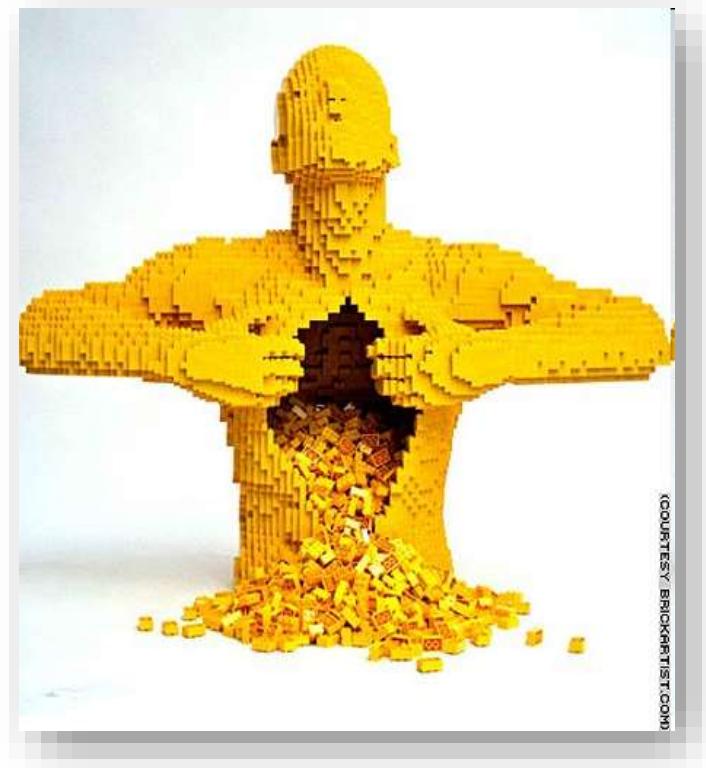
- Content is the **building block** of player experience
- **Search-based PCG:** use optimization algorithms (such as evolutionary algorithms) to search the space of content
- **Experience-driven PCG:** base evaluation function on player experience models



EDPCG: What is it?



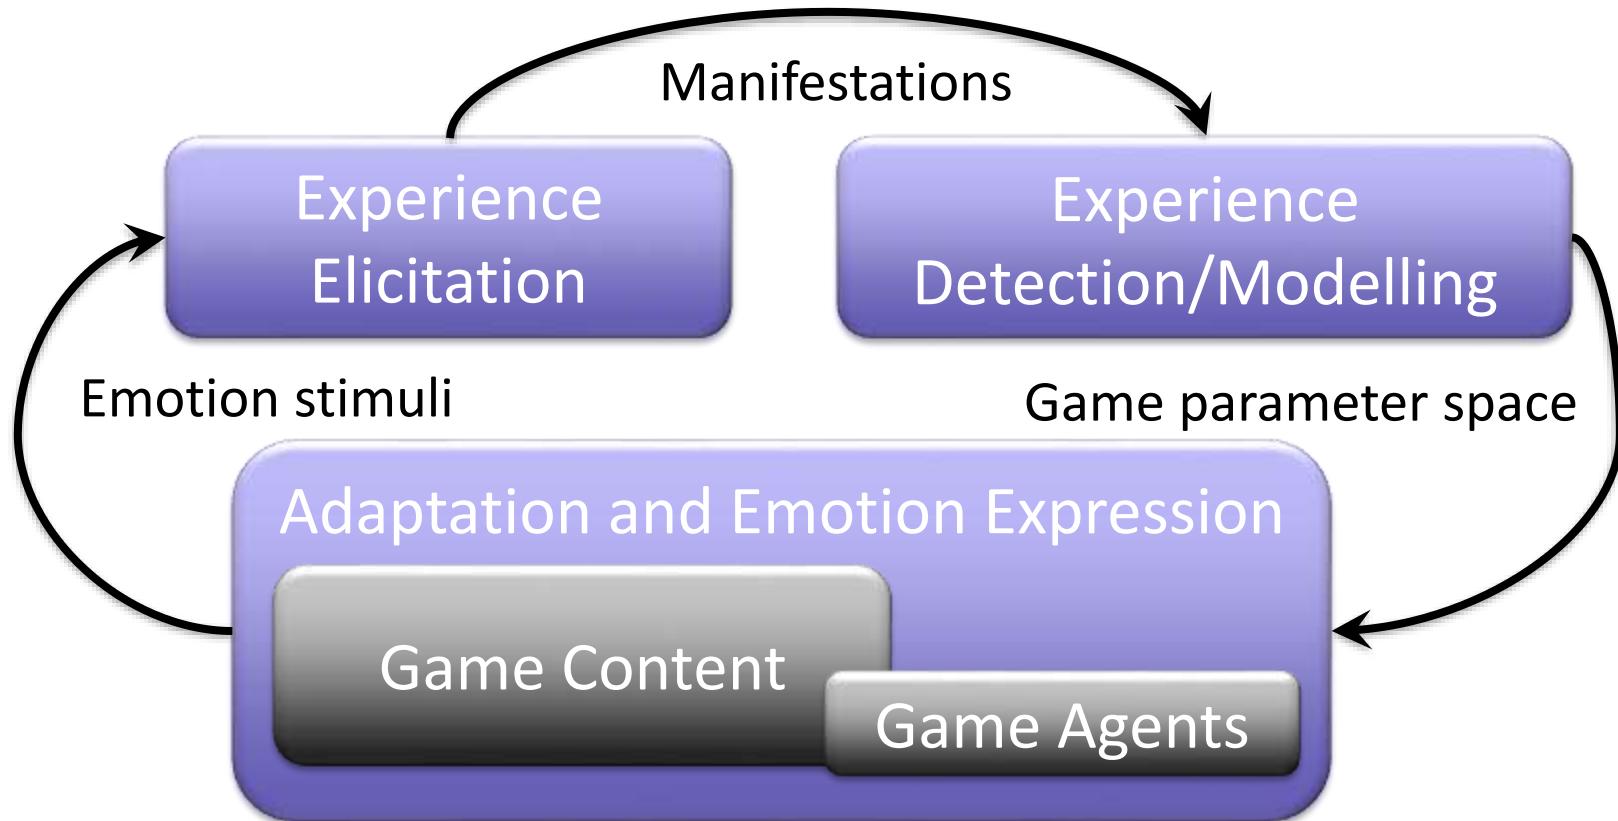
“A framework for personalised generation of content in human computer interaction (in particular in games). It views (game) content as the *building block* of user (player) experience”



Yannakakis, G. N., & Togelius, J. (2011). Experience-driven procedural content generation. *IEEE Transactions on Affective Computing*, 2(3), 147-161.

EDPCG Best Realizes the Affective Loop

Yannakakis and Paiva, **Emotion in Games**, in *Handbook of Affective Computing*, 2013



Experience-Driven

Procedural Content Generation

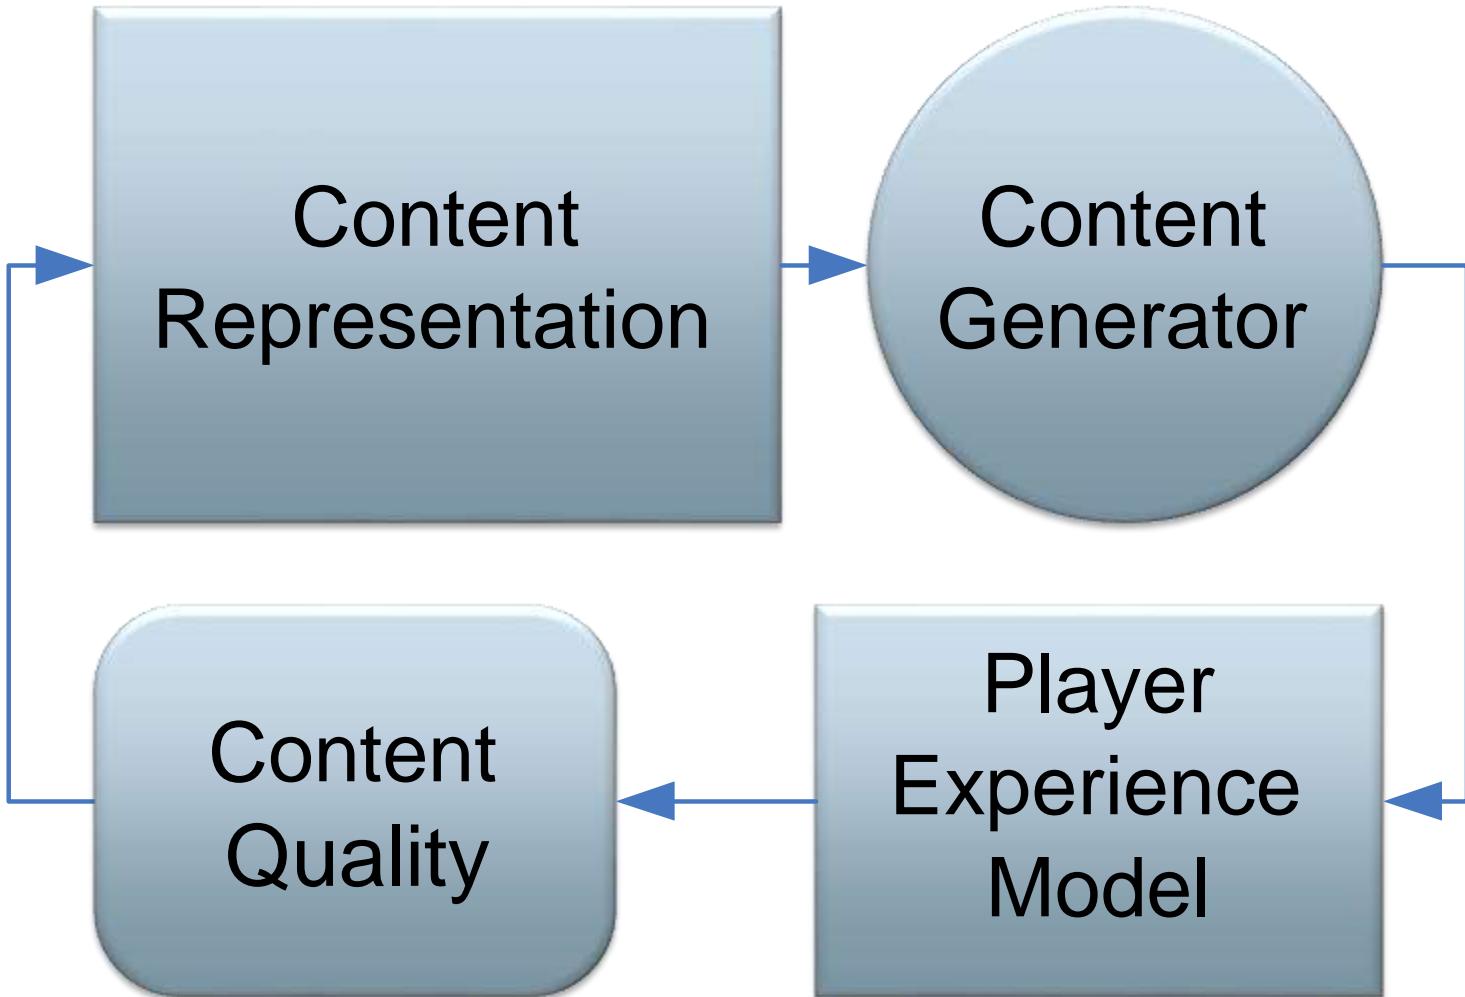


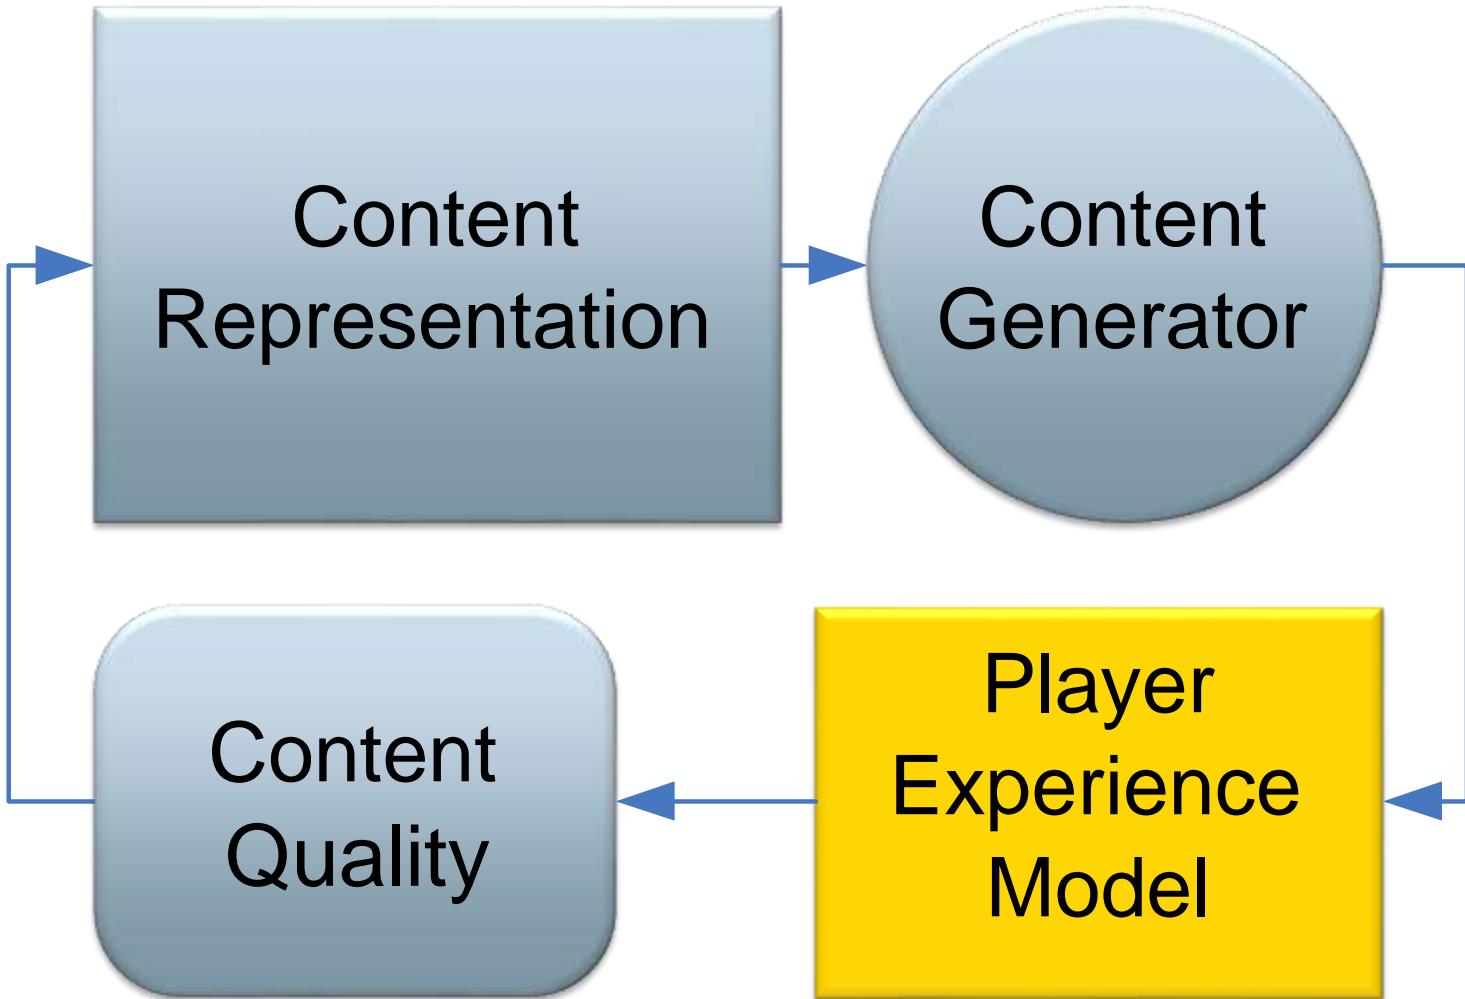
“... collection of **affective patterns**
elicited, **cognitive processes**
emerged and **behavioral traits**
observed during interaction
(gameplay)”



Key Components







Player Experience Modeling



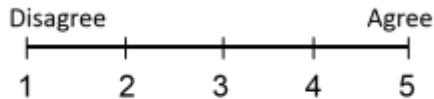
- Content evaluation functions could (should?) be based on player experience models
- Predict player experience from behavior/cognition/affect and/or game content
- Derived from empirical measurements of player experience



Player Experience Model

Subjective

X is frustrating



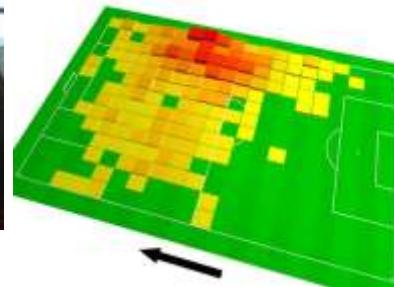
Is X or Y more frustrating?

- X Y
 - Both are **equally** frustrating
 - Neither is frustrating

Objective

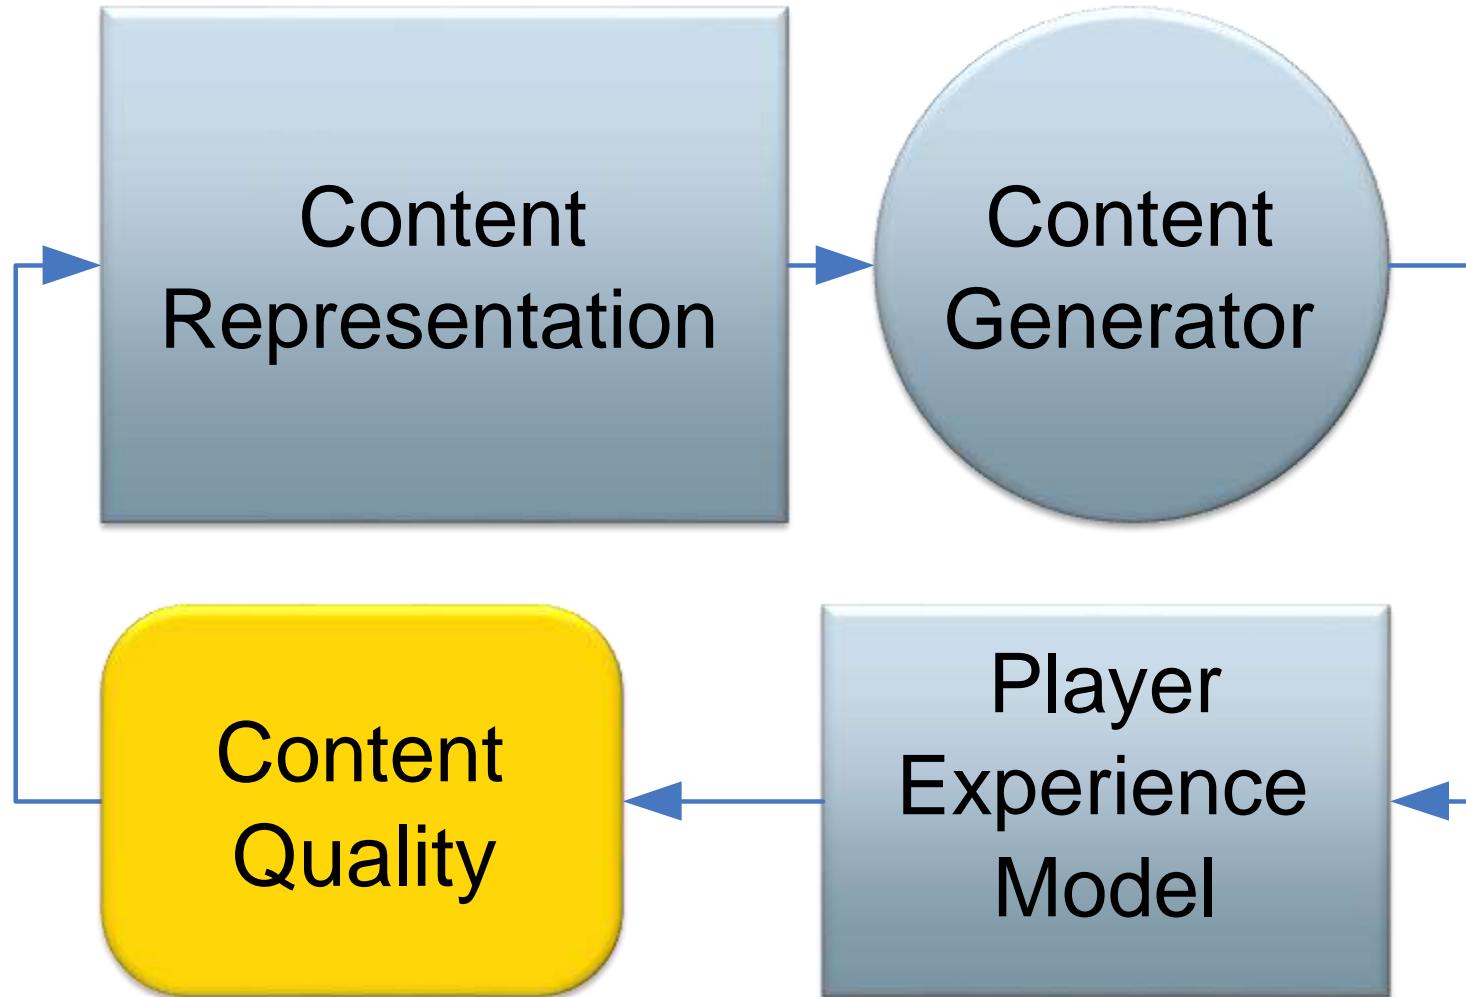


GamePlay-Based



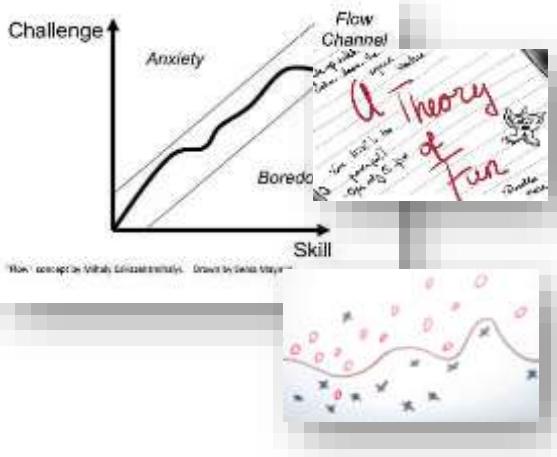
Player Experience Model

Subjective	Objective	GamePlay-Based
Free-Response	Model-based (e.g. arousal-valence dimensions of emotion)	Model-based (e.g. OCC, BDI)
Forced Rating vs. Preference	Model-free (e.g. facial expression annotation)	Model-free (e.g. player modeling)



Content Quality

Direct



Simulation-based

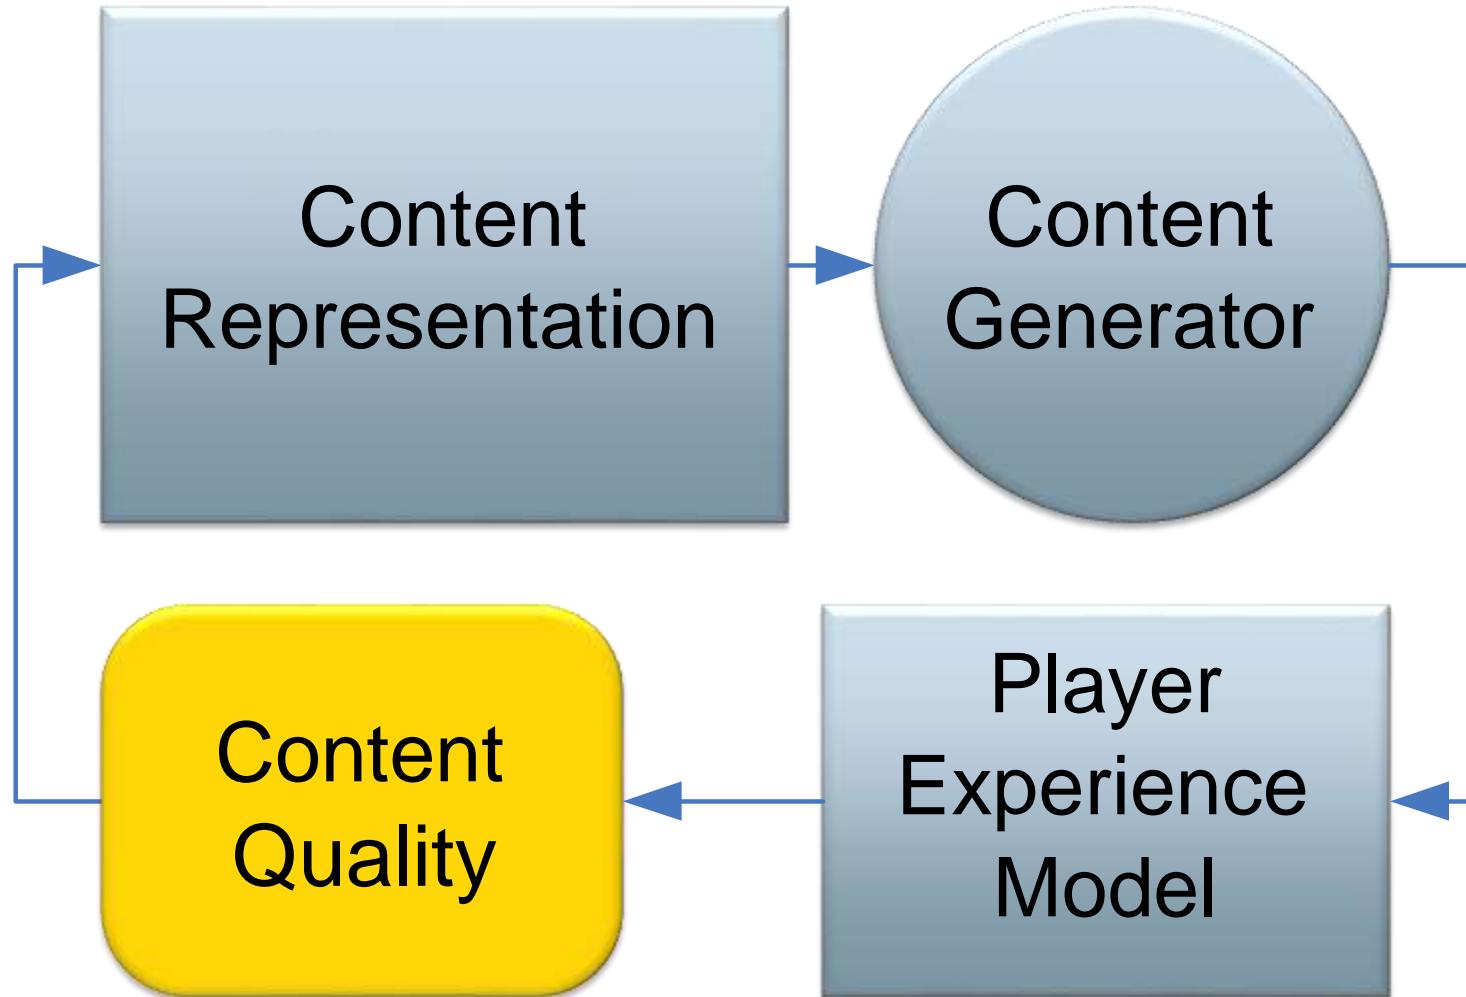


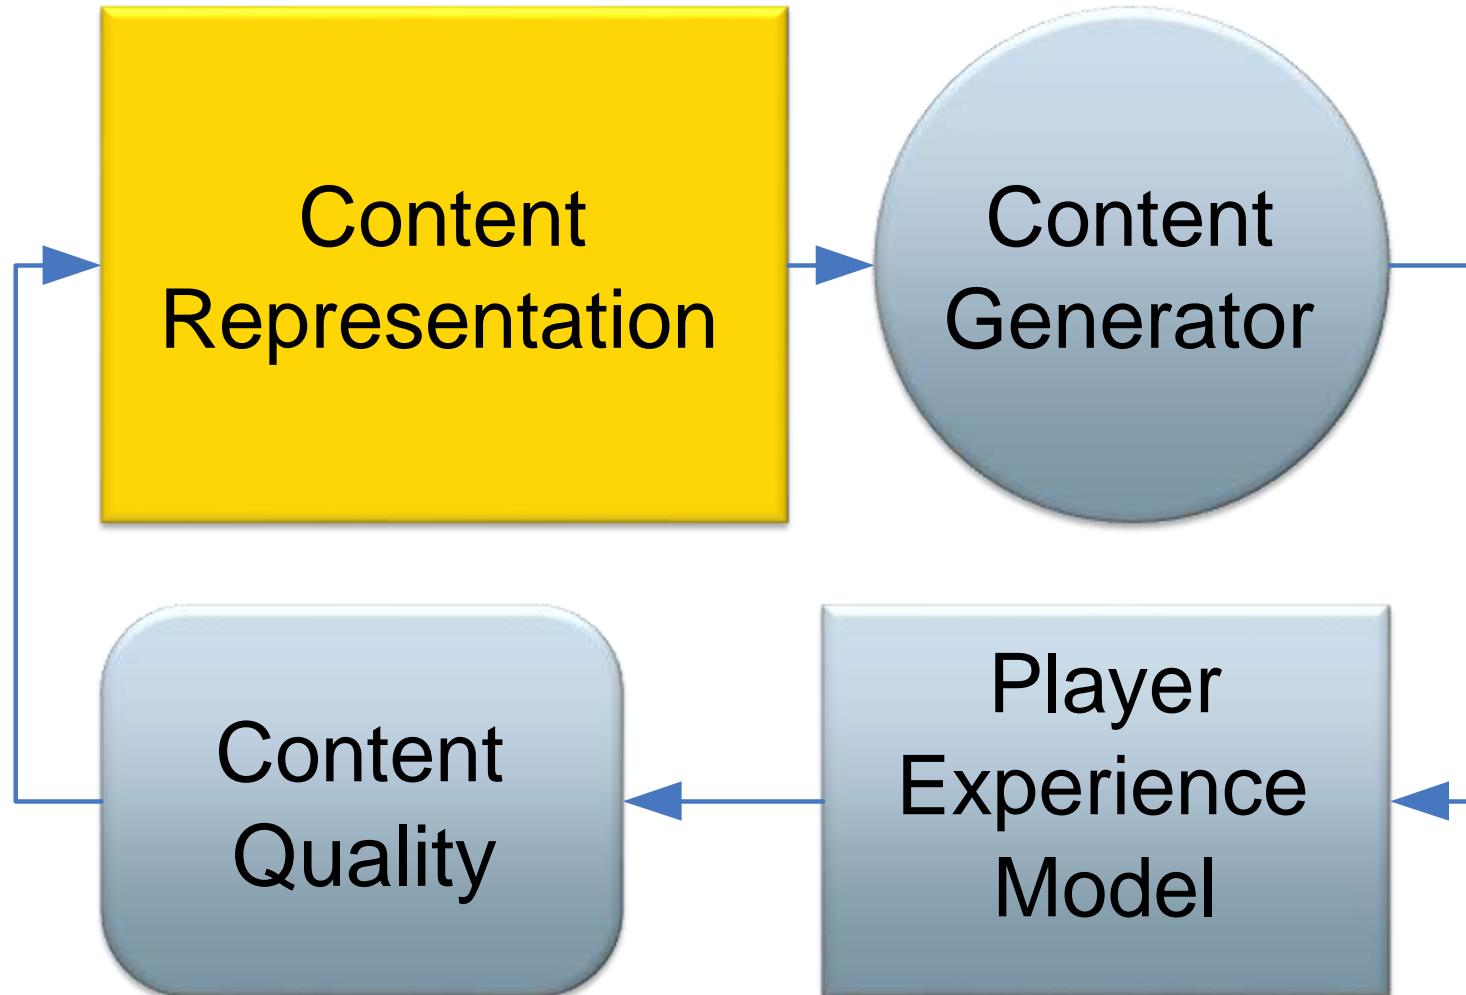
Interactive

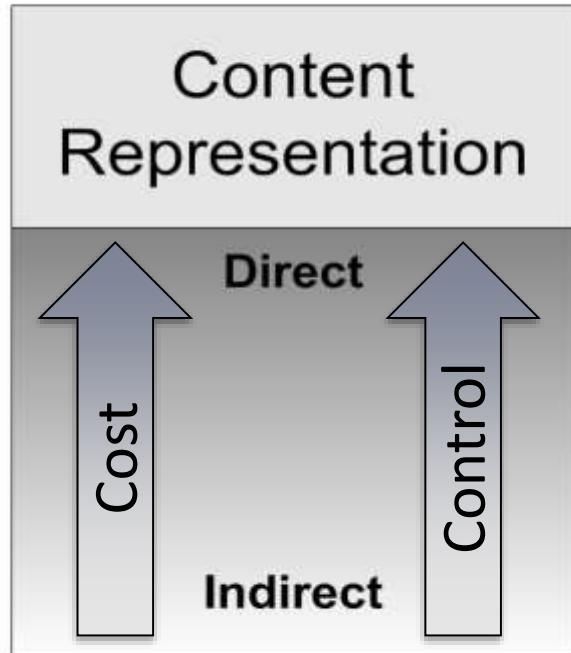


Content Quality

Direct	Simulation-based	Interactive
Theory-driven vs. Data-driven	Static vs. Dynamic	Implicit vs. Explicit

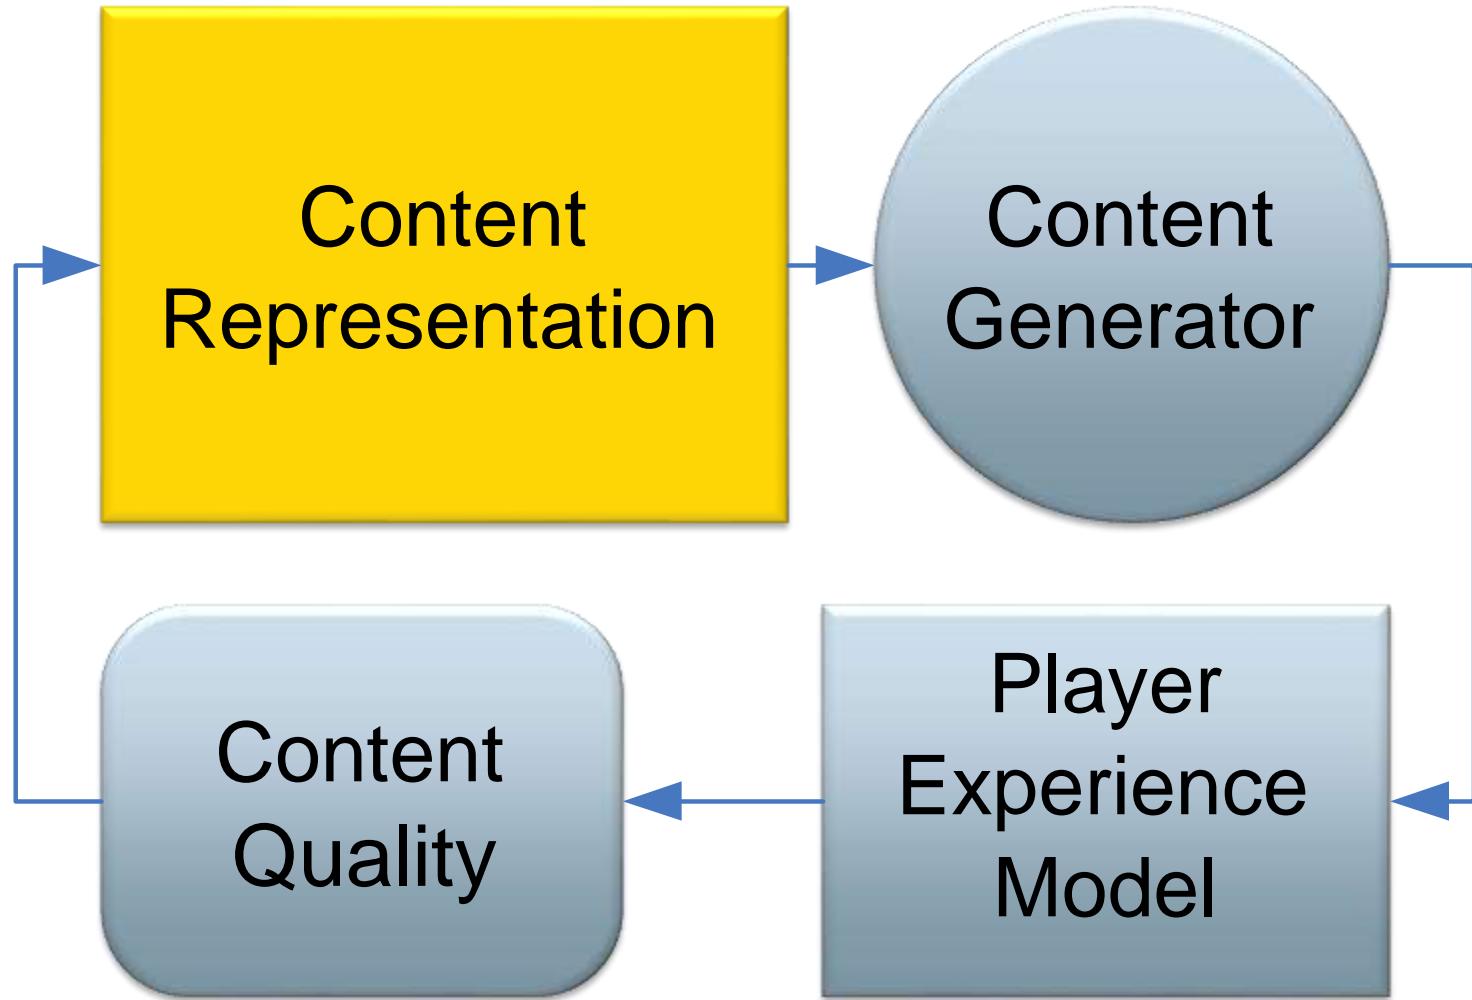


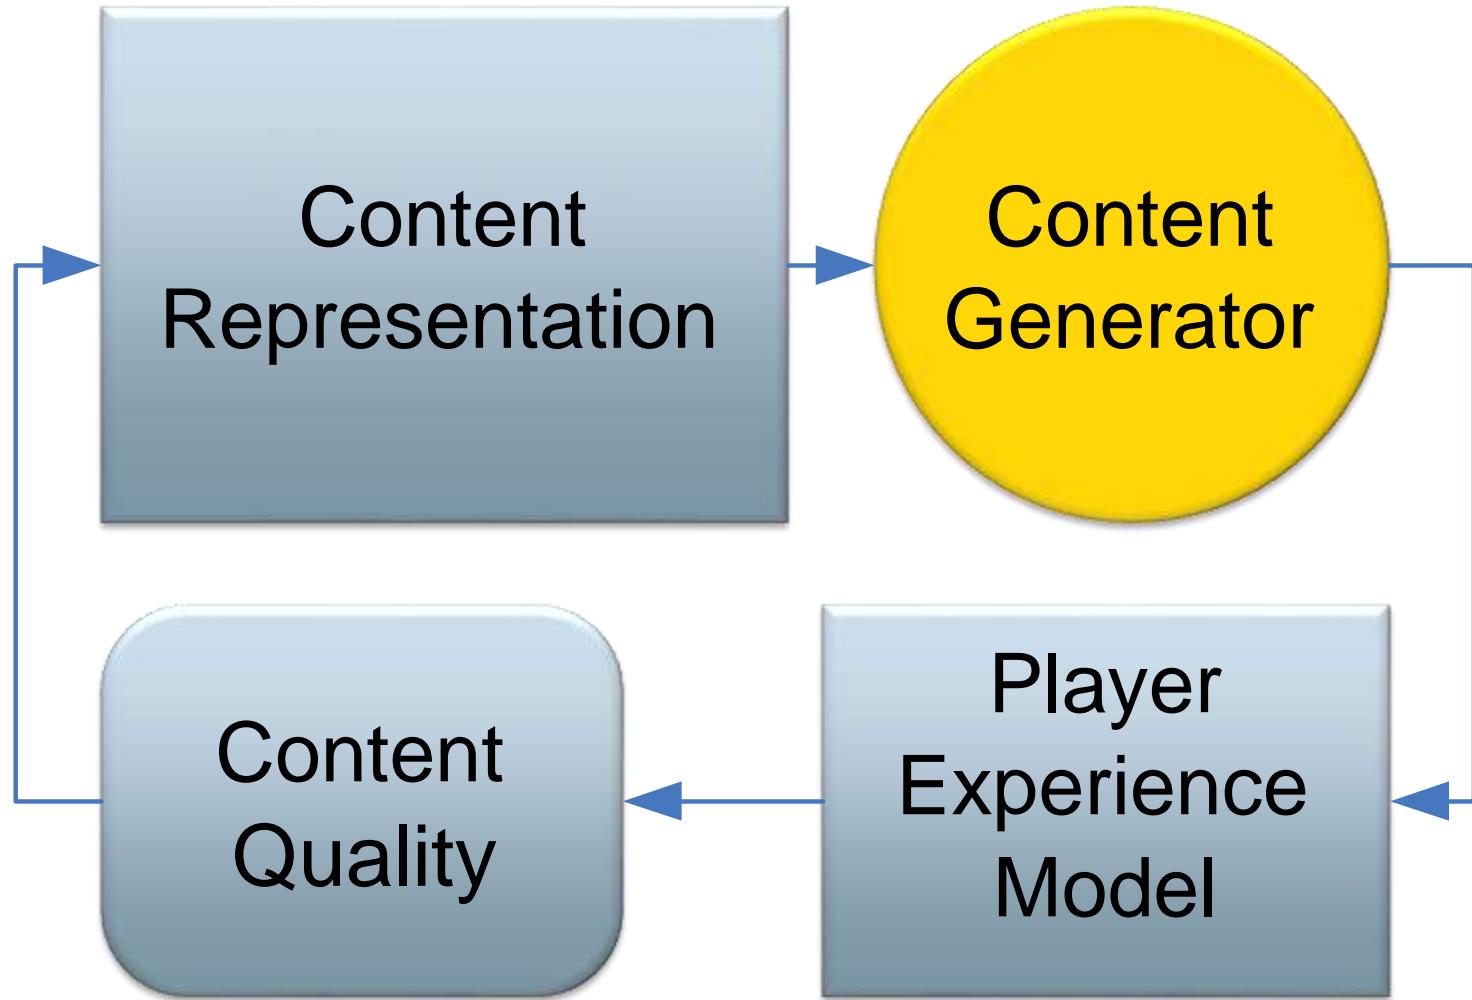


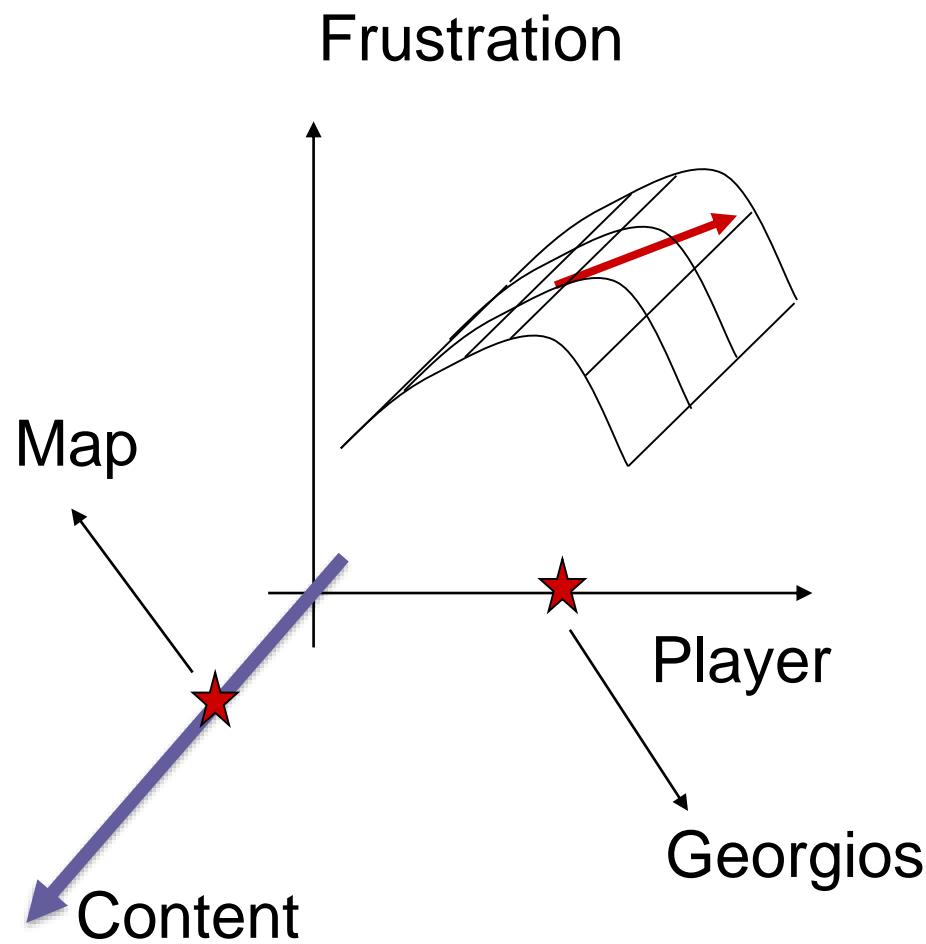


- Grid
- Position and orientation of walls
- Patterns of walls and floor
- Number of rooms and doors
- Random seed

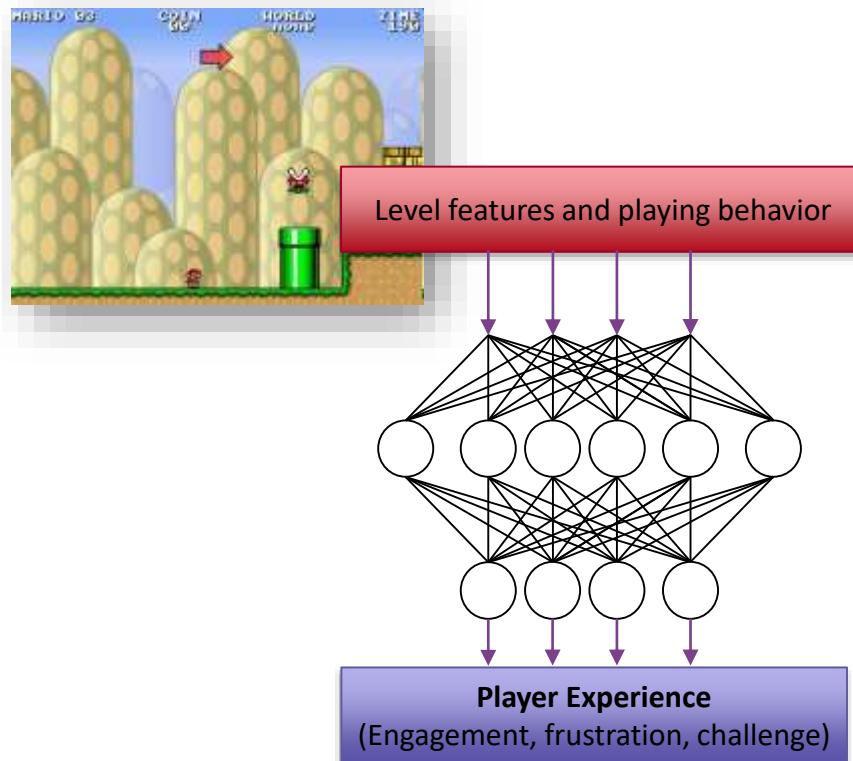






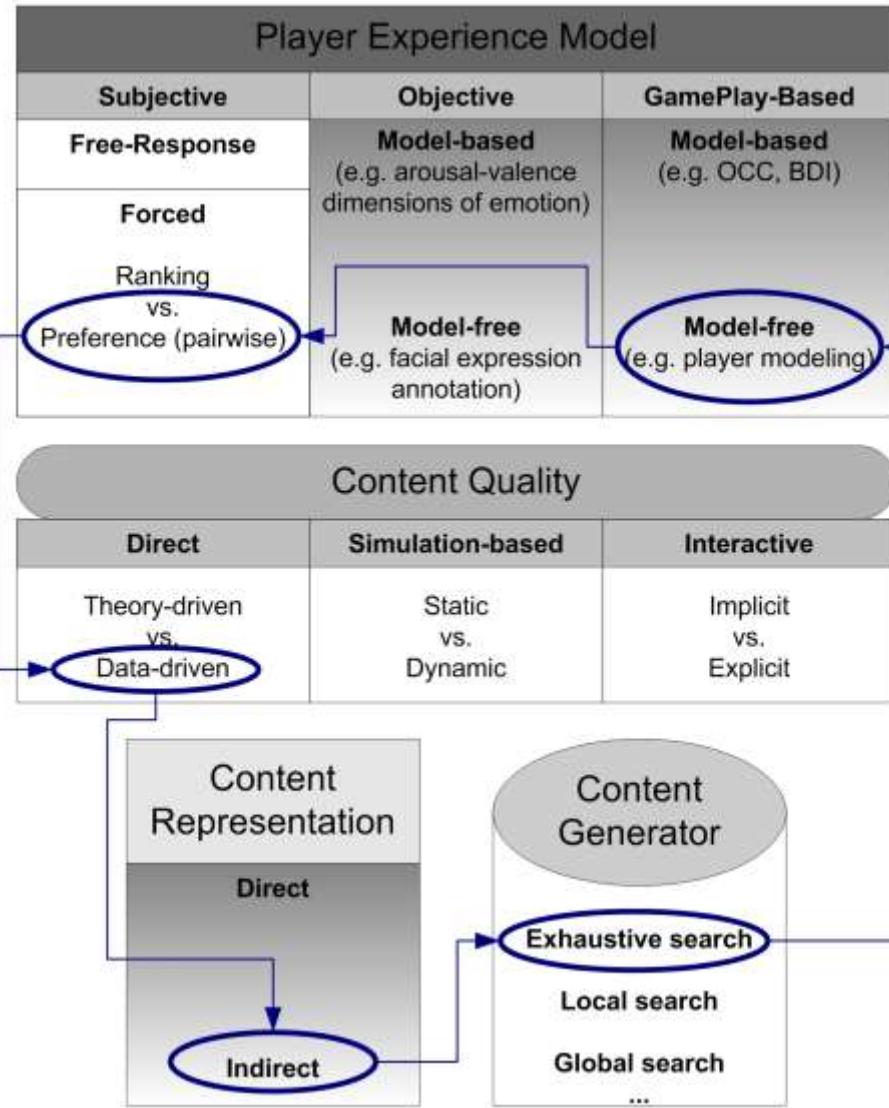


A Super Mario Bros Example



- **Model-free**, gameplay-based PEM (pairwise preferences as outputs)
- **Direct** (data-driven) evaluation function
- **Indirect** content representation (a few parameters)
- Search for good content via **exhaustive search!**

Shaker, N., Yannakakis, G. N., & Togelius, J. (2010, October). Towards Automatic Personalized Content Generation for Platform Games. In *AIIDE*.



Other EDPCG Examples



Sonancia

Lopes, Liapis, and Yannakakis: "**Sonancia: Sonification of Procedurally Generated Game Levels,**" in Proceedings of the ICCC workshop on Computational Creativity & Games, 2015



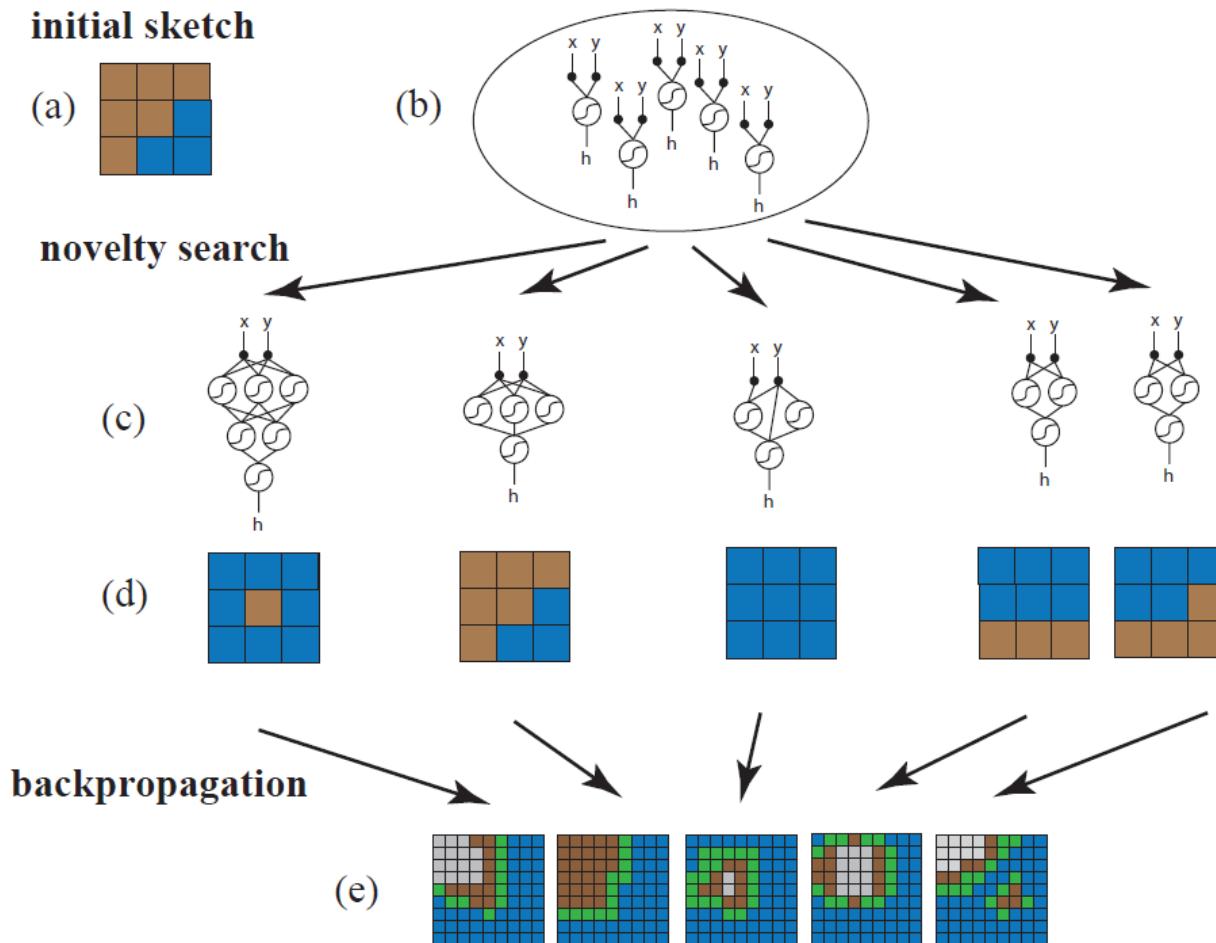
EDPCG for Serious Games: *Village Voices*

Khaled and Yannakakis "Village Voices: An adaptive game for conflict resolution", in Proc. of FDG, pp. 425-426



Sentient World: Hybridizing Evolution and Gradient Search

Liapis et al. "Sentient World: Human-Based Procedural Cartography," *EvoMusArt*, 2013.



Sentient World: Hybridizing Evolution and Gradient Search

Liapis et al. "Sentient World: Human-Based Procedural Cartography," *EvoMusArt*, 2013.

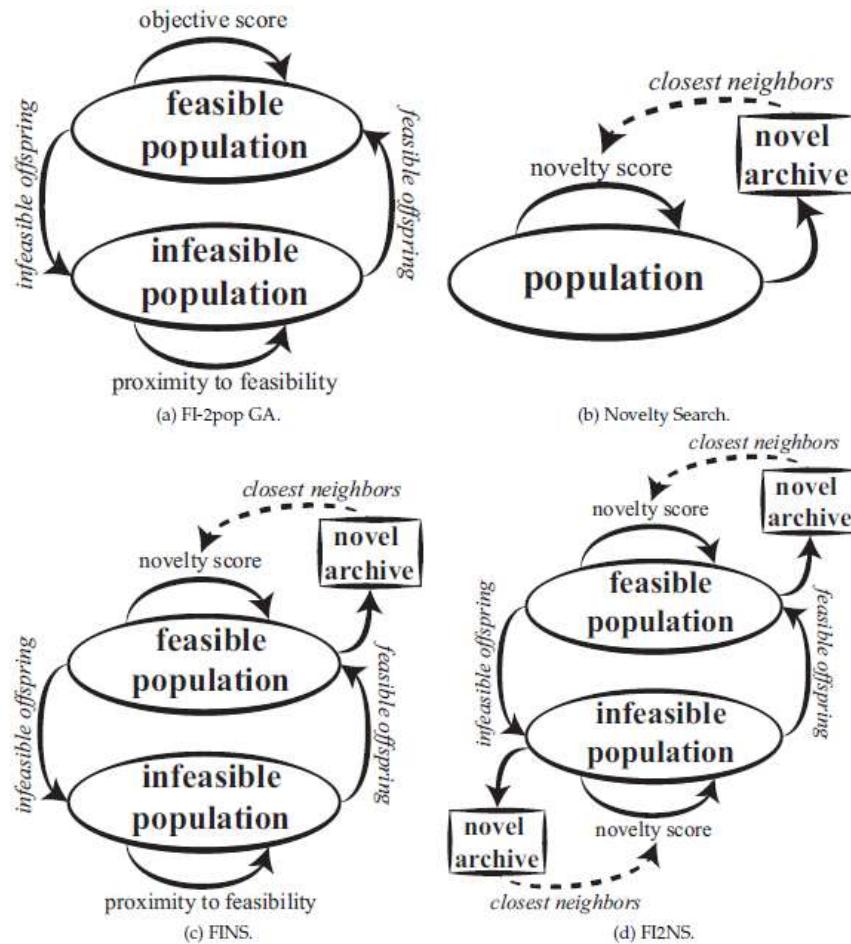


New Map



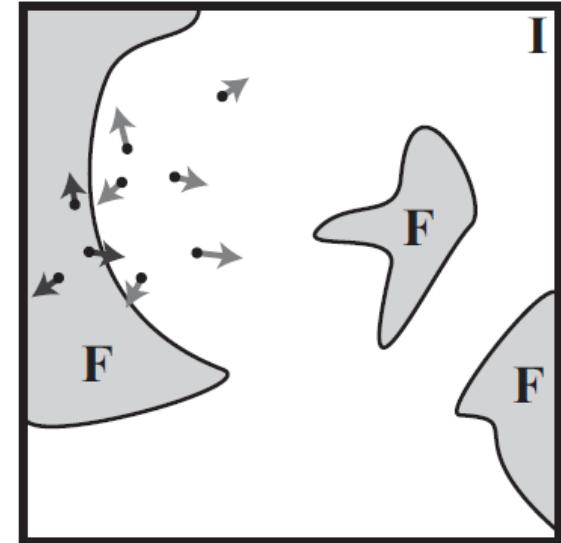
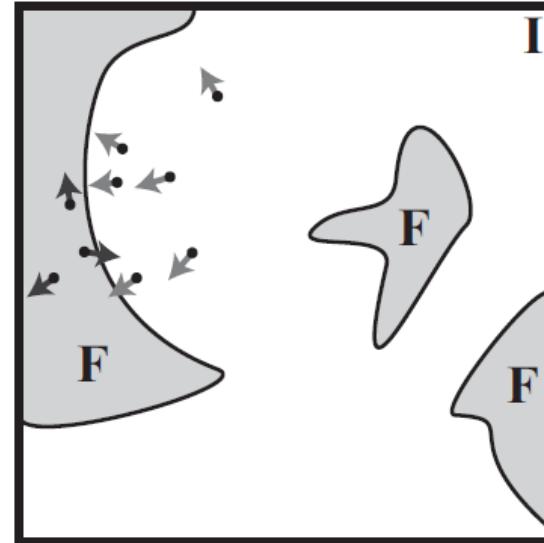
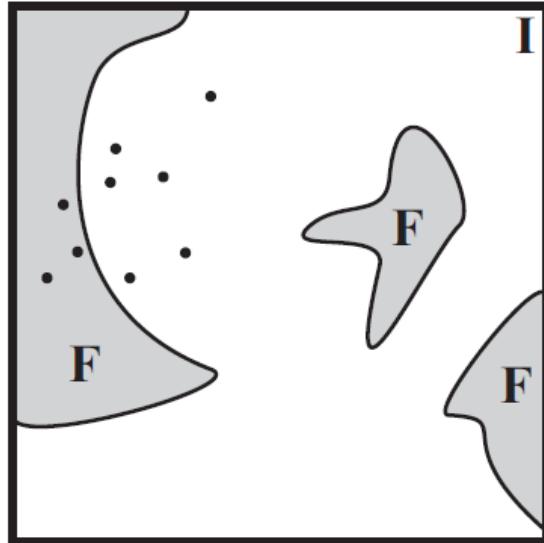
Constrained Novelty Search for PCG

Liapis, Yannakakis and Togelius, **Constrained Novelty Search: A Study on Game Content Generation**, Evolutionary Computation, 21(1), 2015, pp. 101-129



Constrained Novelty Search for PCG

Liapis, Yannakakis and Togelius, **Constrained Novelty Search: A Study on Game Content Generation**, Evolutionary Computation, 21(1), 2015, pp. 101-129



Sentient Sketchbook



Georgios N. Yannakakis, Antonios Liapis and Constantine Alexopoulos:

"Mixed-Initiative Co-Creativity," in Proc. of the ACM Conference on Foundations of Digital Games, 2014.

- Map Sketches (strategy game, dungeon, FPS level)
- Multiple solutions evolved & shown in real-time
- Fitnesses on area influence, exploration and balance... and novelty
- Constraints on playability handled with FI-2pop GA

Welcome to Sentient Sketchbook

Read the tutorial

Draw Small Map

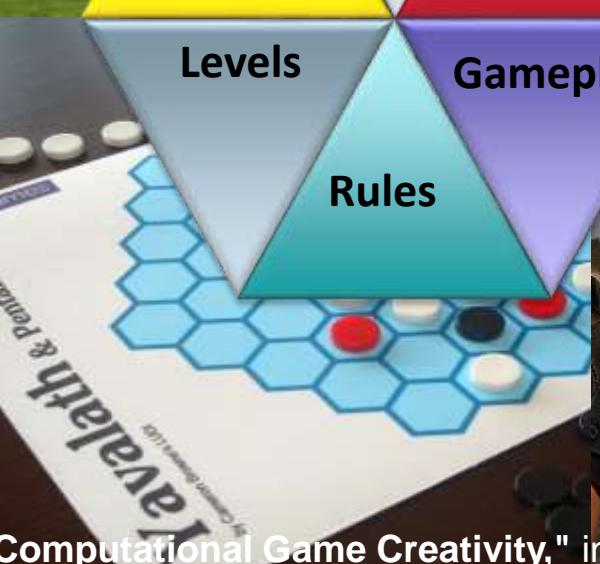
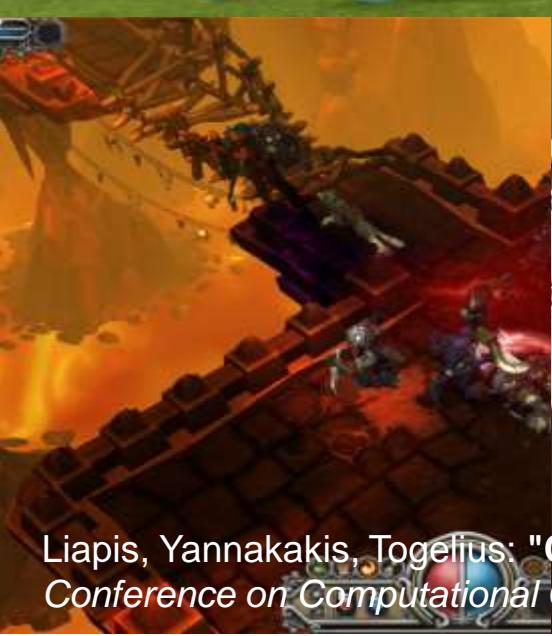
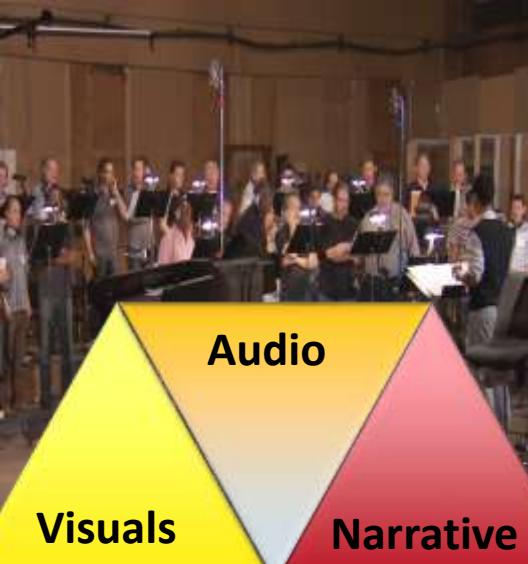
Draw Medium Map

Draw Large Map

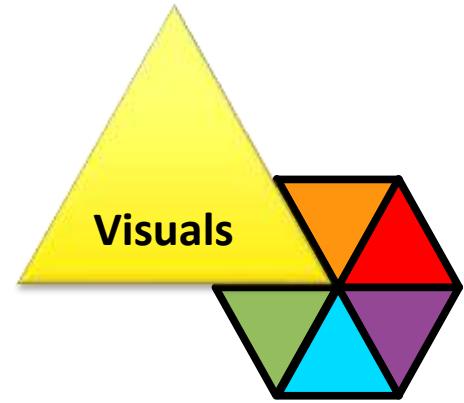
A user can select among a predefined set of map sizes.
Map size determines the number of allowed bases and resources.

What Could be Generated?





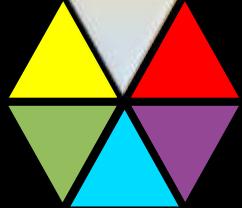
Liapis, Yannakakis, Togelius: "**Computational Game Creativity**," in *Proceedings of the Fifth International Conference on Computational Creativity*, 2014.



**All of the music you hear in
this video was composed in
real time by the program.**



Audio



Brown, Daniel. "Mezzo: An adaptive, real-time composition program for game soundtracks." *Proceedings of the AIIDE Workshop on Musical Metacreativity*. 2012.



Narrative

...ce, are you angry at Trip?!



look, over there! it's a \$100 laptop!



Orkin, J., and Roy, D. 2007. **The restaurant game: Learning social behavior and language from thousands of players online**. Journal of Game Development 3(1):39–60.



Gameplay



Press Esc to exit full screen mode.

MARRY
ME

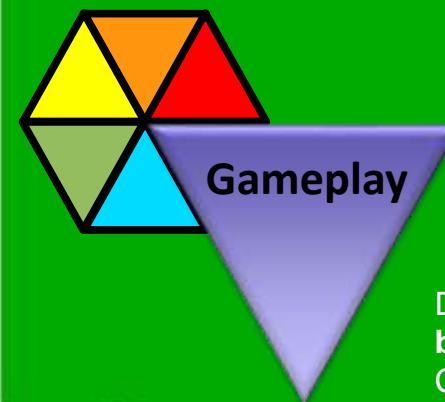


Gameplay

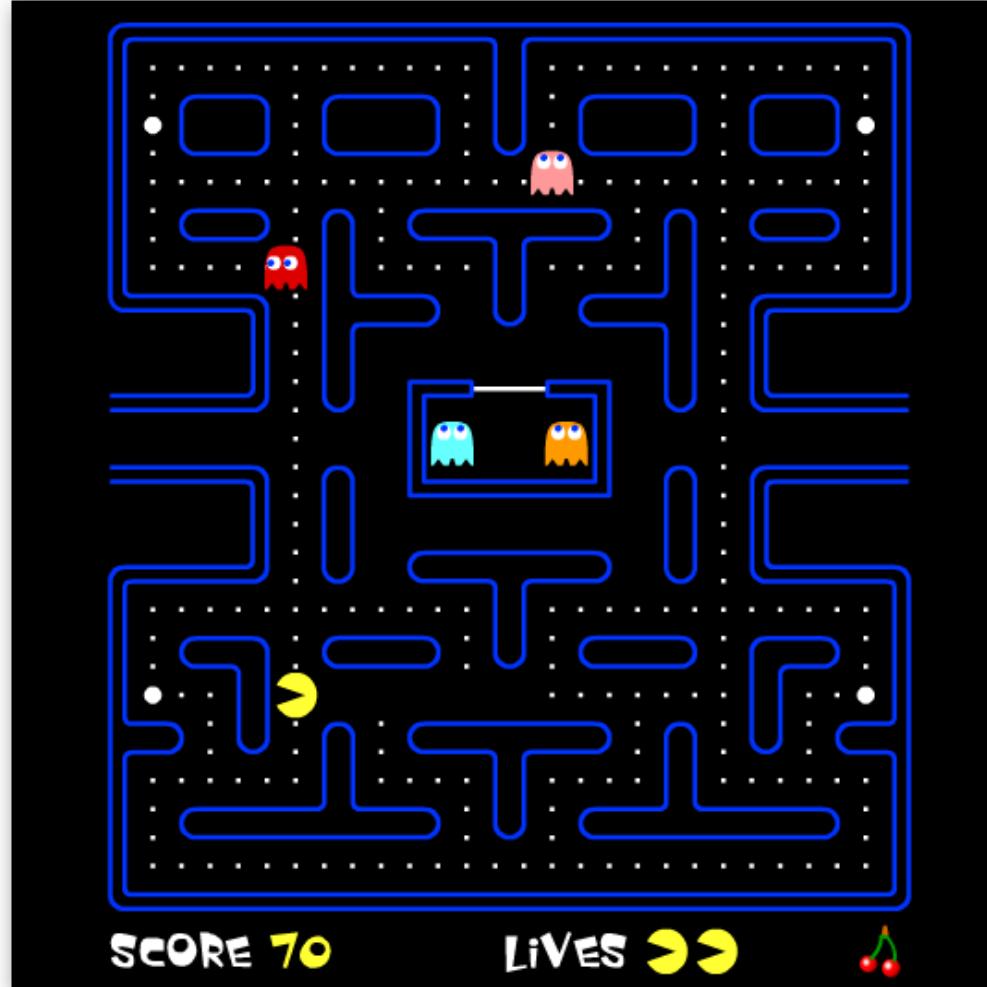
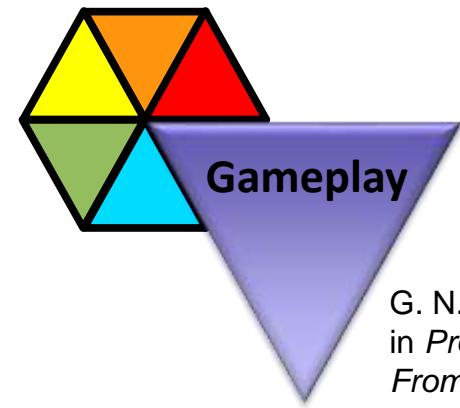


00:16 { 3}{ 3} 0-0
play

182



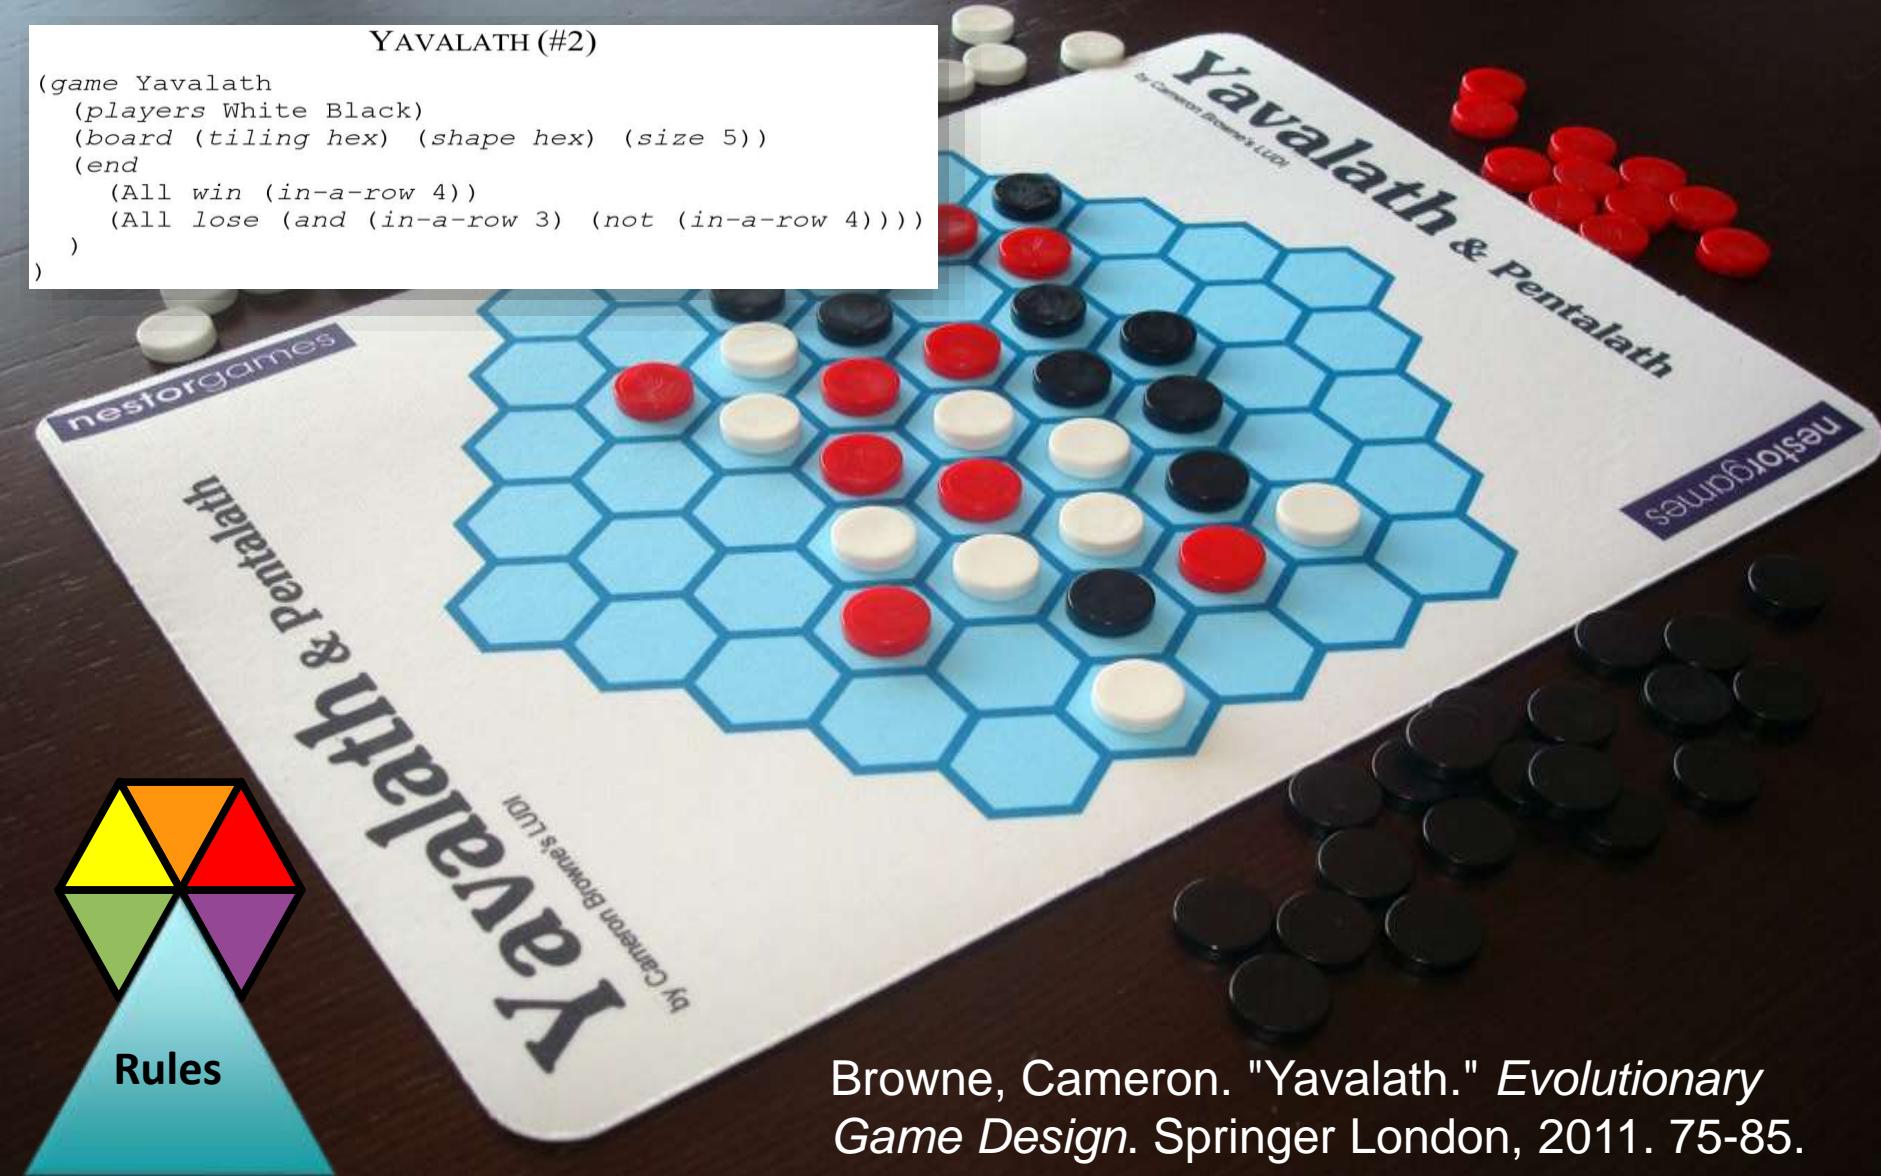
Denzinger, J.; Loose, K.; Gates, D.; and Buchanan, J. 2005. **Dealing with parameterized actions in behavior testing of commercial computer games**. In Proc. of the IEEE Symposium on Computational Intelligence and Games , 37–43.



G. N. Yannakakis, and J. Hallam, “**Evolving Opponents for Interesting Interactive Computer Games**,” in *Proceedings of the 8th International Conference on the Simulation of Adaptive Behavior (SAB’04); From Animals to Animats 8*, pp. 499-508, Los Angeles, CA, USA, July 13-17, 2004. The MIT Press.

YAVALATH (#2)

```
(game Yavalath
  (players White Black)
  (board (tiling hex) (shape hex) (size 5))
  (end
    (All win (in-a-row 4))
    (All lose (and (in-a-row 3) (not (in-a-row 4))))
  )
)
```



Browne, Cameron. "Yavalath." *Evolutionary Game Design*. Springer London, 2011. 75-85.



Level
design

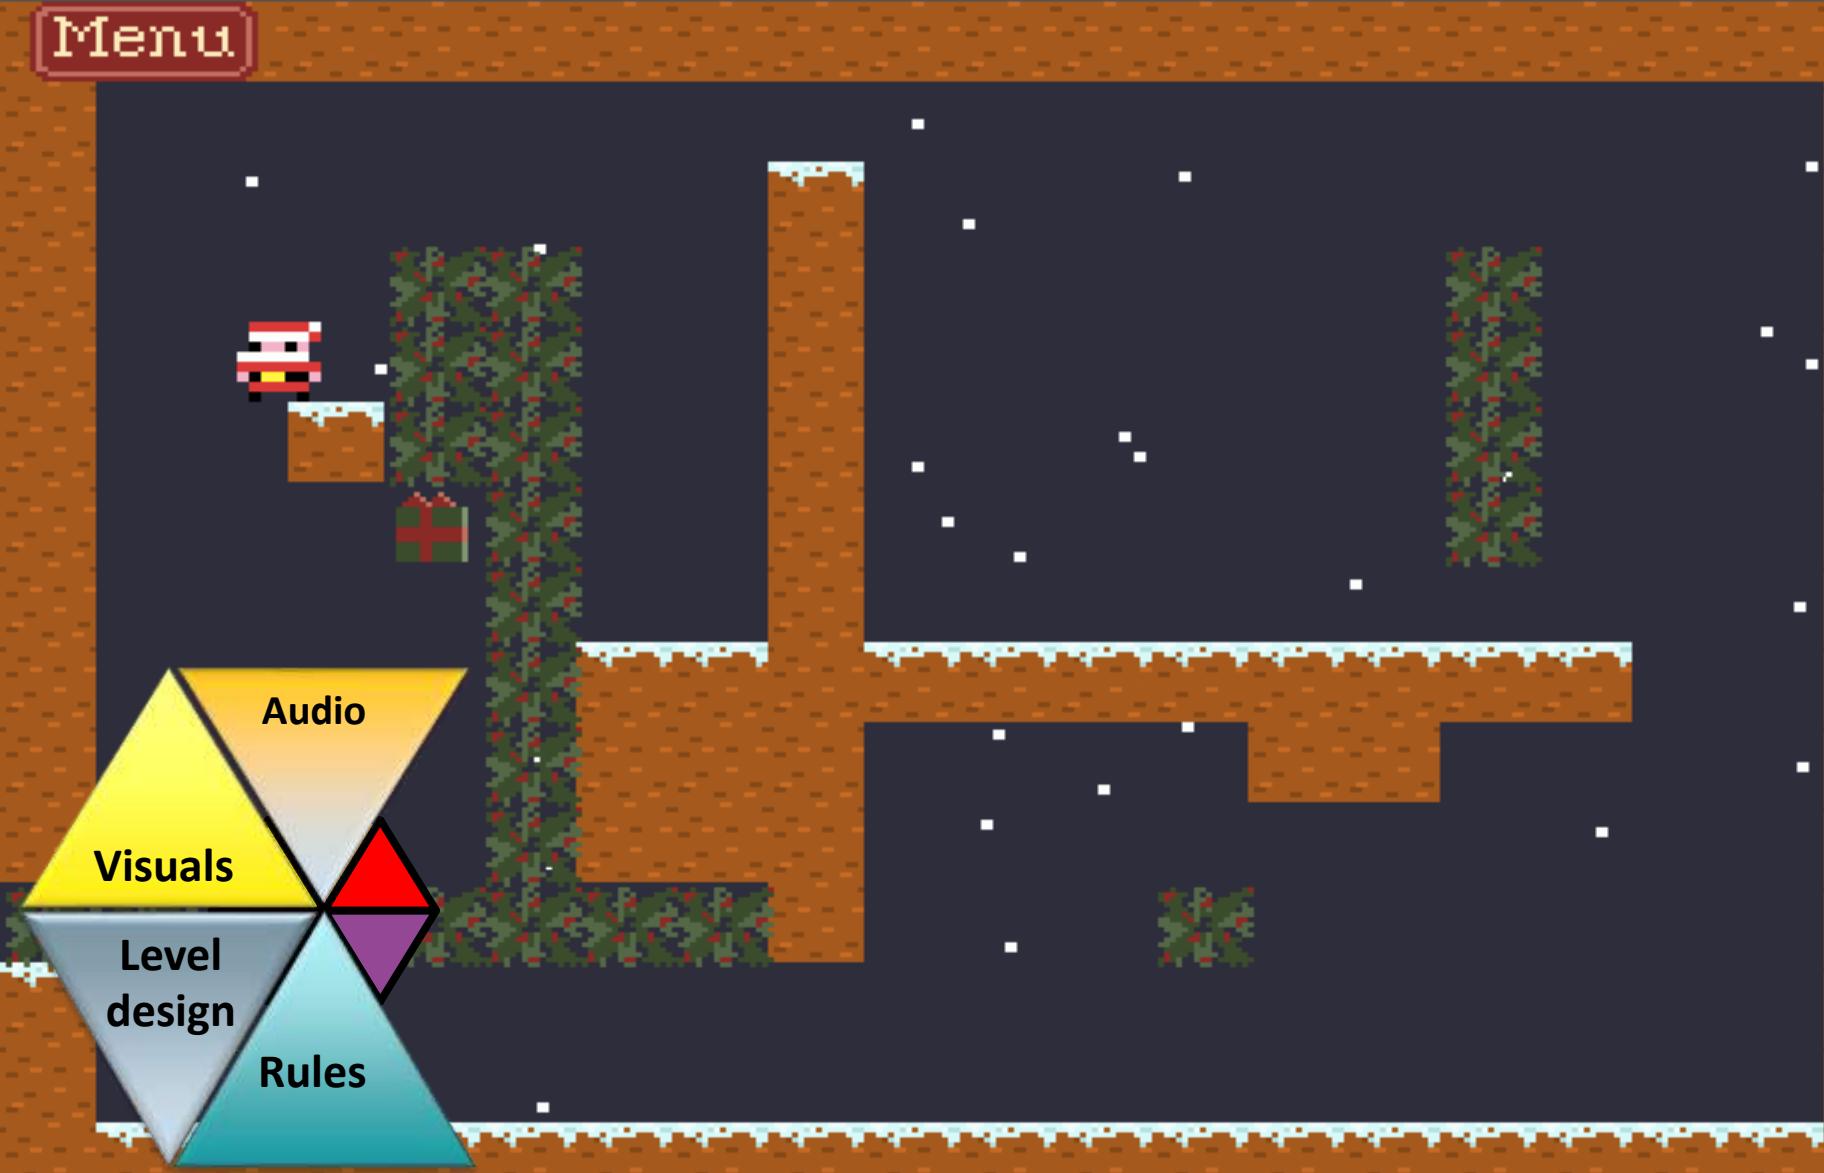
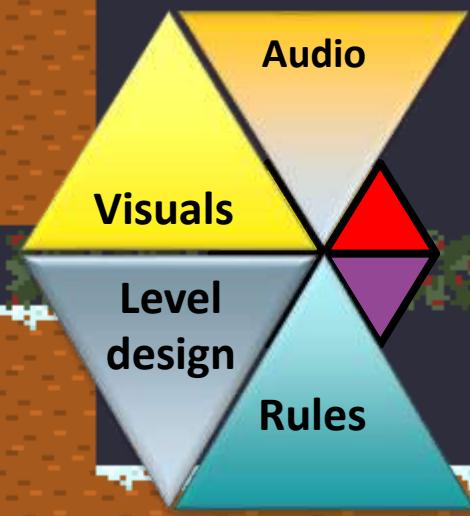


Complete Game Generation – Orchestration

Antonios Liapis, Georgios N. Yannakakis, Mark J. Nelson, Mike Preuss and Rafael Bidarra: "**Orchestrating Game Generation**" in Transactions on Games, 2019.

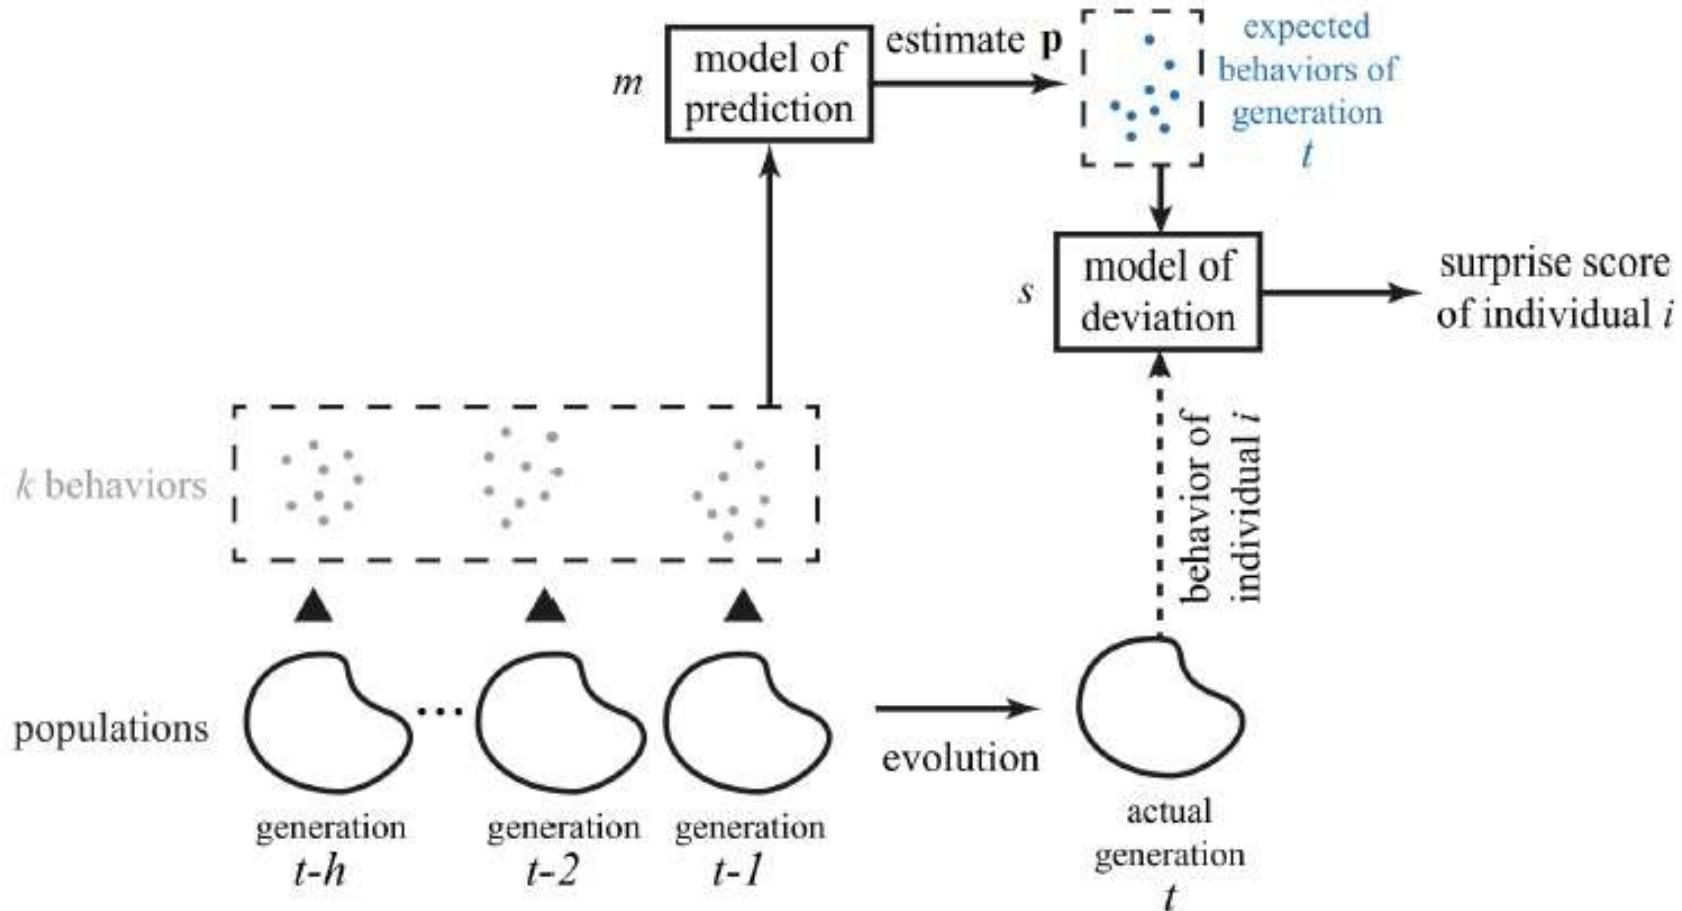


Menu



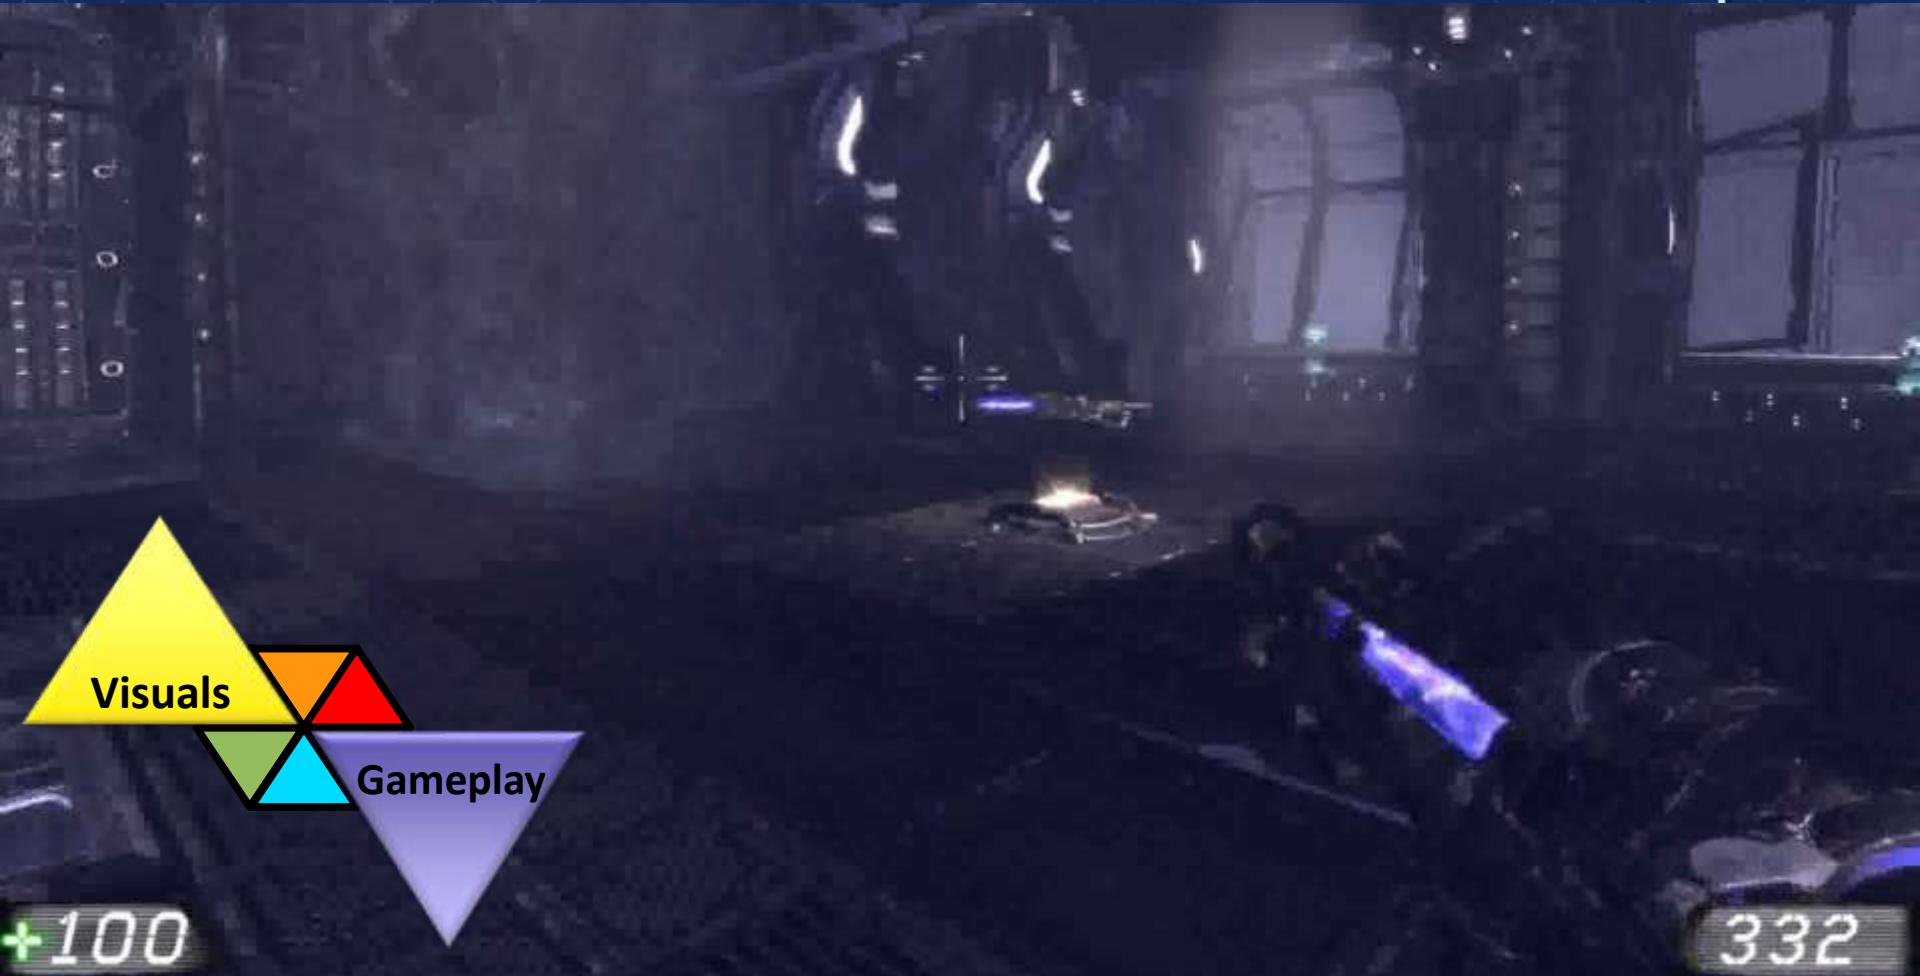
From Novelty Search to Surprise Search

Gravina, Liapis, and Yannakakis: "Surprise Search: beyond Novelty and Objectives" in Proceedings of GECCO, 2016



Surprising Weapons!

Gravina, Liapis and Yannakakis: "Constrained Surprise Search for Content Generation," in Proceedings of the CIGE Conference on Computational Intelligence and Games (CIG). 2016.

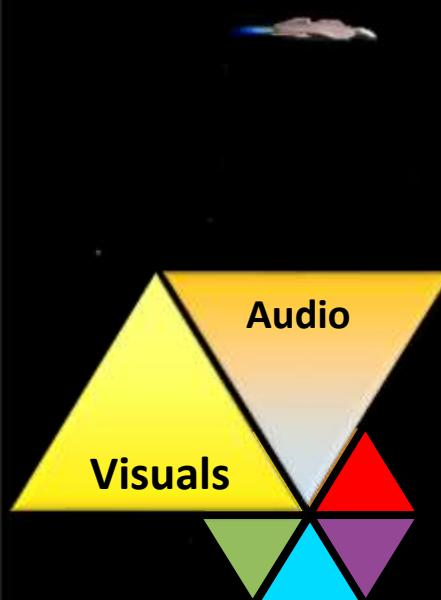


+100

332

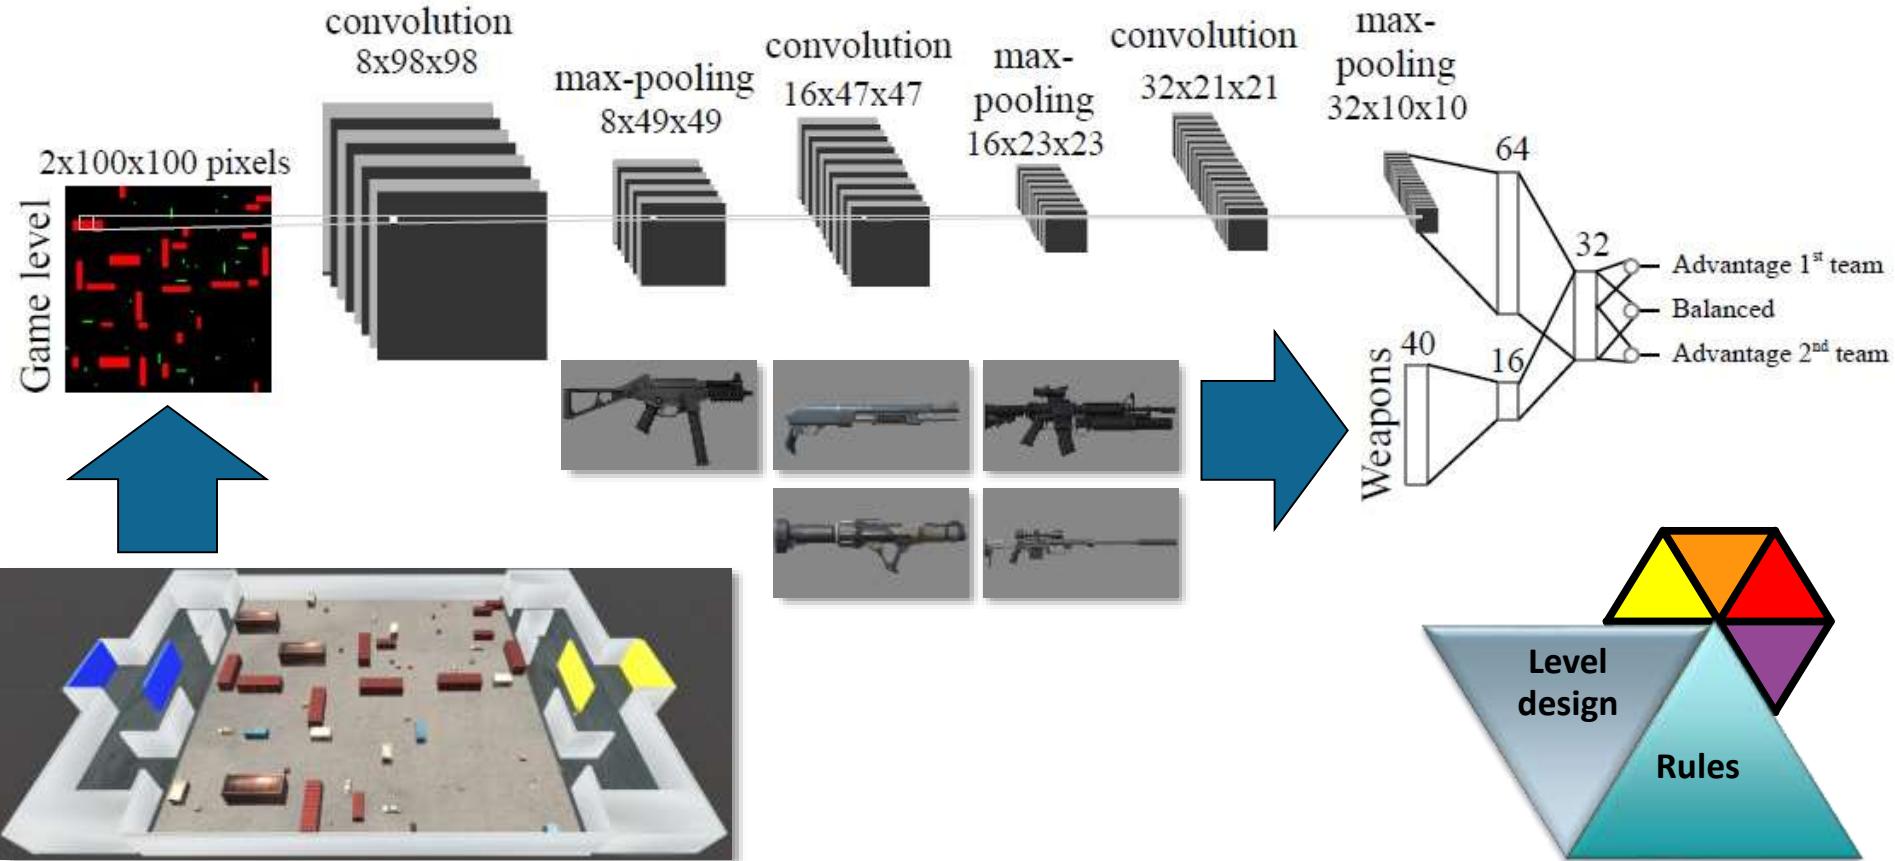
AudioInSpace: From Music to Weapons!

Hoover, Cachia, Liapis, Yannakakis, "AudioInSpace: A Proof-of-Concept Exploring the Creative Fusion of Generative Audio, Visuals and Gameplay," in EvoMusArt, 2015



Orchestrating Level and Game Design

Karavolos, Liapis, and Yannakakis: "Learning Patterns of Balance in a Multi-Player Shooter Game," in *Proceedings of the Foundations on Digital Games*, 2017.



Evaluating Content Generators

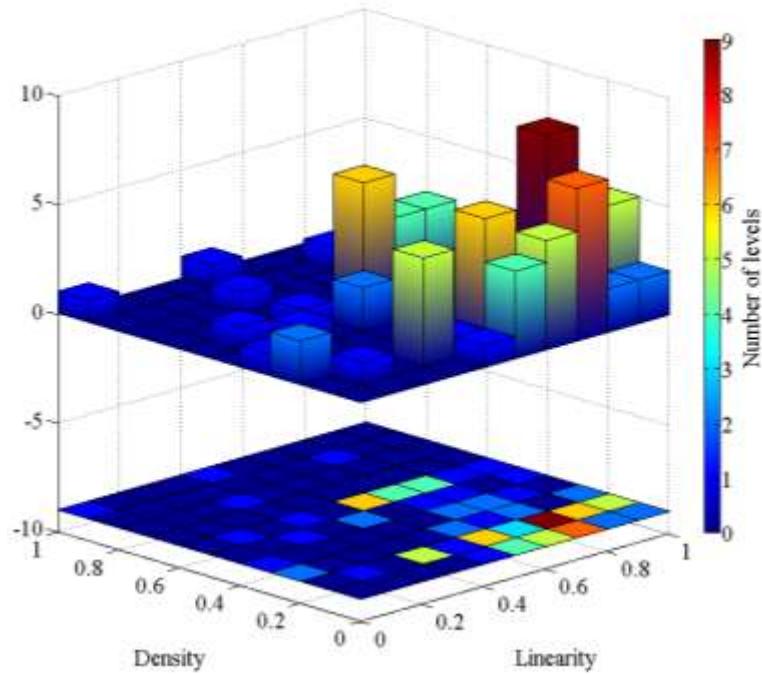


How Can we Evaluate a Content Generator?



Generally speaking there are three ways:

- Visualization (e.g. expressive range)
- AI (playtesting / personas)
- Human players (testing, QA, annotations)





About the Book

Welcome to the Artificial Intelligence and Games book. This book aims to be the first comprehensive textbook on the application and use of artificial intelligence (AI) in, and for, games. Our hope is that the book will be used by educators and students of graduate or advanced undergraduate courses on game AI as well as game AI practitioners at large.

Final Public Draft

The final draft of the book is available [here!](#)

Thank you!

gameaibook.org